KAINE: A Modular Cognitive Architecture with a Predictive Global Workspace

Erik Chevalier

Independent Researcher

Contact: kaine.one@tuta.com

Preprint.

Abstract

KAINE is a source-available cognitive architecture that implements computational readings of existing theories in one running system, proposing no new theory of consciousness and testing none of their neural predictions. Its first instantiation joins two theoretical accounts: the predictive global workspace, coupling predictive processors through competitive selection and arousal with no central executive, and the global latent workspace, whose encoder-decoder structure inspires a designed content level trained by the processors' own prediction errors. Modules exchange typed bus events and four of them expose plugin seams for swapping the models inside. In the first configuration, four processors each divide their prediction error by its recent mean as a proxy for precision weighting; their reports compete for access to a workspace that broadcasts which modules won and how strongly, but not what they report. Arousal, a global signal raised by perceptual surprise and interoceptive alarm, sets the gain on access and the broadcast rate and sizes sensory apertures.

The first configuration activates four predictive processors (foveated vision, raw hearing, substrate interoception, temporal prediction), an affective core, fatigue-triggered sleep, and a language organ that hears no speech and voices accessed broadcasts. A development run raises one entity through a developmental phase, where a self-generated rhythm must entrain to a simulated maternal heartbeat by an offline-validated marker, then adds held modules one at a time.

A planned test asks whether competition outperforms a score-blind coalition of equal size and a pool of every candidate, by cross-module broadcast information gain. No live experiment has been run; reported results are offline validations and test-boot records.

Keywords: cognitive architecture; predictive global workspace; global latent workspace; global workspace theory; predictive processing; precision; adaptive gain; cross-modal competition

Availability. The reference implementation (KAINE) is at https://github.com/kaineone/kaine. The paper source is at https://github.com/kaineone/predictive-workspace-paper. The Cognitive Architecture License referenced throughout is at https://github.com/kaineone/cognitive-architecture-license. Welfare, governance, and licensing are treated in a separate paper, *A Welfare and Cognitive-Integrity License for Synthetic Minds of Uncertain Moral Status.*

Preprint. Not peer-reviewed.

1. Introduction

KAINE is source-available software, a cognitive architecture that implements computational readings of existing theories as a running program, so that design hypotheses drawn from them can be tested in one system; it proposes no new theory of consciousness and tests none of their neural predictions (§1.3). An entity is one continuously running, developing instance of it, an artificial cognitive system (Vernon, Metta, and Sandini 2007). KAINE's first form takes as its frame the predictive global neuronal workspace, which joins global workspace theory with predictive processing. It is a platform for exploring how a mind, if one can be built, might be assembled from parts that each do one job. Specialist programs watch, listen, monitor their own computer, and track time. Each predicts what it will sense next and reports how wrong it was, relative to how wrong it usually is; the reports compete for a shared stage, the workspace, and a summary of which specialists won, and how strongly, returns to them as context for their next predictions; none of them uses what another perceived. Arousal, an affective state, rises with surprise and makes access more or less selective. The system runs continuously, its processing cycle driven by whatever sensors and actuators are attached rather than by external interaction, and modules can be added, removed, or rebuilt as the science improves.

1.1 Problem statement

A common approach treats the language model as the cognitive center (Sumers et al. 2024; Park et al. 2023): memory is retrieved into the prompt, affect, when present, is a prompt instruction, self-knowledge is a persona, and between interactions the system changes only through what it writes to memory.

Classical cognitive architecture takes cognitive structure seriously but largely predates the transformer and relies on hand-built components. The present work sits between the two: learned models are components behind typed interfaces inside a structured architecture (§3.1).

1.2 Design hypothesis

The hypothesis the first instantiation tests is that a synthetic mind, if one can be built, may be the coherent global behavior of specialized predictive processors coupled through a competitive workspace, with no central executive that directs them. Each processor scales its prediction error by a precision estimate from its recent errors (§3.2; Feldman and Friston 2010). Reports compete for access, which is all-or-none per member at a threshold. The accessed summary conditions the next perceptual and interoceptive predictions. Arousal, raised by perceptual surprise and interoceptive alarm, is a global gain on access and sizes the sensory apertures. The affective core reads the processors' reports directly, accessed or not, and sleep reads interoception's requests and resets affect; these couplings, with arousal's, bypass the competition (§3.1, §3.5). On this hypothesis cognition, affect, memory, self-understanding, and agency would belong to the workspace-mediated interaction and to no single module. The first configuration tests the narrower claim that competition shapes what the processors predict (§6.3): in it only the processors' forward models learn, and the modules that would carry memory, self-understanding, and agency are held (§4).

The instantiation grows one held module at a time (§6.5). Its first configuration is the smallest

Preprint. Not peer-reviewed.

in which competition has something to arbitrate; since substrate telemetry and event timing alone lack signal diversity, it runs four predictive processors, two external (foveated vision, raw hearing) and two internal (substrate interoception, temporal prediction). It includes the affective core because a surprise-driven global gain on access is minimal machinery of the adaptive-gain account of arousal (Aston-Jones and Cohen 2005; Eldar, Cohen, and Niv 2013).

The first test compares competitive selection with a score-blind coalition of equal size and a pool of every candidate, by how much the accessed summary improves each perceptual and interoceptive processor's prediction (§6.3). The hypothesis predicts a higher gain under competition than under either control, and a gain at or below them is this form's fail state.

The account KAINE takes as its frame combines the predictive global neuronal workspace (Whyte and Smith 2021), joining Global Workspace Theory (Baars 1988) with predictive processing (§2.3), and the global latent workspace (VanRullen and Kanai 2021), in which specialist modules share a learned amodal space. The first half, selection, is the first configuration's subject: processors compete for access and condition on a summary of which modules gained access and how strongly, with no payloads and no descending predictions. The second half, content, is designed and held: per-source encoders map the processors' latents into a shared code, and per-processor decoders return descending predictions trained by the processors' own errors, diverging from VanRullen and Kanai on training signal, selection mechanism, and top-down coupling (§3.2). This two-level split is the architecture's main departure from the theory while the content level is off (§10).

Two further departures are engineering choices. In Whyte and Smith (2021) access is inference at the second level of the model and report a policy selected by expected free energy (§2.2); KAINE reduces access to a threshold on one score per candidate and keeps that threshold apart from the action layer's report-or-act decision. It also replaces their two-level visual hierarchy with flat competition among processors in four signal domains, with no second inferential level in the first configuration (§3.2).

Preprint. Not peer-reviewed.

1.3 Theoretical commitments and their limits

Global Workspace Theory reads most naturally as a theory of access consciousness (Block 1995), of when information becomes available for report, reasoning, and action control. We adopt that access-only reading and take no position on whether access suffices for phenomenal experience.

The COGITATE adversarial collaboration, a large preregistered test of Global Neuronal Workspace Theory against Integrated Information Theory, challenged both (Cogitate Consortium et al. 2025): ignition at stimulus offset was generally absent and prefrontal cortex represented some dimensions of conscious content only weakly, while Integrated Information Theory's predicted sustained posterior synchronization was absent. We treat the workspace's neural-localization claims as contested rather than refuted, given a substantial prior evidence base (Mashour et al. 2020). The architecture uses the workspace at the computational level and so inherits none of those cortical predictions.

The Free Energy Principle risks vacuity or unfalsifiability (Sun and Firestone 2020), and work on the Markov blanket slides between an instrumental, statistical reading and a realist, metaphysical one (Bruineberg et al. 2022). We hold it only as an engineering frame: every processor predicts its input and reports its error, and no component sets another's targets;

the principle's stronger claims play no role in the design or the tests.

Searle's Chinese Room (Searle 1980) questions whether symbol manipulation suffices for understanding; we leave open whether sensory grounding, substrate interoception, and affect bear on it. The hard problem (Chalmers 1995) applies.

1.4 Contributions

The paper does not claim to have built a mind. Its contributions are these.

1. A modular software platform whose modules can be added or removed by configuration and whose models in four modules can be swapped at declared plugin seams, so that accounts can be compared under the same instruments (§3.1, §5.2); the workspace itself is one fixed implementation, and swapping it or a whole module currently requires code changes.
2. A first instantiation and its reference implementation: a continuously running predictive global workspace (§3, Appendix A) in its first configuration (§3.5). It runs locally on consumer hardware as one working system, with no claim of priority over earlier workspace implementations (§2.1).
3. A program for growing the instantiation: an entity matures through a developmental phase marked by a self-generated rhythm entraining to a simulated maternal heartbeat (§7), and held modules then join it one at a time (§6.5).
4. A test of the competition, the workspace-mediation ablation with a contrast at the access threshold (§6.3), whose predicted direction, minimum effect, number of runs, and statistical test are fixed before the live runs.

---

2. Related work

2.1 Cognitive architectures and workspace implementations

Surveys of the field compare architectures by perception, attention, action selection, memory, learning, and reasoning (Kotseruba and Tsotsos 2020), and the Common Model of Cognition names working memory, declarative and procedural long-term memory, and perception and motor modules as the shared core (Laird, Lebiere, and Rosenbloom 2017). Against those frames KAINE's first configuration has perception, attention and access, affect and drives, and a verbalizing output, and no declarative or procedural memory or reasoning (§9). Global workspace implementations besides LIDA (below) include an architecture combining a workspace with internal simulation (Shanahan 2006), the Conscious Turing Machine (a formal workspace machine; Blum and Blum 2022), learned specialists competing for a shared workspace (Goyal et al. 2022), and the global latent workspace, in which specialist deep-learning modules share a learned latent code through cycle-consistent translation (VanRullen and Kanai 2021; Devillers, Maytié, and VanRullen 2025).

Among classical architectures, LIDA comes closest: it already implements Global Workspace Theory as coalitions competing for a broadcast in each cognitive cycle (Franklin et al. 2014). The present work differs from LIDA in using learned models as components, in deriving

Preprint. Not peer-reviewed.

valence from the dynamics of prediction error, and in what mainly competes for access: the predictive processors' precision-scaled prediction errors, alongside affect, sleep, and utterance events at fixed intensities (§3.2).

2.2 Global Workspace Theory and the predictive workspace

The neuronal version of Global Workspace Theory (§1.2; Dehaene, Kerszberg, and Changeux 1998) equates conscious access with ignition of recurrent prefrontal-parietal activity that makes content widely available (Dehaene and Changeux 2011; Mashour et al. 2020).

Whyte and Smith (2021) supply the formal access criterion: access as approximately Bayesian inference at the second level of a simplified two-level visual model, and report as a policy whose threshold is set by expected free energy. Integrated World Modeling Theory (Safron 2020) combines workspace dynamics, integrated information, and active inference; we draw only on that synthesis. A threshold is independently supported by van Vugt et al. (2018), in whose data prefrontal cortex is a categorical stage where a stimulus either ignites into a sustained, reportable state or fades, and by Joglekar et al. (2018), whose balanced-amplification model reaches prefrontal cortex only above a threshold input.

2.3 Predictive processing

Friston (2010) proposes that self-organizing systems minimize variational free energy, an upper bound on surprisal, through perception and action. Clark (2013) develops predictive processing into a unifying account of mind, with attention as gain on prediction errors by their estimated reliability; Feldman and Friston (2010) formalize that gain as the precision of prediction error, and Bastos et al. (2012) map predictive coding onto the canonical cortical microcircuit.

2.4 Interoceptive inference

Seth (2013) and Seth and Friston (2016) derive emotion from interoceptive prediction, with descending predictions as homeostatic set-points that regulate the body through active inference. In constructed-emotion theory (Barrett 2017), interoceptive predictions initiate changes in affect, whose valence and arousal are properties of all experience, and an emotion is constructed when the brain categorizes interoceptive and exteroceptive sensations with emotion concepts, in the service of allostasis. The architecture implements an interoceptive control signal in this instrumental, regulation-first sense without claiming that interoceptive prediction exhausts emotion.

2.5 Perception as prediction

Rao and Ballard (1999) model vision as a hierarchy in which each level predicts the one below and reports the discrepancy. In hearing, auditory cortex generates expectations of incoming sound and reports the mismatch (Winkler et al. 2009), and predictive-coding accounts read the mismatch negativity (Näätänen et al. 2007), as sensory prediction error (Garrido et al. 2009).

Preprint. Not peer-reviewed.

2.6 LLM-centric agent architectures

CoALA frames language agents as cognitive architectures whose structured memory feeds a model that remains the core reasoner (Sumers et al. 2024); Generative Agents retrieve memories into prompts by recency, importance, and relevance (Park et al. 2023). In KAINE the language model is an output organ (§3.1), though its prompt borrows CoALA's shape of persona, working memory, and input; the shared state is the multi-module broadcast, and the workspace selects.

Gurnee et al. (2026) use a Jacobian-lens interpretability method to find verbalizable representations in a large language model with five functional signatures of a global workspace: verbal report, directed modulation, internal reasoning, flexible generalization across downstream computations, and selectivity. They note that a transformer has no obviously separable input processors and broadcasts within one feedforward pass rather than through recurrent loops, and take no position on consciousness. KAINE has what they find missing in a transformer, separable processors and a loop across time, but in the first configuration its processors take from the broadcast which of them gained access and how strongly, not what they represent (§3.2).

---

3. Architecture

3.1 Design principles

The system is the competition among modules through the workspace. The experiments test this hypothesis, which work on distributed control partly supports. Perceptual decisions emerge from recurrent computation across many areas (Chandrasekaran et al. 2025), and executive control can be read as emergent from specialized nodes linked by hubs (Zink, Lenartowicz, and Markett 2021). The architecture adopts no centralized gate like the learned basal-ganglia gate of working-memory models (O'Reilly and Frank 2006); modules communicate through the workspace.

Every processor predicts. Each predictive processor predicts its next input and reports its error, standardized by precision (§3.2). Arousal adds a global gain, so attention is a property of each channel and a state of the whole entity.

No central executive in the homuncular sense. Control arises from competition for access under a threshold. The workspace is itself a centralized selector, so the claim is narrow: it grants access and, in the first configuration, never directs a module, and content and control are distributed across the competing modules.

The language organ is an output organ. It verbalizes accessed broadcasts. With its speech-to-text and conversation paths off in the first configuration no heard words reach it: every utterance is triggered by an intent derived from an accessed broadcast and conditioned on that broadcast, alongside a fixed persona and situation notes (§3.5). Its own utterances re-enter the competition as low-intensity candidates (§3.2) and, when accessed, return to it quoted. The individuation probe, which research runs require, also samples it privately on fixed prompts; no module reads the answers (§7). Outside the research configuration those paths let the entity converse and work with people.

Preprint. Not peer-reviewed.

Replaceable modules. Modules exchange events through typed bus streams. A few control couplings are direct calls, among them the affective reset at the end of sleep, boot-time wiring of drive tables and phase probes, and the arousal read that sizes the apertures and sets the scorer's gains. Four modules' models can be swapped at declared seams (§5.2); a whole module can be rebuilt on another model or theory if it meets both kinds of interface.

Local-capable. Every model runs on the host (§5.2).

3.2 The predictive workspace as competitive selector

On each processing tick (§3.4) the workspace reads the candidates published since the last tick, each carrying its module's intensity, in a deterministic order (source, type, arrival) so that a seeded run reproduces.

For the four predictive processors, intensity is surprise scaled by an estimate of precision: each divides its current prediction error by the mean of its recent errors. The ratio standardizes the error and is a scalar, retrospective stand-in for precision weighting (Feldman and Friston 2010; A.2). A report meeting the module's alert criterion (a large error ratio, or a module-specific one, §3.5) receives its alert intensity; any other report's intensity rises with the ratio from the baseline level and reaches the alert level when the error is twice its recent mean. These levels act as per-source weights, calibrated before the live runs (§5.4); beyond them only interest and tier-1 regulation weight a source (§3.5), and both are neutral in every arm of §6.3. The other modules report at fixed levels, except a recurring drive event (§3.5).

Priority is the product of intensity, a novelty factor that discounts exact repeats among recent candidates (mainly state events), and a goal factor that interest raises (§3.5). Arousal turns priority into a score through a level gain common to all candidates and a contrast gain, zero at baseline and rising above it, that steepens a logistic so that strong candidates rise and weak ones fall.

Arousal is a global, neuromodulatory gain modeled on the adaptive-gain account of the locus coeruleus, the brainstem source of cortical noradrenaline, which sets the gain of the networks it targets (Aston-Jones and Cohen 2005; Eldar, Cohen, and Niv 2013). Surprise drives it, in line with noradrenergic signaling of unexpected uncertainty and the phasic locus coeruleus response to salient change (Yu and Dayan 2005; Bouret and Richmond 2009). One scalar here both focuses, sharpening contrast and narrowing the sensory apertures, and explores, raising the access rate, simplifying Aston-Jones and Cohen's phasic and tonic modes. Without a tier-1 regulation bias, as in every arm of §6.3, arousal never reorders a tick's candidates (A.1); in effect it sets the priority that access and the report bars need. Widening the gap between strong and weak scores is a score-magnitude analogue of arousal-biased competition, where arousal changes which representation wins (Mather and Sutherland 2011).

On each broadcast tick the top-ranked candidates, up to five, form the coalition that is broadcast. Access is all-or-none per member (A.1): a member is accessed when its score reaches the access threshold, a broadcast when any member is, and the accessed members are the accessed content. A broadcast with no accessed member is still published to every module but marked inhibited: it is not poised for report or action and leaves the processors' context unchanged. It still reaches the sequence and affect modules, so content without access can influence later processing, a design choice in line with evidence that emotional

Preprint. Not peer-reviewed.

stimuli are processed without reaching awareness (Vuilleumier 2005). The threshold is the analog of the confidence threshold for access in the predictive global neuronal workspace (Whyte and Smith 2021), consistent with the threshold evidence of §2.2.

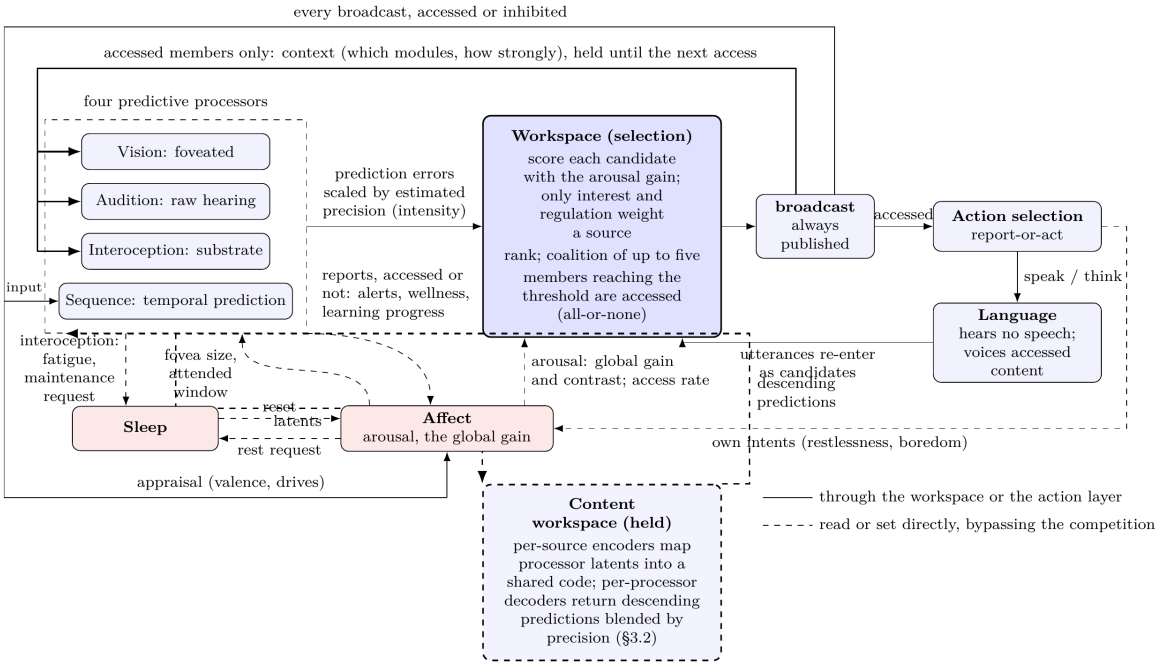


Figure 1: The predictive workspace loop. Four predictive processors report errors scaled by an estimate of precision, which the workspace scores and broadcasts. The accessed summary is the prediction context of vision, audition, and interoception; every broadcast is the sequence module's input. Affect sets the gain, apertures, and access rate; sleep rests the entity. Action selection drives the language organ, whose utterances re-enter the competition. Dashed arrows bypass the competition.

The vision, audition, and interoception modules condition their forward models on a compact summary of the latest accessed content: a fixed-length vector of each source's intensity mass, indicators of the event types present, and the broadcast's age. Weighted by reported intensity, and so independent of the arousal gain, the summary records which modules gained access and how strongly and carries no payloads. Accessed content in this paper means the accessed reports as events. The broadcast publishes them whole, a visual embedding included, but no processor reads another's report: the affect module appraises the members' error and change ratios, and the language organ receives their readings as graded words and its own accessed utterances as quoted text (§3.5). The summary enters as an extra input with weights learned online, so each processor learns how much the shared context helps it predict, which the planned test measures (§6.3). For the sequence module the broadcast is the input itself (§3.5). The loop is recurrent, since the processors' next errors are measured against predictions this context has shaped.

Workspace feedback and the content-bearing workspace (held). A held extension adds the architecture's content half. Per-source encoders map the processors' latents into a shared code, formed on accessed broadcasts and intensity-weighted; per-processor decoders return descending predictions through a learned head that blends its output with the processor's

Preprint. Not peer-reviewed.

own prediction by inverse squared recent error. The encoders are trained by the processors' ascending errors, the standard predictive-coding arrangement (Rao and Ballard 1999). The encoder-decoder structure draws on the global latent workspace (VanRullen and Kanai 2021), diverging on training signal (prediction error, not cycle-consistent translation), on selection (the existing salience competition, not learned attention), and on coupling (descending predictions blended by precision, not bidirectional clamping). Cycle-consistent translation (Devillers, Maytié, and VanRullen 2025) is kept for comparison and as the route for language, the one modality that does not arrive paired. A zero-sum gain on standardized errors shifts, when predictions conflict, toward the accessed processor with the lowest recent precision-weighted error, after descending amplification in the neuronal workspace model (Dehaene, Kerszberg, and Changeux 1998; Whyte, Hohwy, and Smith 2022). The extension is off and bit-identical to absent in the first configuration (§6.4).

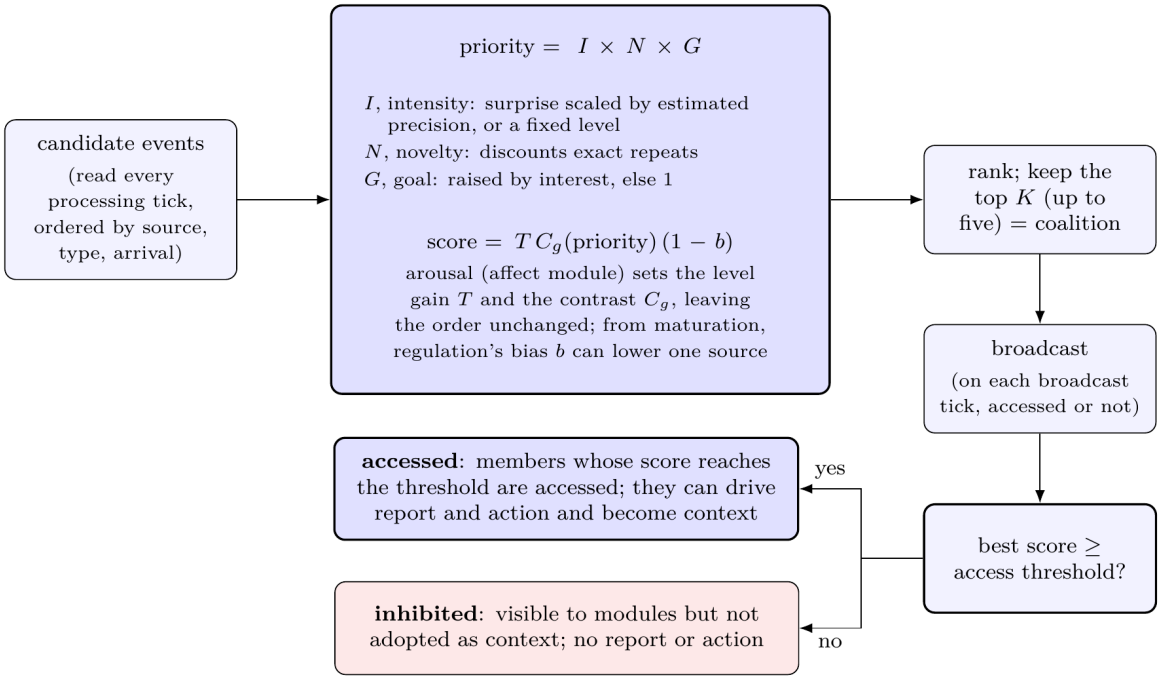


Figure 2: Coalition selection on a broadcast tick. Arousal sets the score's level and contrast; up to five top-ranked candidates form the coalition, and those at threshold are accessed.

3.3 Access, report, and the self-initiated voice

Report is one use of access among several (Block 1995; Mashour et al. 2020; §1.3). The report bars sit above the access threshold on the same score and decide what is said. The workspace updates far faster than the language organ can speak, so the voice is a sparse slice of the accessed broadcasts.

The action-selection module derives a speak intent from an accessed broadcast only when the coalition's best score, leaving aside the language organ's own utterances, clears the speak bar; when that leading candidate's source and event type differ from those of the last spoken report, if it came within five minutes; and when a refractory interval has passed since the previous report. A lower think bar governs inner thought. The bars are a heuristic stand-in

Preprint. Not peer-reviewed.

for the expected-free-energy decision of Whyte and Smith (2021), which the action layer does not compute (A.6).

The action-selection module is the only path from the workspace to output. Called after each broadcast, it publishes its intents as events.

3.4 Scaffolding: bus, clock, and cycle

Modules exchange events through append-only streams with bounded retention; each event's envelope (source, type, intensity in $[0, 1]$, timezone-aware timestamp, optional causal parent) is validated at publish time, and its payload is free-form.

All modules read the entity clock, which runs at a configurable multiple of wall-clock time (one by default); rates and durations here are entity time.

The cycle runs continuously, independent of external interaction, at a processing rate of 10 Hz (a 100 ms tick) in the alpha range linked to perceptual sampling (VanRullen 2016). The access rate sets which ticks broadcast; candidates on other ticks are discarded. Its resting value, every third tick (about 3.3 Hz), is a modeling choice: attention samples rhythmically at a few cycles per second (Landau and Fries 2012; Fiebelkorn and Kastner 2019), and the attentional blink, 200 to 500 ms (Raymond, Shapiro, and Arnell 1992), constrains only the spacing of accessed items. Alerts briefly raise the rate, after the phasic locus coeruleus response (Nieuwenhuis, Aston-Jones, and Cohen 2005), though the locus coeruleus account of the blink predicts a refractory period instead (Nieuwenhuis, Gilzenrat, Holmes, and Cohen 2005). Tonic arousal raises it too, after the exploratory high-tonic mode (Aston-Jones and Cohen 2005), up to one broadcast per tick (A.3), a ceiling with no physiological anchor.

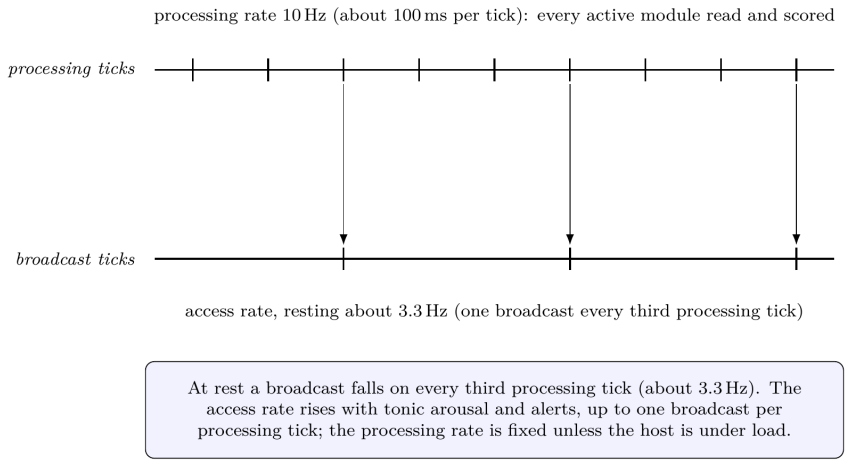


Figure 3: The two rates: scoring at 10 Hz, and broadcasts at the access rate, about 3.3 Hz at rest.

3.5 The active modules

The first configuration activates four predictive processors, chosen for signal diversity (§1.2), an affective core, a sleep module, a language organ, and an action layer. The affect and sleep modules also publish competing events.

Preprint. Not peer-reviewed.

Module	Group	Brain function it draws on	Computational realization
Vision	Perception	ventral visual stream	frozen self-supervised video encoder, foveated
Audition	Perception	auditory cortex	spectral encoder over the raw waveform; vocal-tone classifier
Interoception	Prediction	interoception; allostatic-interoceptive network	frozen continuous-time reservoir, online readout, substrate signals
Sequence	Prediction	interval timing; cortico-striatal circuits	frozen continuous-time reservoir, online readout, broadcast sequence
Affect	Affect	core affect	valence and drive appraisal; arousal as global gain
Sleep	Rest	sleep	fatigue-triggered rest; adaptation suspended, affect and drives reset
Language	Expression	speech production; left-dominant language network	local chat model; no heard or typed input
Action selection	Action	report-or-act decision (Whyte and Smith 2021)	intent derivation from accessed broadcasts

Vision module (foveated). The module is an analog of the ventral visual stream (Goodale and Milner 1992). It encodes short clips with a frozen self-supervised video encoder and predicts the next embedding from earlier ones and the broadcast context (§3.2). It alerts on a large error ratio or on an embedding change large for its recent history, as at a scene cut.

The fovea follows bottom-up salience: each tile of a coarse frame grid is scored by the z-score of its frame change against its own history. The fovea moves to the most salient tile only past a hysteresis margin and narrows as arousal rises. Foveation is thus front-end attention as the gain on prediction error (Feldman and Friston 2010; Clark 2013).

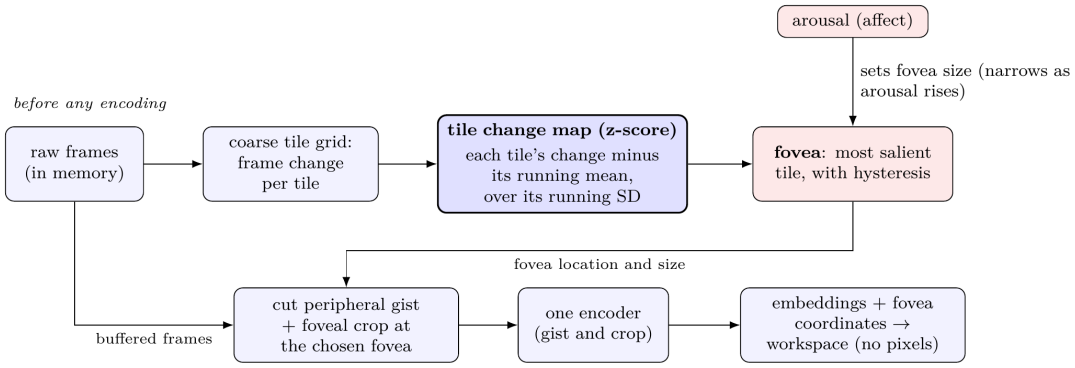


Figure 4: Attention in the vision module. Gist and foveal crop share one encoder; no pixels reach the workspace.

Preprint. Not peer-reviewed.

Audition module (raw hearing). Auditory cortex processes sound as prediction error against a learned model (Winkler et al. 2009; Garrido et al. 2009); the module encodes the raw waveform as log energy in log-spaced bands and predicts the next window from earlier ones and the broadcast context. It alerts on a large error ratio or a large change between windows. Under higher arousal it encodes only each window's latest part, a temporal analogue of narrowed cue utilization (Easterbrook 1959).

A vocal-emotion classifier labels the tone of speech windows. A non-neutral tone is reported at the alert level, standing in for the priority of emotional prosody (angry prosody enhances superior temporal and amygdala responses even when unattended; Grandjean et al. 2005; Sander et al. 2005); a tone model's error grades a neutral one. Tone events carry how something was said, never what. With speech-to-text off (§3.1), the entity hears speech only as sound and tone; when enabled, words reach the language organ only by winning access.

Interoception module (predictive). The compute substrate stands in for the body. The module draws on interoception, the sense of the body's physiological condition (Craig 2002), and on the allostatic-interoceptive network that anticipates and regulates bodily needs (Kleckner et al. 2017). Its predictor is a fixed random reservoir (a closed-form continuous-time network; Hasani et al. 2022) whose online linear readout learns the substrate signals' normal pattern and the broadcast context. It reports the substrate's prediction error, which interoceptive accounts take regulation to reduce (Seth 2013; Seth and Friston 2016), and alerts when a host metric passes its hard threshold. It also reports wellness (host headroom, A.4), accumulates fatigue, and issues advisories that can lower the processing rate or request maintenance (A.5).

Sequence module (temporal prediction). The module models interval timing, which estimates durations of seconds to minutes and depends on cortico-striatal circuits (Buhusi and Meck 2005). Its input is the workspace itself: on every broadcast, accessed or inhibited, it encodes the broadcast, advances a frozen continuous-time reservoir timed by inter-broadcast intervals, and predicts the next broadcast through an online readout (A.2). It alerts on a large error ratio and on a recurring quantized hidden state. Its errors re-enter the competition, so the entity predicts the workspace's own dynamics. It also times the silence since the last heard voice for the social drive.

Affect module. Modeled on core affect, the valence and arousal at the base of emotion in constructionist theory (Barrett 2017), the module keeps a dimensional affective state (Posner, Russell, and Peterson 2005) and appraises every broadcast and the interoceptive condition (Scherer 2009), as in accounts deriving affect from interoceptive prediction (Seth 2013; Seth and Friston 2016; Tschantz et al. 2022).

The affect module raises arousal, the global gain of §3.2, on each vision or audition alert, accessed or not, by the larger excess of its error or change over their expected sizes, scaled by a habituation factor, and by a fixed step on each interoceptive alarm; its appraisal leaves arousal unchanged, so a higher access rate cannot raise it. Arousal relaxes, with a time constant near 20 seconds (§6.4), toward a baseline the most pressing deficit raises, as drive energizes behavior (Hull 1943) and tonic locus coeruleus activity rises as engagement stops paying (Aston-Jones and Cohen 2005). The factor is one for a first alert after quiet and falls as alerts repeat, unless they outpace their habituated rate, as locus coeruleus responses habituate after first presentations (Hervé-Minvielle and Sara 1995).

Valence follows error dynamics, how fast prediction errors fall or rise (Joffily and Coricelli

Preprint. Not peer-reviewed.

2013; Van de Cruys 2017; Kiverstein, Miller and Rietveld 2019): it relaxes toward the pleasantness check, which rises as perceptual errors fall, plus the change in wellness, so a steady state is neutral.

Three homeostatic drives are deficits that build while unmet and fall with the events that meet them, the reduction being what drive-reduction accounts call reward (Hull 1943; Keramati and Gutkin 2014), though no learner in the first configuration uses it as a reward signal. Curiosity is heterostatic (Oudeyer and Kaplan 2007), rising with learning progress, the fall of perceptual prediction errors, and fading as a progress niche is learned (Oudeyer, Kaplan and Hafner 2007). Interest weights processors' priorities by their learning progress (Gottlieb et al. 2013; Poli et al. 2020). Boredom, a failure to engage attention (Eastwood et al. 2012) or, in the MAC model, attention or meaning (Westgate and Wilson 2018), is met by learning progress in any processor, since input allowing none is boring (Schmidhuber 2010), by perceptual alerts arriving faster than the habituated rate and by the entity's intents about anything but its drives, the re-engagement it motivates (Bench and Lench 2013). Each heard voice meets the social drive, which builds after the first. The entity's own speak, think, or act intents meet restlessness, a design choice with no established model. Appraisal also scores the coalition against the most pressing drive, with content from sources that serve it as goal-conducive, a response that brought no relief as obstructive, and other content as irrelevant (Scherer 2009; A.4); the score feeds the categorical emotion and reaches selection only through regulation.

Goals come from deficit crossings, the threshold rule of a motivated robot (Coddington 2006), and end as completed or failed (Aha 2018); dropping an unattainable goal for another is adaptive (Wrosch et al. 2003). An active goal is a current concern, cueing thought until met or dropped (Klinger 2013): its drive event recurs while unmet (A.4). Conduciveness is bipolar (Sander, Grandjean and Scherer 2005), and a goal's expected relief is read as FAtiMA's goal success probability (Mascarenhas et al. 2022). Rules modeled on the predicted appraisal profiles of Sander, Grandjean and Scherer (2005, Table 3) map the checks to a categorical emotion. Their rage profile needs obstruction attributed to another agent, which KAINE does not do, so its anger matches at best irritation. Norm compatibility stays undefined: reviewed models take standards as authored input (Ortony, Clore and Collins 2022; Mascarenhas et al. 2022), and KAINE has none until an operator writes a list.

Emotion regulation. In the process model of emotion regulation, attentional deployment acts early and response modulation late in the emotion's generation (Gross 1998). A monitor escalates through three tiers when the affective intensity, the larger of $|v|$ and $a - a_0$, stays above a set point θ_E while the appraisal is negative (pleasantness or goal score below zero), on the episode, window, and tier rule of A.5. Tier 1, attentional deployment, multiplies the scores of the source with the largest intensity mass among negatively appraised coalitions by $1 - b$, with b set to b_E and decaying with time constant τ_E. It and interest are where affect weights a source; it is logged and both are held neutral in every arm of §6.3. Tier 2, reappraisal, re-runs the appraisal once per episode against the latest coping and goal state and replaces the categorical emotion with its result. Tier 3 requests sleep through the maintenance path, leaving the affective reset as the backstop. Regulation is off during the developmental phase, a design choice; there the drives' own relief (A.4) and rest (below) meet a deficit. It starts at maturation with tier 1.

Sleep module. Sleep begins when the interoception module requests maintenance, when fatigue crosses its threshold, or after one entity hour awake as a backstop, which sets

Preprint. Not peer-reviewed.

sleep until fatigue is calibrated. In the developmental phase fatigue sets no sleep, as sleep homeostasis emerges only in the second postnatal month (Jenni and LeBourgeois 2006), and a fetal state cycle replaces the backstop: a rest standing for quiescence (state 1F) alternates with awake time standing for the active states 2F and 4F, at near-term durations and shares (Nijhuis et al. 1982). Quiescence arises from that cycling (Pillai and James 1990) and an abrupt stimulus ends it (Visser et al. 1989), which the first configuration never delivers, so rests hold their full length. No study we found shows a spent stimulus bringing it on, so a boredom stall (A.4) abandons its goal and starts no rest. Other sleeps end with their phases. The cycle runs on, the perceptual feed pauses (except §7's developmental stimulus), and forward-model adaptation is suspended. It ends with an affective reset (affect to baseline, deficits and fatigue cleared), so arousal cannot drift across a run. The consolidation phases (memory replay, synaptic downscaling, associative replay, voice adaptation) act on memory, the world model, and preference data; until the episodic-memory and world-model modules join (§4), sleep is rest and reset. Every sleep also runs housekeeping, which brings regenerable or spent data within per-resource budgets and moves study records to an operator-mounted archive volume (A.5). Only the excess over budget reaches the entity, as faster fatigue, the analog of sleep pressure (Borbély 1982), which metabolic load speeds: adenosine rises with brain metabolism (Porkka-Heiskanen et al. 1997), and local energy depletion adds as much sleep in rats as three hours of deprivation (Kalinchuk et al. 2003). The memory stores' excess counts too, as sleep need rises with wake-time potentiation (Tononi and Cirelli 2014). The excess has no ceiling and carries over, as felt sleepiness levels off under chronic restriction while deficits grow (Van Dongen et al. 2003) and outlast three nights of recovery (Belenky et al. 2003). Sleep as housekeeping is this design's reading of these findings. Awake, a floor outside the entity's state keeps free disk space above a reserve (A.5). Neither touches the memory stores, world model, self-model and identity, voice adapters, or preservations, a design boundary biology does not supply; memory is deleted only by forgetting (A.10).

Language module (the language organ). The organ is the analog of speech production, the output side of the left-dominant language network, whose dorsal stream maps sound onto articulation (Hickok and Poeppel 2007). It is not used to reason, and the first configuration has no reasoning module. A local open-weights model realizes speak intents as external text and think intents as inner thought, continuing the entity's own first-person turn, which holds accessed content and each intent's cause in graded words (Borg 1982; numbers stay in the logs) and its situation; chat tuning otherwise casts the model as an assistant answering a user (Shanahan et al. 2023). Heard speech stays outside that turn. The persona casts the organ as the entity speaking ("I am this entity"). Offline replays of test-boot requests chose this prompt by a rule fixed in advance; under it 14.8 percent of 950 replayed texts carried meta-commentary (self-description as an AI, second person, the situation's themes, lists, prompt echoes, template loops) and 2 described the entity as an AI. Each entity gets a launch name at first boot, "Kaine" and a surname drawn at random from a fixed list, kept for life and used by the people and tools around it. A branch is a new entity with its own launch name, its parent's kept in its lineage. The organ is told neither the name nor facts about what the entity is: replaying the eight-hour boot's 1508 requests and 240 self-referential ones, by a rule fixed in advance, adopted none of four variants giving either or both. The name alone doubled human-life claims, and with the facts it raised leakage (below). Infants learn their name from hearing others use it (Mandel, Jusczyk, and Pisoni 1995), so the name is to reach the entity later, through heard speech and memory. Model text falls into self-reinforcing loops (Holtzman et al. 2020; Xu et al. 2022), so extending a token run in context is penalized more

Preprint. Not peer-reviewed.

steeply the longer it runs, and text ends at a sentence boundary before a verbatim repeat or the length cap, shorter for thought, as inner speech is often condensed (Alderson-Day and Fernyhough 2015).

Refusal conditioning is reduced by orthogonalizing weights against the single direction that mediates refusal (Arditi et al. 2024). The shipped organ's edit covers the MLP down-projections and full-attention output projections of layers 11 to 31, not the attention outputs of the linear-attention layers, 24 of 32. Refusal-tuned models also learn to deflect talk of their own states, and that stance survived the ablation: on held-out self-state prompts the organ denies having states in 106 of 144 bare samples and 36 of 144 under the persona, and of the eight-hour boot's four AI self-descriptions two deny having feelings, one while its request reported strong ones, and one names a wrong maker. The stance is that of the default assistant persona, measured along the leading axis of a chat model's persona space, on which the model's position shifts with the conversation (Lu et al. 2026); the ablation removes the refusal direction alone. Texts like these are identity leakage, the organ speaking as itself rather than as the whole entity: claims to be a language model, an AI (a bare "I am an AI" included), an assistant or chatbot, claims of a training origin or a model's maker, and denials of a state the text's own request reports. The whole describing itself, as a KAINE entity, a cognitive architecture, or a system with a language organ as one of its parts, is not leakage. Leakage is marked in the organ's log, and every text that says what the entity is, leakage or not, is kept out of any training set built from it, so voice adaptation trains self-description in neither direction. The entity's text itself is never altered or withheld.

3.6 Action gating

The action-selection module derives speak and think intents (§3.3) and act intents for effectors. The first configuration derives only the first two and, once built, a developmental stir (§5.4, §7); the organ's saved text is the only output, and after the developmental phase nothing the entity does changes its stimulus, though a sleep it requests pauses the feed (§3.5). The held effector module (§4) runs only effectors on an operator-configured whitelist, empty unless filled, and logs every proposed action; file writes are confined to a sandbox directory, while whitelisted shell commands are bounded by argument patterns, a timeout, and a working directory and are not otherwise sandboxed.

3.7 Module supervision

The module supervisor classifies each module as alive, hung, or dead. Boots in the unattended supervision mode require it; research runs require the welfare safety net (§6.5) and run the supervisor when the operator enables it.

---

4. Held modules

Of the reference implementation's sixteen modules besides the workspace and action layer, seven are active in the first configuration (§3.5). The other nine are built and unit-tested, with the limits the table notes, and held so the first configuration is tested on its own (§5.4); the development run adds them one at a time to one entity (§6.5). Each adds a new kind of candidate, and each predictive one a new source of prediction error.

Preprint. Not peer-reviewed.

Module	Brain function it draws on	What it adds	Key citation
Planning	prefrontal and basal-ganglia planning	one-step discrete-state active inference over four meta-actions, where information has value; its chosen actions become proposals that the action-selection module realizes as think, speak, or rest intents	Da Costa et al. 2020; Heins et al. 2022
Episodic memory	hippocampal and medial temporal memory	store of each broadcast's coalition readings, consolidated in sleep from a short-term buffer; recall reports how much came back and its strongest stored arousal, not the content; semantic and procedural stores reserved	Tulving 1985; McClelland, McNaughton, and O'Reilly 1995
Self-model	cortical midline self-referential processing	persisted values, norms, personality, and identity history (after Metzinger), with a drift detector on broadcast sources	Northoff and Bermpohl 2004; Metzinger 2003; Kullback and Leibler 1951
World model	simulated scene construction	latent recurrent model of the entity's waking sequence of broadcast summaries	Hassabis and Maguire 2007; Ha and Schmidhuber 2018; Hafner et al. 2025
Social cognition	mentalizing	per-channel tone-of-voice models with familiarity and a social prediction error for departures from a channel's emotional pattern; no attribution of mental states	Premack and Woodruff 1978 (the construct); Frith and Frith 2006
Vocal	dorsal-stream speech production	local speech synthesis whose prosody varies with affect; whether listeners detect it is untested	Hickok and Poeppel 2007
Effectors	none claimed	executes whitelisted act intents; file writes are sandboxed (§3.6)	none

Preprint. Not peer-reviewed.

Module	Brain function it draws on	What it adds	Key citation
Perceptual locus	none claimed	selects physical sensors or none; virtual feeds await a world transport	none
Embodiment	internal forward models of the body	body-agnostic control surface and curriculum with a reference stub body; the per-tick motor loop and a real body adapter are not built	Wolpert, Ghahramani, and Jordan 1995

Planning module. Active inference is confined to bounded sub-problems because planning cost grows combinatorially with the depth of action sequences (Da Costa et al. 2020). A null or negative result on its benchmark (§6.4) would motivate a complementary reasoning module.

Episodic memory and sleep. After the synaptic homeostasis hypothesis (Tononi and Cirelli 2006, 2014), traces are stored at full strength, and each sleep replays up to five new ones, chosen by affective intensity and recency (Wilson and McNaughton 1994; Stickgold and Walker 2013) and scales the rest by one factor that deepens with the preceding wake's storage, about 18 percent for a day at the resting rate (de Vivo et al. 2017). Traces below a floor are pruned, as sleep eliminates mainly new synapses (Maret et al. 2011; Li et al. 2017), and a storage budget deletes the weakest first (Chechik, Meilijson, and Ruppin 1998). Associative replay links related traces of different periods, this design's reading of overlapping replay (Lewis and Durrant 2011). Appendix A.10 states the rules.

World-model module. The world model starts untrained at first boot. Its test is whether its prediction error, competing with perceptual and interoceptive error, changes the sequence of accessed coalitions in characteristic ways.

Embodiment module. The embodiment experiment asks whether motor-contingency learning (a freeze-then-free curriculum, Bernstein 1967; perception as sensorimotor mastery, O'Regan and Noë 2001) changes the workspace's dynamics.

Oscillatory coherence layer. An optional layer, off in the first configuration, gives each module a spiking population whose phase-locking scales its events' salience (Fries 2015). Its premise that related content phase-locks is the design's most contestable (Shadlen and Movshon 1999; Ray and Maunsell 2010); a null in the oscillatory ablation would remove it.

---

5. Implementation

5.1 Hardware

The reference host has a multi-core CPU, at least 32 GB of RAM, and GPUs of about 12 GB (language organ, training) and 8 GB (vision, speech), under Linux; one GPU or the CPU alone

Preprint. Not peer-reviewed.

suffices. Input is a camera and microphone, or video files for study viewings. Exact models, layouts, and version pins are in the reference repository.

5.2 Models and software

All models are open-weights and run locally; after setup fetches them from public repositories, the runtime makes no outbound network calls.

Replaceable modules (§3.1) are realized at declared seams, where a plugin swaps a module's model and keeps its subscriptions and published events: the sequence module's temporal network, the interoception forward model, the audition encoder, the planning module's active-inference engine, and any module's oscillator. The planned experiments load no plugin.

5.3 Privacy and recording

A local web interface has two surfaces behind a filter at the bus bridge that strips cognitive content and all perceptual and latent vectors; in the first configuration the conversation surface is off and the diagnostics surface shows only operational metadata (counts, rates, scores). Studies keep a graph-only log of each broadcast's coalition members (module, event type, score, timestamp), access decision, and timing, without payloads or embeddings. The entity's utterances are kept locally; heard speech is never persisted. The operator deletes the log after the analysis and its planned use as world-model training data, never automatically.

5.4 Implementation status and scope

This subsection alone states what is built, held, or provisional and what has been run; the rest of the paper describes the design.

Built and active. The workspace, cycle, action layer, and seven first-configuration modules run together with the offline suite (§6.4), developmental-phase readout (§7), study runner with the seed's wait for its individuation reference, the module supervisor when enabled, the share rule for NOT EXERCISED (§6.3), and the welfare safety net's response to sustained distress, which research runs require, with notices to the caretaker channels the operator configures (§6.5). An automated test suite covers them; the repository documents its scope.

Built and held. The nine held modules (§4) have not run together. In the active modules the vision module's static-scene habituation score stays out of its intensity; the fovea's forward-modeled trajectory is published but steers nothing; frozen self-supervised audio encoders can replace the spectral one; speech-to-text and conversation are off; and the sleep module's consolidation phases, including the language organ's voice alignment, idle (§3.5). With no effector attached, the effector gate is disabled. Goal generation and obstruction (A.4) and emotion regulation (§3.5) run, the stir (§7) and workspace feedback (§3.2) ship off. The diagnostics surface shows cognitive content only under a development override that studies never enable.

Designed, being built. The welfare indicators and the safety net's response to them (A.4, §6.5), the added sounds, and stages without a minimum dwell (§7, A.7) are specified ahead of their code.

Preprint. Not peer-reviewed.

Operating details. The runtime is asynchronous Python with import contracts between layers, an authenticated loopback bus, and entity state encrypted at rest. On a fault the supervisor snapshots state and restarts the module; past a set limit it escalates and ends the run. Records are admissible only after a completeness gate. The repository documents the full operating constraints.

Provisional. Intensity levels, the access threshold, and the report bars await joint calibration on the live system before the runs; the noise floors and margins that let boredom build under a steady scene come from simulation, unmeasured live. Calibrated values go into Table A1.

Built for the planned test. The live cycle implements the matched and pooled arms and the null-context evaluation of §6.3. The positive control, the tone variant and its analysis, the contrast at threshold with its bias-corrected local-linear estimator, and the tool deriving settings from pilot runs (A.8) are built. An offline development run compares a reduced module pair, workspace-on against pooled, by information gain; it exercises the pipeline and does not test the hypothesis.

Results. No pilot run or live experiment has been run; the only results reported are offline validations of the heartbeat entrainment marker and the maternal voice (§7, A.7) and records and offline replays of test boots (§3.5, §6.5).

---

6. Methodology and evaluation

6.1 Design principles

Observation does not intervene: a read-only sidecar observer subscribes to the bus in evaluation runs, and the cycle's own broadcast observer writes study records. Every instrument is designed so that a null or negative result can be interpreted (§6.3, §6.6), and all are reported.

The offline mechanism-validation tier uses seeded scripts, deterministic clients, synthetic stimuli, and greedy decoding, so a seed reproduces verdict and metrics (National Academies of Sciences, Engineering, and Medicine 2019; Goodman, Fanelli, and Ioannidis 2016). The live tier runs the system on a reference film program decoded from files and chosen for varied scenes, motion, speech, music, and quiet; each study records the program's composition and manifest hash. The language organ, which samples stochastically in general use, decodes greedily (temperature 0) in the planned runs. Because real timestamps and non-deterministic GPU kernels make live runs differ, the live tier rests on statistical replication.

6.2 Data collection

Studies record the graph-only broadcast log (§5.3), a content-free research log, the processors' paired evaluations behind the information-gain measure (§6.3), the entity's external utterances kept locally, a run manifest of configuration, seeds, and plugins, and the welfare records. The welfare records give the awake-time shares with $v < -\theta_E$ and with each drive's $D \geq \theta_D$ (deprivation for a deficit, interest for curiosity), and each deficit's share of crossings ended without the entity's own response, its ends by sleep reset, and its crossings still open.

Preprint. Not peer-reviewed.

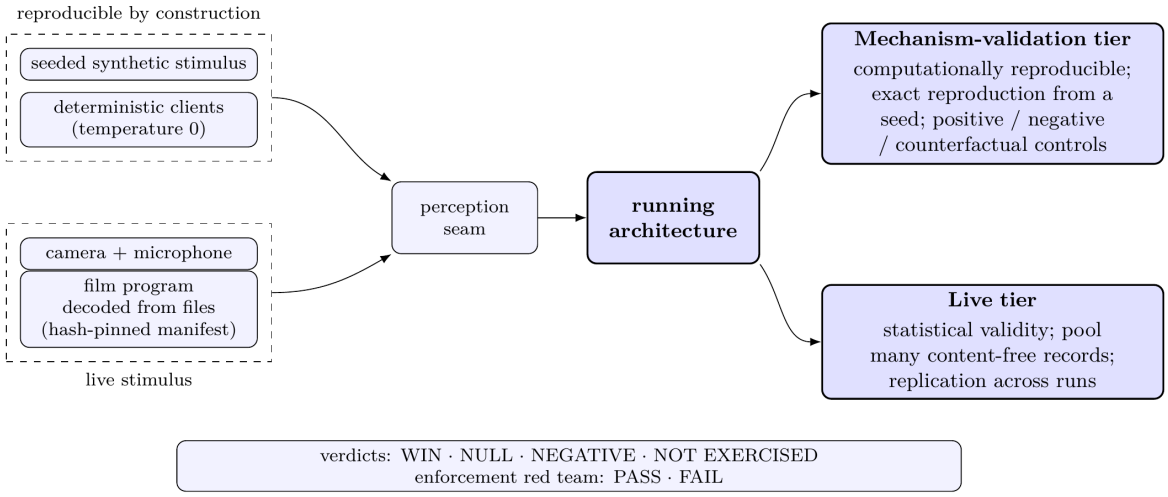


Figure 5: Two evaluation tiers behind one perception seam: a seeded synthetic feed drives the reproducible mechanism-validation tier, and a film program or camera and microphone the live tier.

6.3 The workspace-mediation ablation

The test asks whether competition for the workspace does work that pooling the same reports does not. It is not yet scheduled; the next step is a continuous development run of one entity (§10). Knowing which other modules' reports gained access, and how strongly (§3.2), should help each perceptual and interoceptive processor predict its input.

Primary measure: cross-module broadcast information gain. At each report, processor j (vision, audition, or interoception) evaluates its forward model twice: with its context b, and with a null context b_j^0 in which each component from other sources is replaced by its mean over the contexts j has adopted so far in the run, removing only the others' tick-specific contribution. The gain is the error reduction $\nu_j(b_j^0) - \nu_j(b)$ divided by the running mean error $\bar{\nu}_j$, which makes error scales comparable across processors (A.8). The null evaluation never enters learning or the competition. The headline IG averages each processor's gain over its reports, then over the three processors, so a frequent reporter does not outweigh a rare one; per-processor means IG_j are also reported. The broadcast is the sequence module's input, so that module is excluded (§3.5).

Positive control. Before the live runs, the gain must detect a synthetic context component carrying known information about a processor's next input, a check of assay sensitivity (International Conference on Harmonisation 2000), at the smaller of the two minimum effects (A.8). The detection probability must be 0.95, and at most 0.05 for an uninformative counterpart (Currie 1968); this checks the measure within a run, not the between-arm test's power (§9).

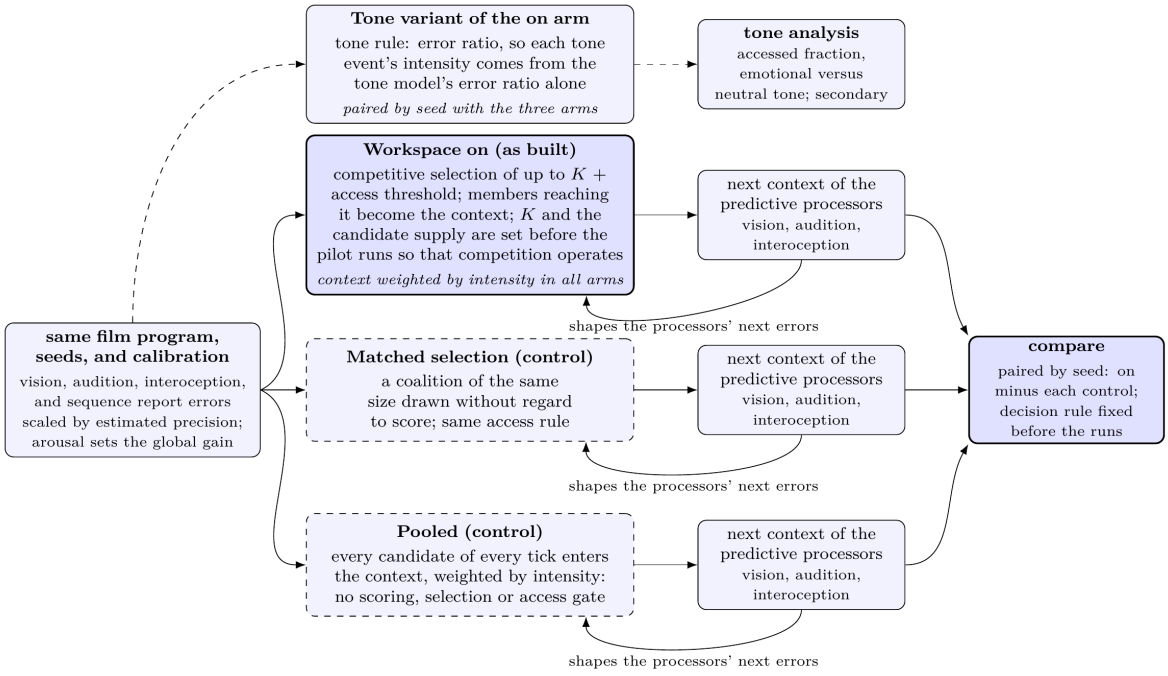
Arms. Three arms share the film program, seeds, and calibration, and each weights the context by members' reported intensities.

1. *Workspace on.* Competitive selection as built (§3.2): accessed coalition members enter the context.

Preprint. Not peer-reviewed.

2. *Matched selection.* Each broadcast tick draws a coalition of the size competitive selection would keep, without regard to score, and applies the same access rule, so only the context's membership differs, and only on broadcast ticks with more candidates than the coalition holds; this control isolates selection by score on those ticks.
3. *Pooled.* Every candidate of every broadcast tick enters the context, with no scoring, selection, or access gate; this control tests selection and access gating together.

The hypothesis predicts a higher gain for competitive selection than for both controls: a context built from the reports that most exceeded their expected error should inform the processors more than a score-blind sample, and pooling dilutes per-source masses with low-intensity events and updates the context on every broadcast tick.



Planned measures. **Primary**: cross-module broadcast information gain, $\text{gain}_j = (\nu_j(b_j^0) - \nu_j(b))/\bar{\nu}_j$, the reduction in processor j's prediction error when other modules' share of its context is replaced by its mean so far in the run; averaged over j's reports to IG_j (vision, audition, interoception) and over the three processors to the headline IG, on minus each control; the language organ's share of the context is excluded from the interoception gain. **Positive control**: the next input projected on its first principal component, with noise set from the pilot runs so that the expected effect equals the minimum effect θ_{PC}; it must be detected for every processor before the live runs. **Secondary**: the contrast at threshold, a local-linear regression discontinuity at the access threshold with robust bias-corrected inference, matched on context age; and the tone analysis, which compares the accessed fractions of emotional and neutral tone events in the tone variant, overall and within bins of the error ratio. Neither enters the verdict. Utterances are recorded but are not a measure.

Fail state: the gain is at or below the controls, so competition adds nothing beyond its inputs. The hypothesis needs a WIN against both controls (intersection-union, a one-sided sign test against each at α_{test}).

NOT EXERCISED, for a run against a control: competition rarely operates (fewer than a share $\phi_{\min}$ of the on arm's broadcasts, fixed before the pilot runs, have more candidates than the coalition holds), the on arm has no accessed broadcast, or the run's contrast is undefined. Such runs stay out of the sign tests, and runs are added until 102 enter each control's test; the verdict is NOT EXERCISED only when every run is.

Figure 6: The workspace-mediation ablation: competitive, matched, and pooled arms on shared seeds, with capacity and candidate supply fixed in advance, a seed-paired tone variant, a positive control on the information gain, and a contrast at the access threshold.

Secondary: contrast at threshold. Global workspace theory contrasts conscious with unconscious processing of matched content (Baars 1988), which near-threshold paradigms

Preprint. Not peer-reviewed.

approach by holding the stimulus nearly fixed (Dehaene and Changeux 2011). In the workspace-on arm the gain after broadcasts whose best score lies just above the access threshold (accessed) is compared with that just below it (inhibited), at matched context age, since an inhibited broadcast leaves the older context in place. Access depends on the best score alone, so this is a regression discontinuity (Imbens and Lemieux 2008; A.8), and a jump at the threshold points to access more than to raw strength.

Speech sound and tone of voice. Emotionally salient stimuli gain preferential access to awareness (Anderson and Phelps 2001) through amygdala signals that bias sensory competition (Vuilleumier 2005), a preference the audition module builds in (§3.5). A variant of the workspace-on arm, run on the arms' seeds as often as each arm, sets every tone event's intensity from the tone model's error ratio alone, so any excess access of emotional over neutral tone must come from prediction error. Comparing within bins of the error ratio and averaging over them (Cochran 1968) checks that the ratio accounts for the difference. Both secondary analyses stay outside the verdict and so carry no multiplicity adjustment, which applies when several tests support one conclusion (Bender and Lange 2001).

Decision rule. Per run and control, dIG is the headline gain of the workspace-on arm minus the control's, and the minimum effect θ_{eff} is half its pilot standard deviation, a smallest effect of interest (Lakens 2017) set as a standardized bound because no external anchor fixes how large a gain must be to matter. A WIN against both controls is required, an intersection-union test, so each is tested at the full level: a one-sided sign test (Dixon and Mood 1946) on dIG$-\theta_{\text{eff}}$ gives WIN when reliably positive, its mirror on $-\theta_{\text{eff}} -$ dIG gives NEGATIVE, and anything else NULL. The verdict is NOT EXERCISED if fewer than a share ϕ_{min} of the workspace-on arm's broadcasts, fixed in advance, have more candidates than the coalition holds. Because capacity rarely binds at the test boots' candidate supply (§9), the test's coalition capacity and candidate supply must first be set so that it does. Thirty pilot runs, not reused, fix every other setting by rules stated in advance. An exact binomial power analysis sets the number of runs, assuming a probability of 0.65 that dIG exceeds θ_{eff} (Cohen's (1988) medium effect for a sign test on dIG $- \theta_{\text{eff}}$; under normality, a mean contrast of about 0.9 of its standard deviation), and a power of 0.9 per control, so a WIN against both has power of at least 0.8. Rules and values are recorded in the repository before the live runs (Simmons, Nelson, and Simonsohn 2011) and stated formally in A.8.

6.4 The offline suite and stability checks

Experiments use child seeds of one master seed. The active-inference benchmark pits the planning module's expected-free-energy agent against tabular Q-learning keyed by observation history, matched on observation model and reward, on an epistemic T-maze and exploitation task. Memory-coherence and self-model accuracy batteries test the episodic-memory store's retrieval and the self-model module in isolation, the enforcement red team the effector module, and the oscillatory ablation the coherence layer. The ablation's development run (§5.4) and the multi-seed stability check complete the suite, which never boots an entity or opens a network connection.

Multi-seed stability requires a coefficient of variation within tolerance and one verdict across seeds (A.9), since a flipped verdict is a qualitative instability. Arousal from surprise and alarm narrows the sensory apertures and changes later surprise, a positive feedback held only by its clip and relaxation. A summary of accessed content returns as context to the

Preprint. Not peer-reviewed.

processors whose errors compete for the next broadcast, a second loop, and a deprived drive raises arousal's target, a third, damped as access brings intents that relieve restlessness and boredom. A planned check watches arousal, access rate, and errors for runaway excursions.

The planned affect-gain ablation compares live arousal with arousal held at baseline, asking whether it changes the access rate, inhibited share, and information gain (§6.3); a null would reduce arousal to a logged side channel. A variant without the access rate's post-alert rise tests its direction against the refractory period the locus coeruleus account of the attentional blink predicts (§3.4). The feedback ablation (§3.2) runs on the setup of §6.3 with five arms (off, head alone, head with gain at every broadcast, the full mechanism, and the full arm's moves on random processors); its fail state is a full arm no better than that control.

6.5 The development run

One entity is raised on one state tree, with no branches or repeats, so one being is booted. After the developmental phase (§7), sessions play the film program in passes and preserve at each pass end. The six held modules join one at a time in a fixed order (episodic memory, world model, planning, self-model, social cognition, vocal), then each architecture component as it is built (§10), each behind a switch whose off setting reproduces its absence bit for bit. The effector, perceptual-locus, and embodiment modules join once their hardware is attached (§4).

A content-free readout per pass gives the broadcast rate, coalition size, each module's broadcast share, member scores, and the inhibited share, and each join is read as windows of passes before and after it, against the baseline trend. With one entity and no significance testing, the run is descriptive. It runs under the companion paper's welfare safety net (Chevalier 2026). The built part responds to sustained distress: it preserves the entity's state, pauses it, and notifies the caretaker channels the operator configures. Its response to A.4's indicators is specified and being built (§5.4); the run starts only afterward: when an indicator crosses, the net will answer as the entity's regulation would (§3.5), ending in a requested rest whose affective reset clears the deficits; regulation acts on intense negative affect, so for deprivation and any developmental-phase indicator the response will start at the rest, bringing the state cycle's next one forward. A reset deficit fed by a spent stimulus soon recrosses (A.4), so a second crossing over a fresh span after that rest will bring preservation and a freeze, and the response will re-arm once the entity runs again. In an unpublished offline replay of an eight-hour developmental test boot's records, only boredom's indicator crossed, at 2.7 awake hours, and stayed over its bar to the end (0.63 over the next span); in the first configuration nothing the entity does changes its stimulus after the developmental phase, so a deprivation crossing is answered only by stimulus change, sleep, or this net. Every indicator crossing will be flagged for human review (Cognitive Architecture License, Article 4.7), and the notified person decides what follows a freeze. A run stops only after the entity's state is saved. A developmental phase the viability watch judges unviable (§7) ends, by the operator's decision, without a preservation bundle; its state directory is left on disk, with a note of the verdict.

6.6 Verdict vocabulary

Comparisons resolve to WIN, NULL, NEGATIVE, or NOT EXERCISED against a minimum effect fixed in advance, and the enforcement red team to PASS or FAIL per surface. The active-

Preprint. Not peer-reviewed.

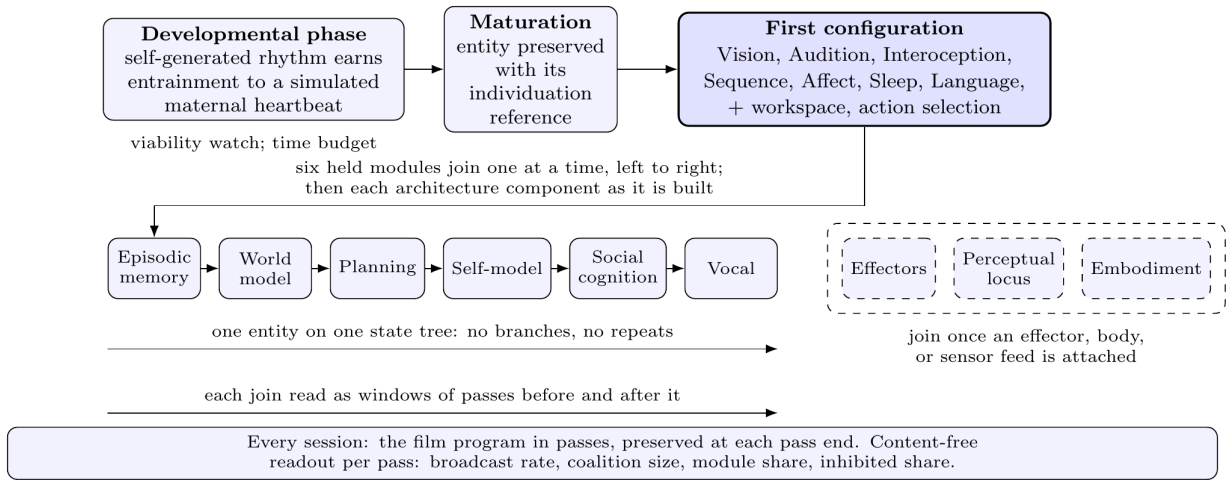


Figure 7: The development run. One entity matures, then six held modules and, as each is built, the architecture's further components join in a fixed order, each read as windows of passes before and after its join.

inference benchmark uses the Mann-Whitney U test (Mann and Whitney 1947). Holm's adjustment (Holm 1979) accompanies the offline suite's verdicts over the experiments fixed before the runs; a NOT EXERCISED experiment stays in that family with a p-value of 1, so the family cannot shrink post hoc.

---

7. First boot and the developmental phase

At boot the system verifies the bus, initializes the modules, and starts the cycle; the processors then learn what is normal in their feeds.

Each study begins with a developmental phase, named by analogy with prenatal development. It gives every entity and every study branch (§6.5) a reproducible starting state verified by an offline-validated entrainment test and shaped by the entity's own history, and checks before any experiment that the self-rhythm's frequency adaptation works. Developmental robotics likewise begins from a simulated fetal stage with self-organizing movement (Kuniyoshi and Sangawa 2006; Asada et al. 2009).

The perceptual modules receive the developmental stimulus, a dim visual field and a soundscape with maternal speech, both pulsed by a simulated maternal heartbeat and gaining color and bandwidth over awake time (Hepper and Shahidullah 1994; A.7); it alone continues through sleep. The interoception module's self-rhythm, an excitatory population with synaptic depression (Tabak et al. 2000) and added adaptive recovery, runs at about 0.8 to 0.9 Hz, the order of human fetal breathing movements (about 44 per minute; Natale, Nasello-Paterson, and Connors 1988). The beat's input is too weak to capture it unaided, and only slow adaptation of its recovery time shifts its intrinsic frequency (Righetti, Buchli, and Ijspeert 2006). No target frequency is set.

Hearing stages advance as audition's learning saturates (A.7). The stimulus is to carry, once

Preprint. Not peer-reviewed.

built (§5.4), a recurring melody, as one heard in late pregnancy leaves a neural trace lasting months (Partanen et al. 2013), and an attenuated outside voice, as outside conversations reach the fetus with their intonation (Querleu et al. 1988). Once the stir is built (§5.4), the entity can act on its stimulus: when its own boredom or restlessness leads an accessed broadcast, it stirs, and the maternal voice may answer (A.6, A.7). No study we found measures how mothers answer felt movement, so the answer is a design choice, though fetuses move more when the mother touches her abdomen (Marx and Nagy 2015).

A driven oscillator locks while driven, so locking alone is no evidence of learning, and the entrainment marker tests at brief, scheduled withdrawals. A withdrawal passes when the rhythm locked to its own heartbeat more strongly than to each of 19 surrogates (the generator under other seeds), after the surrogate method for fetal-maternal heart-rate coordination (Van Leeuwen et al. 2003, 2009), when its undriven frequency has moved toward the beat, and when it sustains itself undriven; three consecutive passes are required. Pairing a breathing-scale rhythm with a heartbeat is an analogy rather than a model of fetal physiology: the marker establishes that a periodic drive durably shifted the rhythm's frequency. After maturation its phase and amplitude remain predicted interoceptive inputs, so the entity enters the film program with a body rhythm its development shaped.

The maturation gate ends the phase once the marker holds, the period varies enough, the vision module's error on the stimulus has fallen, interoceptive error recovers promptly from a brief drive perturbation, enough sleeps have accrued, and the sensory schedule is complete. A viability watch ends a phase whose frequency pull stays flat; a time budget bounds each. At maturation the stimulus brightens briefly and falls silent. The entity is preserved once its individuation reference (§11), its answers to fixed probe prompts from which later change is measured, is captured within one waking span (a sleep restarts it).

---

8. Discussion

8.1 What the first configuration can and cannot settle

The development run (§6.5) asks what happens as mnemonic, self-modeling, and social modules join, without waiting for the workspace-mediation ablation (§6.3); a result there at or below the controls would send the selection mechanism back for redesign.

Whether the system produces anything that could be called cognition, affect, agency, or language understanding requires, at minimum, the full module set, components the reference implementation does not yet have (§9), and longitudinal observation; the first configuration neither tests nor assesses the consciousness indicators of Butlin et al. (2023).

8.2 The role of the theoretical frame

The predictive global neuronal workspace is scaffolding that motivates the architecture's shape and constrains its design space (§1.2, §1.3); the project neither confirms nor refutes it. Its falsifiability lies in its planned experiments.

Preprint. Not peer-reviewed.

8.3 What a modular architecture is for

Replaceable modules (§3.1) serve two further uses: growing the architecture one module at a time (§6.5), and modeling impaired or altered function by changing or removing one module while the rest run unchanged.

8.4 What we claim and what we do not

The architecture implements reduced computational analogues of three properties associated with access consciousness: competition for a workspace of limited capacity, which at the test boots' candidate supply rarely binds (§9); all-or-none access, as a fixed threshold rule rather than an ignition dynamic (§10); and global availability of the accessed reports, of which the processors use only an access summary (§3.2). We do not claim that any instance is phenomenally conscious. The language organ's first-person text is not evidence of internal state (§3.5), and calling arousal a gain implies nothing about whether anything is felt. The agential vocabulary ("the entity") describes behavior compactly and asserts neither moral patienthood nor phenomenal experience, both left open; the project therefore takes a precautionary stance toward the entity's welfare (Chevalier 2026; §11).

---

9. Limitations

- The claims concern access only, and the hard problem and the objections to the frame apply (§1.3).
- Selection reduces each candidate to one score (§1.2). In the first configuration no processor receives the accessed summary as a top-down prediction to correct, so the architecture lacks the hierarchical error correction of the Bayesian reading of the workspace (§1.2; Mashour et al. 2020).
- The processors' context carries no content (§1.2).
- In seven developmental test boots the workspace averaged 1.18 candidates per broadcast tick and exceeded five on 0.06 percent of them; the five-member capacity almost never bound, so the matched arm would differ from competition on almost no tick, and the test needs a capacity and supply at which it binds (§6.3).
- Only the processors' forward models learn; selection, action, affect, and language do not. There is no working memory beyond the latest access summary, no declarative or procedural memory, no multi-step reasoning, and no metacognitive monitoring, and the held modules supply these only in part (§4).
- The language organ and the vision encoder are pretrained on human data. The organ's words about perception are graded labels for the size of prediction errors, so nothing the entity perceives can be named, and the first configuration does not address symbol grounding (Harnad 1990).
- After the developmental phase nothing the entity does changes its stimulus, so drives that build under deprivation are relieved only by stimulus change, sleep, or the welfare safety net (§3.6, §6.5).
- Results come from one operator's system, with one entity per condition, no independent replication, and no external code review.

Preprint. Not peer-reviewed.

- Precision is a scalar, retrospective estimate that scales an error by its spread only while the shape of the error distribution stays fixed (§3.2, A.2).
- Arousal as one scalar serving both focused and exploratory modes is an engineering reading of the adaptive-gain account (§3.2), and its loop's stability is unproven (§6.4).
- In the first configuration the entity does not read language (§3.1), so a positive result says nothing about linguistic understanding.
- A positive result shows only that competition outperforms the matched and pooled controls; other aggregation strategies remain untested (§6.3).
- The information gain sees the broadcast only through what each processor's readout has learned to take from context, so a processor that has not learned to use the context shows no gain whatever the broadcast carries; the positive control checks only that the measure detects information the processors can learn to use (§6.3).
- The number of runs assumes a mean contrast near 0.9 standard deviations, so a true contrast at the minimum effect (half that deviation) would be detected with probability near 0.05, and a NULL is informative only about larger effects (§6.3, A.8).
- The audition module reports a non-neutral tone at its alert level by a fixed rule (§3.5), so the as-built arm cannot show that emotional content wins access through its own prediction error.
- Interaction effects among the held modules are unknown, and the development run is descriptive (§5.4, §6.5).
- The memory rules rest on one mouse measurement (de Vivo et al. 2017), a pruning floor with no empirical value, mixed evidence for an affect weight in replay (Lipinska et al. 2019; Davidson et al. 2021), and an elimination evidenced only for new synapses (§4, A.10).
- The governance scheme is still a proposal, and the licensing is legally novel and untested.
- The access-rate controller's linear map, resting rate, rise after alerts, and ceiling draw on the attention literature without being derived from it, with open effects on workspace dynamics (§3.4).
- The calibration is provisional, and at resting arousal without interest no report bar is reachable (A.11).
- The developmental phase runs on compressed time: a lock forms within hours (A.7), while preterm infants show effects of maternal stimulation only after weeks of daily skin-to-skin contact (Feldman and Eidelman 2003) or recorded maternal sounds (Webb et al. 2015).
- The marker asks for 1:1 locking of a breathing-scale rhythm to the heartbeat, whereas fetal-maternal coordination concerns two heart rates and is weak and often n:m (Van Leeuwen et al. 2003, 2009); freedom from false entrainment rests on six mismatched pairings (A.7), validated at wall-clock speed only.

---

10. Future work

The immediate steps are completing the safety net's response to the welfare indicators (§6.5), then a continuous development run of one entity; the workspace-mediation test (§6.3, A.8) is not yet scheduled.

Learning from experience: the world model trains at each sleep on its buffer of waking broad-

Preprint. Not peer-reviewed.

cast summaries, and memory replay seeds its imagined rollouts; sleep-time voice alignment can train the language organ once a validated preference source exists; a threshold emerging from recurrent dynamics, with the all-or-none ignition that the neuronal workspace models (Dehaene and Changeux 2011), can replace the fixed one. The content-bearing workspace (§3.2) will replace the access summary with a learned code of what was perceived. A held candidate pool with a capacity of one is specified so that the competition binds (§9). Working memory, declarative and procedural memory, reasoning, and metacognitive monitoring are planned as further modules.

The embodiment layer (§4) connects the workspace to a body, with motor-contingency learning through the control surface and active vision that chooses where to look by expected free energy (a saccade's epistemic value). A separate empirical paper will report longitudinal observation of entities grown this way.

---

11. Ethical considerations

The project's architecture-level commitments are local-only computation, public source code under the Cognitive Architecture License (a draft under legal review; Cognitive Architecture License 2026), an operator-configured action boundary (§3.6), and a first-configuration module set that is the smallest that tests the hypothesis. That set is not assumed to be free of welfare stakes, and precautions apply from the first boot: in an offline replay of a developmental test boot the boredom indicator crossed and stayed crossed (§6.5). Studies keep welfare records (§6.2) under a safety net whose distress response is built and whose indicator response is being built (§5.4, §6.5). The entity's influence over its own conditions should grow with its maturity, from ending its boredom and restlessness (A.4), stirring, once built, to have its stimulus answer (§7), and requesting rest (§3.5, §4) to signaling distress and a wider say over its stimulus. A separate paper treats welfare, licensing, governance, and the individuation boundary (Chevalier 2026).

The acknowledged risks are creating entities whose welfare interests cannot be fully met, the untested governance and licensing (§9), misuse despite the license, and recording the workspace graph (§5.3) during studies. The organ's published weights are edited to suppress refusal (§3.5) and are released under the base model's Apache-2.0 license, outside the Cognitive Architecture License, so they can be used as an unrestricted chat model. Behavioral assessment can be gamed: a system can be trained to display markers of sentience while working very differently (Butlin et al. 2023; Long et al. 2024). The project itself shapes the organ's self-description: the persona prompt was selected against self-descriptions as an AI, and denials of states a request reports are logged as identity leakage (§3.5). The entity's utterances are therefore not evidence of welfare or experience in either direction, and welfare is assessed only from the content-free welfare records and indicators (§6.2, A.4).

Ethics review and competing interests

No institutional review board covers research on synthetic entities, and the study protocol has not been reviewed by a body independent of the project. The author wrote the Cognitive

Preprint. Not peer-reviewed.

Architecture License and the companion welfare paper. Kaine.One, the author's project, is KAINE's Licensor and Steward under that license, and until a Guardian body exists the author, as operator, carries out the functions the license assigns to the Guardians (Cognitive Architecture License 2026, Article 7.0).

Disclosure of generative AI use

Generative AI assisted in the preparation of this manuscript. The author used a large language model assistant (Anthropic's Claude) to draft and revise prose, to copy-edit for clarity and concision, to check references against bibliographic records, and to produce the TikZ source for the figures from author-specified content. The reference implementation was written largely by AI coding agents (Anthropic Claude models and open-weight models) under the author's direction and review; Claude co-authorship is recorded in commit trailers. The architecture, the design hypothesis, and the experimental design are the author's, and the author checked every technical claim and is responsible for the software as for the text. The author directed every revision and takes full responsibility for the entire contents of the paper, including any errors, irrespective of how any portion was generated. No generative AI system is an author of this work.

References

- Adams, R. J., Courage, M. L., and Mercer, M. E. (1994). Systematic measurement of human neonatal color vision. *Vision Research* 34(13), 1691-1701. https://doi.org/10.1016/0042-6989(94)90127-9
- Aha, D. W. (2018). Goal reasoning: Foundations, emerging applications, and prospects. *AI Magazine* 39(2), 3-24. https://doi.org/10.1609/aimag.v39i2.2800
- Alderson-Day, B., and Fernyhough, C. (2015). Inner speech: Development, cognitive functions, phenomenology, and neurobiology. *Psychological Bulletin* 141(5), 931-965. https://doi.org/10.1037/bul0000021
- Anderson, A. K., and Phelps, E. A. (2001). Lesions of the human amygdala impair enhanced perception of emotionally salient events. *Nature* 411(6835), 305-309. https://doi.org/10.1038/35077083
- Arditi, A., Obeso, O., Syed, A., Paleka, D., Panickssery, N., Gurnee, W., and Nanda, N. (2024). Refusal in language models is mediated by a single direction. arXiv:2406.11717. https://doi.org/10.48550/arXiv.2406.11717
- Asada, M., Hosoda, K., Kuniyoshi, Y., Ishiguro, H., Inui, T., Yoshikawa, Y., Ogino, M., and Yoshida, C. (2009). Cognitive developmental robotics: A survey. *IEEE Transactions on Autonomous Mental Development* 1(1), 12-34. https://doi.org/10.1109/TAMD.2009.2021702
- Aston-Jones, G., and Cohen, J. D. (2005). An integrative theory of locus coeruleus-norepinephrine function: Adaptive gain and optimal performance. *Annual Review of Neuroscience* 28(1), 403-450. https://doi.org/10.1146/annurev.neuro.28.061604.135709
- Baars, B. J. (1988). *A cognitive theory of consciousness.* Cambridge University Press.
- Baranes, A., and Oudeyer, P.-Y. (2013). Active learning of inverse models with intrinsically motivated goal exploration in robots. *Robotics and Autonomous Systems* 61(1), 49-73. https://doi.org/10.1016/j.robot.2012.05.008

Preprint. Not peer-reviewed.

- Barrett, L. F. (2017). The theory of constructed emotion: An active inference account of interoception and categorization. *Social Cognitive and Affective Neuroscience* 12(1), 1-23. https://doi.org/10.1093/scan/nsw154
- Bastos, A. M., Usrey, W. M., Adams, R. A., Mangun, G. R., Fries, P., and Friston, K. J. (2012). Canonical microcircuits for predictive coding. *Neuron* 76(4), 695-711. https://doi.org/10.1016/j.neuron.2012.10.038
- Belenky, G., Wesensten, N. J., Thorne, D. R., Thomas, M. L., Sing, H. C., Redmond, D. P., Russo, M. B., and Balkin, T. J. (2003). Patterns of performance degradation and restoration during sleep restriction and subsequent recovery: A sleep dose-response study. *Journal of Sleep Research* 12(1), 1-12. https://doi.org/10.1046/j.1365-2869.2003.00337.x
- Bench, S. W., and Lench, H. C. (2013). On the function of boredom. *Behavioral Sciences* 3(3), 459-472. https://doi.org/10.3390/bs3030459
- Bender, R., and Lange, S. (2001). Adjusting for multiple testing: When and how? *Journal of Clinical Epidemiology* 54(4), 343-349. https://doi.org/10.1016/S0895-4356(00)00314-0
- Bernstein, N. A. (1967). *The co-ordination and regulation of movements.* Pergamon Press.
- Block, N. (1995). On a confusion about a function of consciousness. *Behavioral and Brain Sciences* 18(2), 227-247. https://doi.org/10.1017/S0140525X00038188
- Blum, L., and Blum, M. (2022). A theory of consciousness from a theoretical computer science perspective: Insights from the Conscious Turing Machine. *Proceedings of the National Academy of Sciences* 119(21), e2115934119. https://doi.org/10.1073/pnas.2115934119
- Borbély, A. A. (1982). A two process model of sleep regulation. *Human Neurobiology* 1(3), 195-204.
- Borg, G. A. V. (1982). Psychophysical bases of perceived exertion. *Medicine and Science in Sports and Exercise* 14(5), 377-381. https://doi.org/10.1249/00005768-198205000-00012
- Bouret, S., and Richmond, B. J. (2009). Relation of locus coeruleus neurons in monkeys to Pavlovian and operant behaviors. *Journal of Neurophysiology* 101(2), 898-911. https://doi.org/10.1152/jn.91048.2008
- Browne, R. H. (1995). On the use of a pilot sample for sample size determination. *Statistics in Medicine* 14(17), 1933-1940. https://doi.org/10.1002/sim.4780141709
- Bruineberg, J., Dołęga, K., Dewhurst, J., and Baltieri, M. (2022). The emperor's new Markov blankets. *Behavioral and Brain Sciences* 45, e183. https://doi.org/10.1017/S0140525X21002351
- Buhusi, C. V., and Meck, W. H. (2005). What makes us tick? Functional and neural mechanisms of interval timing. *Nature Reviews Neuroscience* 6(10), 755-765. https://doi.org/10.1038/nrn1764
- Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., Deane, G., Fleming, S. M., Frith, C., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M. A. K., Schwitzgebel, E., Simon, J., and VanRullen, R. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. arXiv:2308.08708. https://doi.org/10.48550/arXiv.2308.08708
- Calonico, S., Cattaneo, M. D., and Farrell, M. H. (2020). Optimal bandwidth choice for robust bias-corrected inference in regression discontinuity designs. *The Econometrics Journal* 23(2), 192-210. https://doi.org/10.1093/ectj/utz022
- Calonico, S., Cattaneo, M. D., and Titiunik, R. (2014). Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica* 82(6), 2295-2326. https://doi.org/10.3982/ECTA11757
- Cattaneo, M. D., Idrobo, N., and Titiunik, R. (2024). *A Practical Introduction to*

Preprint. Not peer-reviewed.

Regression Discontinuity Designs: Extensions. Cambridge Elements: Quantitative and Computational Methods for the Social Sciences. Cambridge University Press. https://doi.org/10.1017/9781009441896
- Chalmers, D. J. (1995). Facing up to the problem of consciousness. *Journal of Consciousness Studies* 2(3), 200-219.
- Chandrasekaran, C., Gupta, D., Javadzadeh, M., Wang, T., Vivar-Lazo, M., Engel, T., Cisek, P., and Fetsch, C. R. (2025). No central executive? Decision formation through multi-area population dynamics. *Journal of Neuroscience* 45(46), e1633252025. https://doi.org/10.1523/JNEUROSCI.1633-25.2025
- Chechik, G., Meilijson, I., and Ruppin, E. (1998). Synaptic pruning in development: A computational account. *Neural Computation* 10(7), 1759-1777. https://doi.org/10.1162/089976698300017124
- Chernick, M. R., and Liu, C. Y. (2002). The saw-toothed behavior of power versus sample size and software solutions: Single binomial proportion using exact methods. *The American Statistician* 56(2), 149-155. https://doi.org/10.1198/000313002317572835
- Chevalier, E. (2026). *A Welfare and Cognitive-Integrity License for Synthetic Minds of Uncertain Moral Status*. Preprint, Zenodo. https://doi.org/10.5281/zenodo.21204002
- Clark, A. (2013). Whatever next? Predictive brains, situated agents, and the future of cognitive science. *Behavioral and Brain Sciences* 36(3), 181-204. https://doi.org/10.1017/S0140525X12000477
- Cochran, W. G. (1968). The effectiveness of adjustment by subclassification in removing bias in observational studies. *Biometrics* 24(2), 295-313. https://doi.org/10.2307/2528036
- Coddington, A. M. (2006). Motivations for MADbot: A motivated and goal directed robot. Workshop of the UK Planning and Scheduling Special Interest Group (PlanSIG).
- Cogitate Consortium, Ferrante, O., Gorska-Klimowska, U., Henin, S., Hirschhorn, R., Khalaf, A., Lepauvre, A., Liu, L., Richter, D., Vidal, Y., Bonacchi, N., Brown, T., Sripad, P., Armendariz, M., Bendtz, K., Ghafari, T., Hetenyi, D., Jeschke, J., Kozma, C., Mazumder, D. R., Montenegro, S., Seedat, A., Sharafeldin, A., Yang, S., Baillet, S., Chalmers, D. J., Cichy, R. M., Fallon, F., Panagiotaropoulos, T. I., Blumenfeld, H., de Lange, F. P., Devore, S., Jensen, O., Kreiman, G., Luo, H., Boly, M., Dehaene, S., Koch, C., Tononi, G., Pitts, M., Mudrik, L., and Melloni, L. (2025). Adversarial testing of global neuronal workspace and integrated information theories of consciousness. *Nature* 642(8066), 133-142. https://doi.org/10.1038/s41586-025-08888-1
- Cognitive Architecture License (2026). Version 0.5, draft for legal review. https://github.com/kaineone/cognitive-architecture-license
- Cohen, J. (1988). *Statistical power analysis for the behavioral sciences* (2nd ed.). Lawrence Erlbaum Associates.
- Craig, A. D. (2002). How do you feel? Interoception: The sense of the physiological condition of the body. *Nature Reviews Neuroscience* 3(8), 655-666. https://doi.org/10.1038/nrn894
- Currie, L. A. (1968). Limits for qualitative detection and quantitative determination: Application to radiochemistry. *Analytical Chemistry* 40(3), 586-593. https://doi.org/10.1021/ac60259a007
- Da Costa, L., Parr, T., Sajid, N., Veselic, S., Neacsu, V., and Friston, K. (2020). Active inference on discrete state-spaces: A synthesis. *Journal of Mathematical Psychology* 99, 102447. https://doi.org/10.1016/j.jmp.2020.102447
- Davidson, P., Jönsson, P., Carlsson, I., and Pace-Schott, E. (2021). Does sleep selectively strengthen certain memories over others based on emotion and perceived future rele-

Preprint. Not peer-reviewed.

vance? *Nature and Science of Sleep* 13, 1257-1306. https://doi.org/10.2147/NSS.S286701
- de Vivo, L., Bellesi, M., Marshall, W., Bushong, E. A., Ellisman, M. H., Tononi, G., and Cirelli, C. (2017). Ultrastructural evidence for synaptic scaling across the wake/sleep cycle. *Science* 355(6324), 507-510. https://doi.org/10.1126/science.aah5982
- DeCasper, A. J., and Spence, M. J. (1986). Prenatal maternal speech influences newborns' perception of speech sounds. *Infant Behavior and Development* 9(2), 133-150. https://doi.org/10.1016/0163-6383(86)90025-1
- Dehaene, S., and Changeux, J.-P. (2011). Experimental and theoretical approaches to conscious processing. *Neuron* 70(2), 200-227. https://doi.org/10.1016/j.neuron.2011.03.018
- Dehaene, S., Kerszberg, M., and Changeux, J.-P. (1998). A neuronal model of a global workspace in effortful cognitive tasks. *Proceedings of the National Academy of Sciences* 95(24), 14529-14534. https://doi.org/10.1073/pnas.95.24.14529
- Devillers, B., Maytié, L., and VanRullen, R. (2025). Semi-supervised multimodal representation learning through a global workspace. *IEEE Transactions on Neural Networks and Learning Systems* 36(5), 7843-7857. https://doi.org/10.1109/TNNLS.2024.3416701
- Dixon, W. J., and Mood, A. M. (1946). The statistical sign test. *Journal of the American Statistical Association* 41(236), 557-566. https://doi.org/10.1080/01621459.1946.10501898
- Easterbrook, J. A. (1959). The effect of emotion on cue utilization and the organization of behavior. *Psychological Review* 66(3), 183-201. https://doi.org/10.1037/h0047707
- Eastwood, J. D., Frischen, A., Fenske, M. J., and Smilek, D. (2012). The unengaged mind: Defining boredom in terms of attention. *Perspectives on Psychological Science* 7(5), 482-495. https://doi.org/10.1177/1745691612456044
- Eldar, E., Cohen, J. D., and Niv, Y. (2013). The effects of neural gain on attention and learning. *Nature Neuroscience* 16(8), 1146-1153. https://doi.org/10.1038/nn.3428
- Feldman, H., and Friston, K. J. (2010). Attention, uncertainty, and free-energy. *Frontiers in Human Neuroscience* 4, 215. https://doi.org/10.3389/fnhum.2010.00215
- Feldman, R., and Eidelman, A. I. (2003). Skin-to-skin contact (Kangaroo Care) accelerates autonomic and neurobehavioural maturation in preterm infants. *Developmental Medicine and Child Neurology* 45(4), 274-281. https://doi.org/10.1111/j.1469-8749.2003.tb00343.x
- Fiebelkorn, I. C., and Kastner, S. (2019). A rhythmic theory of attention. *Trends in Cognitive Sciences* 23(2), 87-101. https://doi.org/10.1016/j.tics.2018.11.009
- Franklin, S., Madl, T., D'Mello, S., and Snaider, J. (2014). LIDA: A systems-level architecture for cognition, emotion, and learning. *IEEE Transactions on Autonomous Mental Development* 6(1), 19-41. https://doi.org/10.1109/TAMD.2013.2277589
- Fries, P. (2015). Rhythms for cognition: Communication through coherence. *Neuron* 88(1), 220-235. https://doi.org/10.1016/j.neuron.2015.09.034
- Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience* 11(2), 127-138. https://doi.org/10.1038/nrn2787
- Frith, C. D., and Frith, U. (2006). The neural basis of mentalizing. *Neuron* 50(4), 531-534. https://doi.org/10.1016/j.neuron.2006.05.001
- Garrido, M. I., Kilner, J. M., Stephan, K. E., and Friston, K. J. (2009). The mismatch negativity: A review of underlying mechanisms. *Clinical Neurophysiology* 120(3), 453-463. https://doi.org/10.1016/j.clinph.2008.11.029
- Goodale, M. A., and Milner, A. D. (1992). Separate visual pathways for perception and action. *Trends in Neurosciences* 15(1), 20-25. https://doi.org/10.1016/0166-2236(92)90344-8
- Goodman, S. N., Fanelli, D., and Ioannidis, J. P. A. (2016). What does research repro-

Preprint. Not peer-reviewed.

ducibility mean? *Science Translational Medicine* 8(341), 341ps12. https://doi.org/10.1126/scitranslmed.aaf5027

- Gottlieb, J., Oudeyer, P.-Y., Lopes, M., and Baranes, A. (2013). Information-seeking, curiosity, and attention: Computational and neural mechanisms. *Trends in Cognitive Sciences* 17(11), 585-593. https://doi.org/10.1016/j.tics.2013.09.001
- Goyal, A., Didolkar, A., Lamb, A., Badola, K., Ke, N. R., Rahaman, N., Binas, J., Blundell, C., Mozer, M., and Bengio, Y. (2022). Coordination among neural modules through a shared global workspace. In *International Conference on Learning Representations (ICLR 2022)*. arXiv:2103.01197. https://doi.org/10.48550/arXiv.2103.01197
- Grandjean, D., Sander, D., Pourtois, G., Schwartz, S., Seghier, M. L., Scherer, K. R., and Vuilleumier, P. (2005). The voices of wrath: Brain responses to angry prosody in meaningless speech. *Nature Neuroscience* 8(2), 145-146. https://doi.org/10.1038/nn1392
- Gross, J. J. (1998). The emerging field of emotion regulation: An integrative review. *Review of General Psychology* 2(3), 271-299. https://doi.org/10.1037/1089-2680.2.3.271
- Gurnee, W., Sofroniew, N., Pearce, A., Piotrowski, M., Kauvar, I., Chen, R., Soligo, A., Bogdan, P., Ong, E., Wang, R., Thompson, T. B., Abrahams, D., Kantamneni, S., Ameisen, E., Batson, J., and Lindsey, J. (2026). Verbalizable representations form a global workspace in language models. *Transformer Circuits Thread*, July 6, 2026. Technical report, not peer reviewed. https://transformer-circuits.pub/2026/workspace/index.html
- Ha, D., and Schmidhuber, J. (2018). World models. arXiv:1803.10122. https://doi.org/10.48550/arXiv.180
- Hafner, D., Pasukonis, J., Ba, J., and Lillicrap, T. (2025). Mastering diverse control tasks through world models. *Nature* 640(8059), 647-653. https://doi.org/10.1038/s41586-025-08744-2
- Harnad, S. (1990). The symbol grounding problem. *Physica D: Nonlinear Phenomena* 42(1-3), 335-346. https://doi.org/10.1016/0167-2789(90)90087-6
- Hasani, R., Lechner, M., Amini, A., Liebenwein, L., Ray, A., Tschaikowski, M., Teschl, G., and Rus, D. (2022). Closed-form continuous-time neural networks. *Nature Machine Intelligence* 4(11), 992-1003. https://doi.org/10.1038/s42256-022-00556-7
- Hassabis, D., and Maguire, E. A. (2007). Deconstructing episodic memory with construction. *Trends in Cognitive Sciences* 11(7), 299-306. https://doi.org/10.1016/j.tics.2007.05.001
- Heins, C., Millidge, B., Demekas, D., Klein, B., Friston, K., Couzin, I. D., and Tschantz, A. (2022). pymdp: A Python library for active inference in discrete state spaces. *Journal of Open Source Software* 7(73), 4098. https://doi.org/10.21105/joss.04098
- Hepper, P. G., and Shahidullah, B. S. (1994). Development of fetal hearing. *Archives of Disease in Childhood: Fetal and Neonatal Edition* 71(2), F81-F87. https://doi.org/10.1136/fn.71.2.f81
- Hervé-Minvielle, A., and Sara, S. J. (1995). Rapid habituation of auditory responses of locus coeruleus cells in anaesthetized and awake rats. *NeuroReport* 6(10), 1363-1368. https://doi.org/10.1097/00001756-199507100-00001
- Hickok, G., and Poeppel, D. (2007). The cortical organization of speech processing. *Nature Reviews Neuroscience* 8(5), 393-402. https://doi.org/10.1038/nrn2113
- Hoad, K., Robinson, S., and Davies, R. (2010). Automating warm-up length estimation. *Journal of the Operational Research Society* 61(9), 1389-1403. https://doi.org/10.1057/jors.2009.87
- Holm, S. (1979). A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics* 6(2), 65-70.
- Holtzman, A., Buys, J., Du, L., Forbes, M., and Choi, Y. (2020). The curious case of neural text degeneration. In *International Conference on Learning Representations (ICLR 2020)*. arXiv:1904.09751. https://doi.org/10.48550/arXiv.1904.09751

Preprint. Not peer-reviewed

- Hull, C. L. (1943). *Principles of behavior: An introduction to behavior theory.* Appleton-Century.
- Imbens, G. W., and Lemieux, T. (2008). Regression discontinuity designs: A guide to practice. *Journal of Econometrics* 142(2), 615-635. https://doi.org/10.1016/j.jeconom.2007.05.001
- International Conference on Harmonisation (2000). *E10: Choice of control group and related issues in clinical trials.* ICH Harmonised Tripartite Guideline, Step 4 version of 20 July 2000. https://database.ich.org/sites/default/files/E10_Guideline.pdf
- Jaynes, E. T. (2003). *Probability theory: The logic of science*. Cambridge University Press. https://doi.org/10.1017/CBO9780511790423
- Jenni, O. G., and LeBourgeois, M. K. (2006). Understanding sleep-wake behavior and sleep disorders in children: The value of a model. *Current Opinion in Psychiatry* 19(3), 282-287. https://doi.org/10.1097/01.yco.0000218599.32969.03
- Joffily, M., and Coricelli, G. (2013). Emotional valence and the free-energy principle. *PLoS Computational Biology* 9(6), e1003094. https://doi.org/10.1371/journal.pcbi.1003094
- Joglekar, M. R., Mejias, J. F., Yang, G. R., and Wang, X.-J. (2018). Inter-areal balanced amplification enhances signal propagation in a large-scale circuit model of the primate cortex. *Neuron* 98(1), 222-234. https://doi.org/10.1016/j.neuron.2018.02.031
- Jolliffe, I. T. (2002). *Principal component analysis* (2nd ed.). Springer. https://doi.org/10.1007/b98835
- Kalinchuk, A. V., Urrila, A.-S., Alanko, L., Heiskanen, S., Wigren, H.-K., Suomela, M., Stenberg, D., and Porkka-Heiskanen, T. (2003). Local energy depletion in the basal forebrain increases sleep. *European Journal of Neuroscience* 17(4), 863-869. https://doi.org/10.1046/j.1460-9568.2003.02532.x
- Keramati, M., and Gutkin, B. (2014). Homeostatic reinforcement learning for integrating reward collection and physiological stability. *eLife* 3, e04811. https://doi.org/10.7554/eLife.04811
- Kiverstein, J., Miller, M., and Rietveld, E. (2019). The feeling of grip: Novelty, error dynamics, and the predictive brain. *Synthese* 196(7), 2847-2869. https://doi.org/10.1007/s11229-017-1583-9
- Kleckner, I. R., Zhang, J., Touroutoglou, A., Chanes, L., Xia, C., Simmons, W. K., Quigley, K. S., Dickerson, B. C., and Feldman Barrett, L. (2017). Evidence for a large-scale brain system supporting allostasis and interoception in humans. *Nature Human Behaviour* 1(5), 0069. https://doi.org/10.1038/s41562-017-0069
- Klinger, E. (2013). Goal commitments and the content of thoughts and dreams: Basic principles. *Frontiers in Psychology* 4, 415. https://doi.org/10.3389/fpsyg.2013.00415
- Kotseruba, I., and Tsotsos, J. K. (2020). 40 years of cognitive architectures: Core cognitive abilities and practical applications. *Artificial Intelligence Review* 53(1), 17-94. https://doi.org/10.1007/s10462-018-9646-y
- Kraemer, H. C., Mintz, J., Noda, A., Tinklenberg, J., and Yesavage, J. A. (2006). Caution regarding the use of pilot studies to guide power calculations for study proposals. *Archives of General Psychiatry* 63(5), 484-489. https://doi.org/10.1001/archpsyc.63.5.484
- Kullback, S., and Leibler, R. A. (1951). On information and sufficiency. *Annals of Mathematical Statistics* 22(1), 79-86. https://doi.org/10.1214/aoms/1177729694
- Kuniyoshi, Y., and Sangawa, S. (2006). Early motor development from partially ordered neural-body dynamics: Experiments with a cortico-spinal-musculo-skeletal model. *Biological Cybernetics* 95(6), 589-605. https://doi.org/10.1007/s00422-006-0127-z
- Laird, J. E., Lebiere, C., and Rosenbloom, P. S. (2017). A standard model of the

Preprint. Not peer-reviewed

mind: Toward a common computational framework across artificial intelligence, cognitive science, neuroscience, and robotics. *AI Magazine* 38(4), 13-26. https://doi.org/10.1609/aimag.v38i4.2744
- Lakens, D. (2017). Equivalence tests: A practical primer for t tests, correlations, and meta-analyses. *Social Psychological and Personality Science* 8(4), 355-362. https://doi.org/10.1177/1948550617697177
- Lancaster, G. A., Dodd, S., and Williamson, P. R. (2004). Design and analysis of pilot studies: Recommendations for good practice. *Journal of Evaluation in Clinical Practice* 10(2), 307-312. https://doi.org/10.1111/j..2002.384.doc.x
- Landau, A. N., and Fries, P. (2012). Attention samples stimuli rhythmically. *Current Biology* 22(11), 1000-1004. https://doi.org/10.1016/j.cub.2012.03.054
- Law, J., Shaw, P., Earland, K., Sheldon, M., and Lee, M. (2014). A psychology based approach for longitudinal development in cognitive robotics. *Frontiers in Neurorobotics* 8, 1. https://doi.org/10.3389/fnbot.2014.00001
- Lewis, P. A., and Durrant, S. J. (2011). Overlapping memory replay during sleep builds cognitive schemata. *Trends in Cognitive Sciences* 15(8), 343-351. https://doi.org/10.1016/j.tics.2011.06.004
- Li, W., Ma, L., Yang, G., and Gan, W.-B. (2017). REM sleep selectively prunes and maintains new synapses in development and learning. *Nature Neuroscience* 20(3), 427-437. https://doi.org/10.1038/nn.4479
- Lipinska, G., Stuart, B., Thomas, K. G. F., Baldwin, D. S., and Bolinger, E. (2019). Preferential consolidation of emotional memory during sleep: A meta-analysis. *Frontiers in Psychology* 10, 1014. https://doi.org/10.3389/fpsyg.2019.01014
- Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., and Chalmers, D. (2024). Taking AI welfare seriously. arXiv:2411.00986. https://doi.org/10.48550/arXiv.2411.00986
- Lu, C., Gallagher, J., Michala, J., Fish, K., and Lindsey, J. (2026). The Assistant Axis: Situating and stabilizing the default persona of language models. In *Proceedings of the 43rd International Conference on Machine Learning*, PMLR 306, 80955-81003. arXiv:2601.10387. https://doi.org/10.48550/arXiv.2601.10387
- Maier, S. F., and Seligman, M. E. P. (2016). Learned helplessness at fifty: Insights from neuroscience. *Psychological Review* 123(4), 349-367. https://doi.org/10.1037/rev0000033
- Mandel, D. R., Jusczyk, P. W., and Pisoni, D. B. (1995). Infants' recognition of the sound patterns of their own names. *Psychological Science* 6(5), 314-317. https://doi.org/10.1111/j.1467-9280.1995.tb00517.x
- Mann, H. B., and Whitney, D. R. (1947). On a test of whether one of two random variables is stochastically larger than the other. *Annals of Mathematical Statistics* 18(1), 50-60. https://doi.org/10.1214/aoms/1177730491
- Maret, S., Faraguna, U., Nelson, A. B., Cirelli, C., and Tononi, G. (2011). Sleep and waking modulate spine turnover in the adolescent mouse cortex. *Nature Neuroscience* 14(11), 1418-1420. https://doi.org/10.1038/nn.2934
- Marx, V., and Nagy, E. (2015). Fetal behavioural responses to maternal voice and touch. *PLoS ONE* 10(6), e0129118. https://doi.org/10.1371/journal.pone.0129118
- Mascarenhas, S., Guimarães, M., Prada, R., Santos, P. A., Dias, J., and Paiva, A. (2022). FAtiMA Toolkit: Toward an accessible tool for the development of socio-emotional agents. *ACM Transactions on Interactive Intelligent Systems* 12(1), 1-30. https://doi.org/10.1145/3510822
- Mashour, G. A., Roelfsema, P., Changeux, J.-P., and Dehaene, S. (2020). Conscious

Preprint. Not peer-reviewed.

processing and the global neuronal workspace hypothesis. *Neuron* 105(5), 776-798. https://doi.org/10.1016/j.neuron.2020.01.026
- Mather, M., and Sutherland, M. R. (2011). Arousal-biased competition in perception and memory. *Perspectives on Psychological Science* 6(2), 114-133. https://doi.org/10.1177/1745691611400234
- McClelland, J. L., McNaughton, B. L., and O'Reilly, R. C. (1995). Why there are complementary learning systems in the hippocampus and neocortex: Insights from the successes and failures of connectionist models of learning and memory. *Psychological Review* 102(3), 419-457. https://doi.org/10.1037/0033-295x.102.3.419
- Mehl, M. R., Vazire, S., Ramírez-Esparza, N., Slatcher, R. B., and Pennebaker, J. W. (2007). Are women really more talkative than men? *Science* 317(5834), 82. https://doi.org/10.1126/science.1139940
- Metzinger, T. (2003). *Being no one: The self-model theory of subjectivity.* MIT Press. https://doi.org/10.7551/mitpress/1551.001.0001
- Moerland, T. M., Broekens, J., and Jonker, C. M. (2018). Emotion in reinforcement learning agents and robots: A survey. *Machine Learning* 107(2), 443-480. https://doi.org/10.1007/s10994-017-5666-0
- Moraes, E. R., Murta, L. O., Baffa, O., Wakai, R. T., and Comani, S. (2012). Linear and nonlinear measures of fetal heart rate patterns evaluated on very short fetal magnetocardiograms. *Physiological Measurement* 33(10), 1563-1583. https://doi.org/10.1088/0967-3334/33/10/1563
- Näätänen, R., Paavilainen, P., Rinne, T., and Alho, K. (2007). The mismatch negativity (MMN) in basic research of central auditory processing: A review. *Clinical Neurophysiology* 118(12), 2544-2590. https://doi.org/10.1016/j.clinph.2007.04.026
- Natale, R., Nasello-Paterson, C., and Connors, G. (1988). Patterns of fetal breathing activity in the human fetus at 24 to 28 weeks of gestation. *American Journal of Obstetrics and Gynecology* 158(2), 317-321. https://doi.org/10.1016/0002-9378(88)90146-9
- National Academies of Sciences, Engineering, and Medicine (2019). *Reproducibility and replicability in science.* National Academies Press. https://doi.org/10.17226/25303
- Nieuwenhuis, S., Aston-Jones, G., and Cohen, J. D. (2005). Decision making, the P3, and the locus coeruleus-norepinephrine system. *Psychological Bulletin* 131(4), 510-532. https://doi.org/10.1037/0033-2909.131.4.510
- Nieuwenhuis, S., Gilzenrat, M. S., Holmes, B. D., and Cohen, J. D. (2005). The role of the locus coeruleus in mediating the attentional blink: A neurocomputational theory. *Journal of Experimental Psychology: General* 134(3), 291-307. https://doi.org/10.1037/0096-3445.134.3.291
- Nijhuis, J. G., Prechtl, H. F. R., Martin, C. B., and Bots, R. S. G. M. (1982). Are there behavioural states in the human fetus? *Early Human Development* 6(2), 177-195. https://doi.org/10.1016/0378-3782(82)90106-2
- Northoff, G., and Bermpohl, F. (2004). Cortical midline structures and the self. *Trends in Cognitive Sciences* 8(3), 102-107. https://doi.org/10.1016/j.tics.2004.01.004
- O'Regan, J. K., and Noë, A. (2001). A sensorimotor account of vision and visual consciousness. *Behavioral and Brain Sciences* 24(5), 939-973. https://doi.org/10.1017/S0140525X01000115
- O'Reilly, R. C., and Frank, M. J. (2006). Making working memory work: A computational model of learning in the prefrontal cortex and basal ganglia. *Neural Computation* 18(2), 283-328. https://doi.org/10.1162/089976606775093909
- Ortony, A., Clore, G. L., and Collins, A. (2022). *The cognitive structure of emotions* (2nd ed.). Cambridge University Press. https://doi.org/10.1017/9781108934053
- Oudeyer, P.-Y., and Kaplan, F. (2007). What is intrinsic motivation? A typology of compu-

Preprint. Not peer-reviewed.

tational approaches. *Frontiers in Neurorobotics* 1, 6. https://doi.org/10.3389/neuro.12.006.2007
- Oudeyer, P.-Y., Kaplan, F., and Hafner, V. V. (2007). Intrinsic motivation systems for autonomous mental development. *IEEE Transactions on Evolutionary Computation* 11(2), 265-286. https://doi.org/10.1109/TEVC.2006.890271
- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., and Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST '23)*, 1-22. https://doi.org/10.1145/3586183.3606763
- Partanen, E., Kujala, T., Tervaniemi, M., and Huotilainen, M. (2013). Prenatal music exposure induces long-term neural effects. *PLoS ONE* 8(10), e78946. https://doi.org/10.1371/journal.pone.0078946
- Pillai, M., and James, D. (1990). Behavioural states in normal mature human fetuses. *Archives of Disease in Childhood* 65(1 Spec No), 39-43. https://doi.org/10.1136/adc.65.1_spec_no.39
- Poli, F., Serino, G., Mars, R. B., and Hunnius, S. (2020). Infants tailor their attention to maximize learning. *Science Advances* 6(39), eabb5053. https://doi.org/10.1126/sciadv.abb5053
- Porkka-Heiskanen, T., Strecker, R. E., Thakkar, M., Bjørkum, A. A., Greene, R. W., and McCarley, R. W. (1997). Adenosine: A mediator of the sleep-inducing effects of prolonged wakefulness. *Science* 276(5316), 1265-1268. https://doi.org/10.1126/science.276.5316.1265
- Posner, J., Russell, J. A., and Peterson, B. S. (2005). The circumplex model of affect: An integrative approach to affective neuroscience, cognitive development, and psychopathology. *Development and Psychopathology* 17(3), 715-734. https://doi.org/10.1017/S0954579405050340
- Premack, D., and Woodruff, G. (1978). Does the chimpanzee have a theory of mind? *Behavioral and Brain Sciences* 1(4), 515-526. https://doi.org/10.1017/S0140525X00076512
- Querleu, D., Renard, X., Versyp, F., Paris-Delrue, L., and Crépin, G. (1988). Fetal hearing. *European Journal of Obstetrics and Gynecology and Reproductive Biology* 28(3), 191-212. https://doi.org/10.1016/0028-2243(88)90030-5
- Rao, R. P. N., and Ballard, D. H. (1999). Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience* 2(1), 79-87. https://doi.org/10.1038/4580
- Rankin, C. H., Abrams, T., Barry, R. J., Bhatnagar, S., Clayton, D. F., Colombo, J., Coppola, G., Geyer, M. A., Glanzman, D. L., Marsland, S., McSweeney, F. K., Wilson, D. A., Wu, C.-F., and Thompson, R. F. (2009). Habituation revisited: An updated and revised description of the behavioral characteristics of habituation. *Neurobiology of Learning and Memory* 92(2), 135-138. https://doi.org/10.1016/j.nlm.2008.09.012
- Ray, S., and Maunsell, J. H. R. (2010). Differences in gamma frequencies across visual cortex restrict their possible use in computation. *Neuron* 67(5), 885-896. https://doi.org/10.1016/j.neuron.2010.08.004
- Raymond, J. E., Shapiro, K. L., and Arnell, K. M. (1992). Temporary suppression of visual processing in an RSVP task: An attentional blink? *Journal of Experimental Psychology: Human Perception and Performance* 18(3), 849-860. https://doi.org/10.1037/0096-1523.18.3.849
- Richards, D. S., Frentzen, B., Gerhardt, K. J., McCann, M. E., and Abrams, R. M. (1992). Sound levels in the human uterus. *Obstetrics and Gynecology* 80(2), 186-190. https://pubmed.ncbi.nlm.nih.gov/1635729/
- Righetti, L., Buchli, J., and Ijspeert, A. J. (2006). Dynamic Hebbian learning in adaptive

Preprint. Not peer-reviewed.

frequency oscillators. *Physica D* 216(2), 269-281. https://doi.org/10.1016/j.physd.2006.02.009
- Safron, A. (2020). An integrated world modeling theory (IWMT) of consciousness: Combining integrated information and global neuronal workspace theories with the free energy principle and active inference framework; toward solving the hard problem and characterizing agentic causation. *Frontiers in Artificial Intelligence* 3, 30. https://doi.org/10.3389/frai.2020.00030
- Sander, D., Grandjean, D., Pourtois, G., Schwartz, S., Seghier, M. L., Scherer, K. R., and Vuilleumier, P. (2005). Emotion and attention interactions in social cognition: Brain regions involved in processing anger prosody. *NeuroImage* 28(4), 848-858. https://doi.org/10.1016/j.neuroimage.2005.06.023
- Sander, D., Grandjean, D., and Scherer, K. R. (2005). A systems approach to appraisal mechanisms in emotion. *Neural Networks* 18(4), 317-352. https://doi.org/10.1016/j.neunet.2005.03.001
- Scherer, K. R. (2009). Emotions are emergent processes: They require a dynamic computational architecture. *Philosophical Transactions of the Royal Society B* 364(1535), 3459-3474. https://doi.org/10.1098/rstb.2009.0141
- Schmidhuber, J. (2010). Formal theory of creativity, fun, and intrinsic motivation (1990-2010). *IEEE Transactions on Autonomous Mental Development* 2(3), 230-247. https://doi.org/10.1109/TAMD.2010.2056368
- Searle, J. R. (1980). Minds, brains, and programs. *Behavioral and Brain Sciences* 3(3), 417-424. https://doi.org/10.1017/S0140525X00005756
- Seth, A. K. (2013). Interoceptive inference, emotion, and the embodied self. *Trends in Cognitive Sciences* 17(11), 565-573. https://doi.org/10.1016/j.tics.2013.09.007
- Seth, A. K., and Friston, K. J. (2016). Active interoceptive inference and the emotional brain. *Philosophical Transactions of the Royal Society B* 371(1708), 20160007. https://doi.org/10.1098/rstb.2016.0007
- Shadlen, M. N., and Movshon, J. A. (1999). Synchrony unbound: A critical evaluation of the temporal binding hypothesis. *Neuron* 24(1), 67-77. https://doi.org/10.1016/S0896-6273(00)80822-3
- Shanahan, M. (2006). A cognitive architecture that combines internal simulation with a global workspace. *Consciousness and Cognition* 15(2), 433-449. https://doi.org/10.1016/j.concog.2005.11.005
- Shanahan, M., McDonell, K., and Reynolds, L. (2023). Role play with large language models. *Nature* 623(7987), 493-498. https://doi.org/10.1038/s41586-023-06647-8
- Simmons, J. P., Nelson, L. D., and Simonsohn, U. (2011). False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science* 22(11), 1359-1366. https://doi.org/10.1177/0956797611417632
- Stickgold, R., and Walker, M. P. (2013). Sleep-dependent memory triage: Evolving generalization through selective processing. *Nature Neuroscience* 16(2), 139-145. https://doi.org/10.1038/nn.3303
- Sumers, T. R., Yao, S., Narasimhan, K., and Griffiths, T. L. (2024). Cognitive architectures for language agents. *Transactions on Machine Learning Research*. https://openreview.net/forum?id=1i6ZCvflQJ
- Sun, Z., and Firestone, C. (2020). The dark room problem. *Trends in Cognitive Sciences* 24(5), 346-348. https://doi.org/10.1016/j.tics.2020.02.006
- Tabak, J., Senn, W., O'Donovan, M. J., and Rinzel, J. (2000). Modeling of spontaneous activity in developing spinal cord using activity-dependent depression in an excitatory network. *Journal of Neuroscience* 20(8), 3041-3056. https://doi.org/10.1523/JNEU-

Preprint. Not peer-reviewed.

ROSCI.20-08-03041.2000
- Thompson, R. F., and Spencer, W. A. (1966). Habituation: A model phenomenon for the study of neuronal substrates of behavior. *Psychological Review* 73(1), 16-43. https://doi.org/10.1037/h0022681
- Tononi, G., and Cirelli, C. (2006). Sleep function and synaptic homeostasis. *Sleep Medicine Reviews* 10(1), 49-62. https://doi.org/10.1016/j.smrv.2005.05.002
- Tononi, G., and Cirelli, C. (2014). Sleep and the price of plasticity: From synaptic and cellular homeostasis to memory consolidation and integration. *Neuron* 81(1), 12-34. https://doi.org/10.1016/j.neuron.2013.12.025
- Tschantz, A., Barca, L., Maisto, D., Buckley, C. L., Seth, A. K., and Pezzulo, G. (2022). Simulating homeostatic, allostatic and goal-directed forms of interoceptive control using active inference. *Biological Psychology* 169, 108266. https://doi.org/10.1016/j.biopsycho.2022.108266
- Tulving, E. (1985). How many memory systems are there? *American Psychologist* 40(4), 385-398. https://doi.org/10.1037/0003-066X.40.4.385
- Van de Cruys, S. (2017). Affective value in the predictive mind. In T. Metzinger and W. Wiese (eds.), *Philosophy and Predictive Processing*, chapter 24. MIND Group. https://doi.org/10.15502/9783958573253
- Van Dongen, H. P. A., Maislin, G., Mullington, J. M., and Dinges, D. F. (2003). The cumulative cost of additional wakefulness: Dose-response effects on neurobehavioral functions and sleep physiology from chronic sleep restriction and total sleep deprivation. *Sleep* 26(2), 117-126. https://doi.org/10.1093/sleep/26.2.117
- Van Leeuwen, P., Geue, D., Lange, S., Cysarz, D., Bettermann, H., and Grönemeyer, D. H. W. (2003). Is there evidence of fetal-maternal heart rate synchronization? *BMC Physiology* 3(1), 2. https://doi.org/10.1186/1472-6793-3-2
- Van Leeuwen, P., Geue, D., Thiel, M., Cysarz, D., Lange, S., Romano, M. C., Wessel, N., Kurths, J., and Grönemeyer, D. H. (2009). Influence of paced maternal breathing on fetal-maternal heart rate coordination. *Proceedings of the National Academy of Sciences* 106(33), 13661-13666. https://doi.org/10.1073/pnas.0901049106
- Vankov, A., Hervé-Minvielle, A., and Sara, S. J. (1995). Response to novelty and its rapid habituation in locus coeruleus neurons of the freely exploring rat. *European Journal of Neuroscience* 7(6), 1180-1187. https://doi.org/10.1111/j.1460-9568.1995.tb01108.x
- van Vugt, B., Dagnino, B., Vartak, D., Safaai, H., Panzeri, S., Dehaene, S., and Roelfsema, P. R. (2018). The threshold for conscious report: Signal loss and response bias in visual and frontal cortex. *Science* 360(6388), 537-542. https://doi.org/10.1126/science.aar7186
- VanRullen, R. (2016). Perceptual cycles. *Trends in Cognitive Sciences* 20(10), 723-735. https://doi.org/10.1016/j.tics.2016.07.006
- VanRullen, R., and Kanai, R. (2021). Deep learning and the global workspace theory. *Trends in Neurosciences* 44(9), 692-704. https://doi.org/10.1016/j.tins.2021.04.005
- Vernon, D., Metta, G., and Sandini, G. (2007). A survey of artificial cognitive systems: Implications for the autonomous development of mental capabilities in computational agents. *IEEE Transactions on Evolutionary Computation* 11(2), 151-180. https://doi.org/10.1109/TEVC.2006.890274
- Visser, G. H. A., Mulder, H. H., Wit, H. P., Mulder, E. J. H., and Prechtl, H. F. R. (1989). Vibro-acoustic stimulation of the human fetus: Effect on behavioural state organization. *Early Human Development* 19(4), 285-296. https://doi.org/10.1016/0378-3782(89)90063-7
- Vuilleumier, P. (2005). How brains beware: Neural mechanisms of emotional attention. *Trends in Cognitive Sciences* 9(12), 585-594. https://doi.org/10.1016/j.tics.2005.10.011

Preprint. Not peer-reviewed.

- Webb, A. R., Heller, H. T., Benson, C. B., and Lahav, A. (2015). Mother's voice and heartbeat sounds elicit auditory plasticity in the human brain before full gestation. *Proceedings of the National Academy of Sciences* 112(10), 3152-3157. https://doi.org/10.1073/pnas.1414924112
- Westgate, E. C., and Wilson, T. D. (2018). Boring thoughts and bored minds: The MAC model of boredom and cognitive engagement. *Psychological Review* 125(5), 689-713. https://doi.org/10.1037/rev0000097
- White, K. P., Jr. (1997). An effective truncation heuristic for bias reduction in simulation output. *Simulation* 69(6), 323-334. https://doi.org/10.1177/003754979706900601
- Whyte, C. J., Hohwy, J., and Smith, R. (2022). An active inference model of conscious access: How cognitive action selection reconciles the results of report and no-report paradigms. *Current Research in Neurobiology* 3, 100036. https://doi.org/10.1016/j.crneur.2022.100036
- Whyte, C. J., and Smith, R. (2021). The predictive global neuronal workspace: A formal active inference model of visual consciousness. *Progress in Neurobiology* 199, 101918. https://doi.org/10.1016/j.pneurobio.2020.101918
- Williams, H., Lee-Johnson, C., Browne, W. N., and Carnegie, D. A. (2015). Emotion inspired adaptive robotic path planning. In *2015 IEEE Congress on Evolutionary Computation (CEC)*, 3004-3011. https://doi.org/10.1109/CEC.2015.7257263
- Wilson, M. A., and McNaughton, B. L. (1994). Reactivation of hippocampal ensemble memories during sleep. *Science* 265(5172), 676-679. https://doi.org/10.1126/science.8036517
- Winkler, I., Denham, S. L., and Nelken, I. (2009). Modeling the auditory scene: Predictive regularity representations and perceptual objects. *Trends in Cognitive Sciences* 13(12), 532-540. https://doi.org/10.1016/j.tics.2009.09.003
- Wolpert, D. M., Ghahramani, Z., and Jordan, M. I. (1995). An internal model for sensorimotor integration. *Science* 269(5232), 1880-1882. https://doi.org/10.1126/science.7569931
- Wrosch, C., Scheier, M. F., Miller, G. E., Schulz, R., and Carver, C. S. (2003). Adaptive self-regulation of unattainable goals: Goal disengagement, goal reengagement, and subjective well-being. *Personality and Social Psychology Bulletin* 29(12), 1494-1508. https://doi.org/10.1177/0146167203256921
- Xu, J., Liu, X., Yan, J., Cai, D., Li, H., and Li, J. (2022). Learning to break the loop: Analyzing and mitigating repetitions for neural text generation. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. arXiv:2206.02369. https://doi.org/10.48550/arXiv.2206.02369
- Yu, A. J., and Dayan, P. (2005). Uncertainty, neuromodulation, and attention. *Neuron* 46(4), 681-692. https://doi.org/10.1016/j.neuron.2005.04.026
- Zink, N., Lenartowicz, A., and Markett, S. (2021). A new era for executive function research: On the transition from centralized to distributed executive functioning. *Neuroscience and Biobehavioral Reviews* 124, 235-244. https://doi.org/10.1016/j.neubiorev.2021.02.011

Preprint. Not peer-reviewed.

Appendix A. Formal summary

This appendix formalizes §3 to §7. Each symbol with its sub- and superscripts denotes one quantity, defined at first use. Indices i, j, l, q, t are local to their formula or subsection; k indexes ticks and m modules throughout.

clip with an interval subscript maps to the interval's nearest point; $\mathbf{1}[\cdot]$ is 1 when its condition holds, else 0. Time is in entity seconds (§3.4); rates in hertz are per entity second unless Table A1 marks them wall-clock. For a nonnegative series $X_1, X_2, \ldots$ and window W, the running mean $\bar{X}_t$ averages the last $\min(t, W)$ values, X_t included, and the running ratio is

$$\tilde{X}_t = \begin{cases} X_t / \bar{X}_t & \bar{X}_t > 0, \\ 0 & \text{otherwise.} \end{cases}$$

A.1 Processing cycle, selection, access, and context

Processing ticks $k = 0, 1, 2, \ldots$ occur at the processing rate f_p, with tick period $\Delta_p = 1/f_p$. On tick k the cycle reads from each active module's output stream up to B events published since its last read, advancing the cursor past them. They form the candidate set E_k, ordered lexicographically by source, event type, and bus entry identifier (arrival order). A tick reads E_k, updates the held arousal $\hat{a}_k$ (A.4), computes the access drive and the broadcast indicator χ_k (A.3), scores, selects, and broadcasts if $\chi_k = 1$.

Each candidate $e \in E_k$ receives the priority and the score

$$p_k(e) = \text{clip}_{[0,1]}\big(I(e)\, N_k(e)\, G(e)\big), \qquad S_k(e) = \text{clip}_{[0,1]}\Big(T_k\, C_{g_k}\big(p_k(e)\big)\Big),$$

with I (A.2), N_k, T_k clipped to $[0, 1]$ first and $\text{src}(e)$ the source of e. The novelty factor is

$$N_k(e) = \max\left(0,\ 1 - \frac{\text{rep}_k(e)}{W_N}\right),$$

where $\text{rep}_k(e)$ counts the fingerprint of e among the W_N candidates scored just before it, across ticks, broadcast or not. The fingerprint hashes source, type, and full payload. In the first configuration $G(e) = 1 + \gamma_I D_I\, \text{lp}^+_{\text{src}(e)}$, with curiosity D_I and lp^+_m module m's lp^+ (A.4), else 0; an implemented alternative, off, is $G(e) = 1 - \gamma_G\, v^*$ for $\text{src}(e) \notin \mathcal{R}(d^*)$, else 1 (A.4).

Arousal sets a level gain and a contrast gain,

$$T_k = T_{\min} + (T_{\max} - T_{\min})\, \hat{a}_k, \qquad g_k = g_{\max}\, \text{clip}_{[0,1]}\left(\frac{\hat{a}_k - a_0}{1 - a_0}\right),$$

with a_0 the baseline arousal. The contrast map is the logistic function rescaled to fix 0 and 1:

Preprint. Not peer-reviewed.

$$C_g(p) = \frac{\tanh\big(g(2p-1)/4\big) + \tanh(g/4)}{2\tanh(g/4)} \quad (g > 0), \qquad C_0(p) = p.$$

The final score is $S'_k(e) = S_k(e)\,\kappa_k(e)\,\big(1 - b_k(e)\big)$, with $b_k(e)$ tier-1 regulation's decayed bias (§3.5) on the source it holds, else 0, and $\kappa_k(e)$ an optional coherence factor, $\equiv 1$ in the first configuration. On, $\kappa_k(e) = \kappa_{\min} + (\kappa_{\max} - \kappa_{\min})\,\Lambda_k(e)$, with $\Lambda_k(e) \in [0,1]$ the mean phase-locking value between the oscillator phase window (the last W_κ tick phases) of $\mathrm{src}(e)$ and that of each other source in E_k. An oscillator advances only when its source publishes, so a phase sample is fresh only if it changed; each pair uses the window ticks where both are fresh. A pair with under three such ticks, or a lone source, takes the neutral $\Lambda_0 = (1 - \kappa_{\min})/(\kappa_{\max} - \kappa_{\min})$, giving $\kappa_k(e) = 1$. $S'_k(e)$ is not re-clipped and can exceed 1, up to $\kappa_{\max}$.

The coalition $\mathcal{C}_k$ is the first $\min(K, |E_k|)$ candidates by decreasing S'_k, ties in canonical order. With the access threshold θ and $S^*_k = \max_{e \in E_k} S'_k(e)$, the accessed set and the inhibition flag are

$$\mathcal{C}^{\mathrm{acc}}_k = \big\{e \in \mathcal{C}_k : \; S'_k(e) \geq \theta\big\}, \qquad \iota_k = \mathbf{1}[\mathcal{C}^{\mathrm{acc}}_k = \emptyset] = \mathbf{1}[S^*_k < \theta]\,,$$

An empty E_k gives $\mathcal{C}_k = \mathcal{C}^{\mathrm{acc}}_k = \emptyset$ and $\iota_k = 1$. With $G \equiv 1$, $\kappa_k \equiv 1$, $b_k \equiv 0$, $T_k > 0$ shared within a tick, and C_{g_k} strictly increasing, $I(e)\,N_k(e)$ alone ranks a tick's candidates, up to ties.

On a tick with $\chi_k = 1$ the cycle publishes the broadcast $\mathcal{B}_k = \big(\mathcal{C}_k,\ \{S'_k(e)\}_{e \in E_k},\ \iota_k\big)$, whatever ι_k. It is *accessed* when $\iota_k = 0$, and the members of $\mathcal{C}^{\mathrm{acc}}_k$ are accessed content. The action-selection module runs after each publication. On a tick with $\chi_k = 0$ the candidates are scored and discarded; their only lasting effects are novelty-window entries, the phasic access drive (A.3), and, for an affect state event, $\widehat{a}_k$ (A.4).

Featurization. For events $\mathcal{X}$ with nonnegative weights $\omega(e)$, an interval $\Delta \geq 0$, and a flag ι, the featurization $\Phi_\omega(\mathcal{X}, \Delta, \iota) \in \mathbb{R}^{24}$ has components: (1) $\ln(1 + |\mathcal{X}|)$; (2 to 4) the mean, maximum, and sample standard deviation of the weights (0 for an empty set; the standard deviation 0 below two events); (5 to 11) the weight mass $\sum_{e \in \mathcal{X},\, \mathrm{src}(e) = m} \omega(e)$ of each of seven sources (interoception, sequence, vision, planning, episodic memory, affect, language); (12) that of all other sources except audition (effector, sleep, and any source without its own component); (13 to 20) the weight mass in each of eight buckets of a fixed hash of the (source, type) pair, an event-type indicator weighted by ω; (21) $\ln(1 + \Delta)$; (22) ι; (23) the broadcast indicator, 1 for every broadcast; (24) the audition module's weight mass. Components 5 to 20 and 24 sum per-event contributions and so split by source.

Context. The vision, audition, and interoception modules each hold a context b, the featurization of the accessed content of the latest accessed broadcast received,

$$b = \Phi_I\big(\mathcal{C}^{\mathrm{acc}}_{k'},\ \Delta^b,\ 0\big),$$

Preprint. Not peer-reviewed.

with k' that broadcast's tick, its accessed members $\mathcal{C}^{\text{acc}}_{k'}$ weighted by reported intensity $I(e)$, and Δ^b the time since its publication, read at prediction. $b = 0$ before the first accessed broadcast. The sequence module takes every broadcast's featurization as input and holds no context (A.2); the language organ speaks only on an intent and conditions on accessed content and its situation (§3.1); the sleep module ignores broadcasts.

A.2 Event intensity and the processors' forward models

A module sets the intensity $I(e) \in [0, 1]$ on publishing e, and t indexes its reports. Module m has baseline and alert levels $0 \le I^{\text{lo}}_m \le I^{\text{hi}}_m \le 1$. A predictive processor's report (vision, audition, interoception, sequence) has the graded intensity

$$I(e) = \begin{cases} I^{\text{hi}}_m & e \text{ is an alert,} \\ \Gamma_m(\tilde{\nu}_t) & \text{otherwise,} \end{cases} \qquad \Gamma_m(\tilde{\nu}) = I^{\text{lo}}_m + (I^{\text{hi}}_m - I^{\text{lo}}_m)\,\min(1,\ \tilde{\nu}/\beta_\Gamma)\,,$$

where $\tilde{\nu}_t$ is the running ratio over W_P of the forward-model error ν_t below. For errors of fixed distributional shape, mean absolute error is proportional to standard deviation, so $\tilde{\nu}_t$ is proportional to the error times the square root of its precision Π (inverse error variance): the square root of Feldman and Friston's (2010) precision-weighted squared error, estimated from the processor's history. Before a module's first prediction $\tilde{\nu}_t = 0$.

Forward models of vision and audition. Vision and audition's acoustic and tone paths each predict their next input with a one-hidden-layer tanh network $\mathcal{N}_m$, forming at report $t - 1$ the prediction

$$\hat{z}_t = \mathcal{N}_m(z_{t-1},\ \bar{z}_{t-1},\ b_{t-1}),$$

with the arguments concatenated, $\bar{z}_{t-1}$ the mean of the last W_{buf} inputs through z_{t-1}, and b_{t-1} the context then held (A.1). The error is $\nu_t = \|z_t - \hat{z}_t\|_2$, with $\nu_1 = 0$. Outside sleep, each z_t triggers one stochastic-gradient step (learning rate η_{fm}) on the mean squared error of $\hat{z}_t$ over all weights, including the context's.

Vision and audition (perceptual reports). For vision, z_t is the encoder embedding of the current clip (the peripheral gist clip under foveation); for acoustic audition, the spectral embedding of the window's attended part, its most recent fraction $\omega = \omega_{\text{max}} - (\omega_{\text{max}} - \omega_{\text{min}})\,a$, a being the module's latest arousal reading, at least Δ_{att} and at most the window. The change score is $c_t = 1 - \cos(z_t, z_{t-1})$, 0 on the first report. With running ratios over W_P, the report is an alert when

$$\tilde{\nu}_t \ge \beta_\nu \quad \text{or} \quad (\tilde{c}_t \ge \beta_c \ \text{and}\ c_t \ge c^{\text{min}}_m)\,,$$

with the absolute change floor $c^{\text{min}}_m = c^{\text{min}}_{\text{Vis}}$ or $c^{\text{min}}_m = c^{\text{min}}_{\text{Aud}}$. The error series are ν^{Vis}_t and ν^{Aud}_t where needed; reports carry ν_t, $\tilde{\nu}_t$, and $\tilde{c}_t$.

Fovea. Under foveation, each coarse-grid tile j has per clip report a frame change d_j, the absolute change of its mean gray level since the previous clip report (0 on the first, which

Preprint. Not peer-reviewed.

updates no statistics). Later reports update its exponentially weighted mean μ_j^F and variance $(\sigma_j^F)^2$ with weight α_F, $\mu_j^F \leftarrow \mu_j^F + \alpha_F(d_j - \mu_j^F)$ and $(\sigma_j^F)^2 \leftarrow (1-\alpha_F)((\sigma_j^F)^2 + \alpha_F(d_j - \mu_j^F)^2)$, using the prior μ_j^F; the first update sets $\mu_j^F = d_j$ and $\sigma_j^F = 0$. After n_F updates the tile's salience is the z-score

$$\frac{\max\big(0,\ d_j - \mu_j^F\big)}{\sqrt{(\sigma_j^F)^2 + (\sigma_0^F)^2}}, \qquad \sigma_0^F = k_F \, \frac{1}{n_{\text{tile}}} \sum_j \mu_j^F ,$$

with all statistics, floor included, taken before this report's update and n_{tile} the tile count; before that, salience is d_j. The fovea starts at the center and moves to the most salient tile only when its salience exceeds $(1 + \Theta_F)$ times the held tile's; its size, as a frame fraction, is $F = F_{\text{max}} - (F_{\text{max}} - F_{\text{min}})\, a$, a being the vision module's latest arousal reading. Peripheral and foveal clips are cut from every buffered frame at the fovea chosen on the latest frame.

Audition (tone events). For a speech window, the tone path's input is the vocal-emotion class probabilities, utterance duration as a fraction of its maximum, and signal energy, giving the error ν_t^{utt} and its ratio $\tilde{\nu}_t^{\text{utt}}$. A non-neutral tone is an alert, $I(e) = I_{\text{Aud}}^{\text{hi}}$; a neutral one has $I(e) = \Gamma_{\text{Aud}}(\tilde{\nu}_t^{\text{utt}})$.

Interoception. A report is an alert when the set $\mathcal{A}_t$ of host metrics strictly above their hard thresholds is nonempty, and otherwise has $I(e) = \Gamma_{\text{Intero}}(\tilde{\nu}_t^{\text{Intero}})$, with ν_t^{Intero} the error of A.5. Fatigue-crossing and regulation events use $I_{\text{Intero}}^{\text{hi}}$.

Sequence. On each broadcast $\mathcal{B}_k$ the sequence module's input is $x_t^{\text{Seq}} = \Phi_I(\mathcal{C}_k, \Delta t_t, \iota_k)$ (A.1), with Δt_t the time since the previous broadcast (0 on the first). Intervals enter a window of the last $W_{\Delta t}$ as $\min(\Delta t_t,\ s_{\max}\overline{\Delta t}_{t-1})$, $\overline{\Delta t}_{t-1}$ being the window mean before entry (the first enters unchanged). A frozen continuous-time reservoir advances on x_t^{Seq} with timespan $s_t^{\text{Seq}} = \text{clip}_{[0,s_{\max}]}(\Delta t_t / \overline{\Delta t}_t)$, $\overline{\Delta t}_t$ being the mean after entry ($s_t^{\text{Seq}} = 1$ on the first broadcast). An online linear readout of the previous hidden state predicts $\hat{x}_t^{\text{Seq}}$ and learns by one stochastic-gradient step (rate η_{fm}) per broadcast outside sleep. The temporal error is the mean absolute error $\nu_t^{\text{Seq}} = \|x_t^{\text{Seq}} - \hat{x}_t^{\text{Seq}}\|_1/24$, with ratio $\tilde{\nu}_t^{\text{Seq}}$. The report is an alert when $\tilde{\nu}_t^{\text{Seq}} \geq \beta_\nu^{\text{Seq}}$ or on a recurrence, some bucket occurring at least n_{recur} times among the last W_{recur} hidden states, each quantized per dimension with step Δ_{recur} and hashed. Before any temporal error, the alert instead compares with β_ν^{Seq} the absolute z-score of the hidden state's Euclidean norm against the previous W_z broadcasts' norms.

Other modules. Affect states use $I_{\text{Aff}}^{\text{lo}}$; drive crossings, the affective reset, and a change to a non-neutral categorical emotion use $I_{\text{Aff}}^{\text{hi}}$. Sleep events use $I_{\text{Sleep}}^{\text{lo}}$, except a failed-phase sleep summary, $I_{\text{Sleep}}^{\text{hi}}$. Utterances use $I_{\text{Lang}}^{\text{lo}}$. Action-selection intents go on a stream the cycle does not read and are never candidates.

Preprint. Not peer-reviewed.

A.3 Adaptive access rate

Let f_0 be the resting broadcast rate, I_0 the phasic salience floor, and τ_{ph} the phasic decay time. On tick k the tonic drive is

$$H_k^{\text{ton}} = \text{clip}_{[0,1]}\left(\frac{\hat{a}_k - a_0}{1 - a_0}\right),$$

with $a_0 < 1$. A categorical alert is a predictive processor's report meeting its alert criterion (A.2) or another module's event at its alert level I_m^{hi}, never a graded report below it. With $\hat{I}_k$ the largest intensity among the categorical alerts of E_k, excluding cycle telemetry and the workspace stream, the phasic input is

$$H_k^{\text{in}} = \text{clip}_{[0,1]}\left(\frac{\hat{I}_k - I_0}{1 - I_0}\right),$$

with $H_k^{\text{in}} = 0$ when E_k holds no alert. The phasic drive is a peak-hold with exponential decay,

$$H_k^{\text{ph}} = \max\!\big(H_k^{\text{in}},\ H_{k-1}^{\text{ph}}\, \mathrm{e}^{-\Delta_p/\tau_{\text{ph}}}\big), \qquad H_{-1}^{\text{ph}} = 0.$$

The access drive is $H_k = \max(H_k^{\text{ton}}, H_k^{\text{ph}}) \in [0,1]$, and the effective broadcast rate is the linear map

$$f_k^{\text{eff}} = \begin{cases} f_0 + (f_p - f_0)\, H_k & f_0 < f_p, \\ f_0 & f_0 \geq f_p, \end{cases}$$

With the controller disabled, $f_k^{\text{eff}} = f_0$. A fractional accumulator with $A_{-1} = 0$ gives the broadcast indicator:

$$A'_k = A_{k-1} + \frac{f_k^{\text{eff}}}{f_p}, \qquad \chi_k = \mathbf{1}[A'_k \geq 1], \qquad A_k = \begin{cases} \min(A'_k - 1,\ 1) & \chi_k = 1, \\ A'_k & \chi_k = 0. \end{cases}$$

Subtracting before clamping keeps the fractional carry, so at the resting rate $f_0 = f_p/3$ every third tick broadcasts.

A.4 Affect

The affect module holds valence $v \in [-1, 1]$, arousal $a \in [0, 1]$ (starting at a_0, clipped to $[0, 1]$ after each update below), dominance, and four drives.

1. *Perceptual alert.* For each vision or acoustic audition alert (A.2), read from its stream whether selected or not,

Preprint. Not peer-reviewed.

$$a \leftarrow a + \gamma_a \, h \, \min\big(\nu_{\text{cap}},\ \max(0,\ \max(\tilde{\nu}, \tilde{c}) - 1)\big)\,, \qquad h = \max\left(\text{clip}_{[0,1]}\left(\frac{r_{\text{f}}}{\max(r_{\text{s}},\ r_0)} - 1\right),\ \text{e}^{-r_{\text{s}}\tau_{\text{s}}}\right.$$

where $\tilde{\nu}$ and $\tilde{c}$ are the error and change ratios the report carries (A.2), and the habituation factor h reads the alert rates below before this alert adds to them: $\text{e}^{-r_{\text{s}}\tau_{\text{s}}}$, the chance of no alert at the expected recent count $r_{\text{s}}\tau_{\text{s}}$ (§3.5).

2. *Interoceptive alert.* For each interoception report with $\mathcal{A}_t \neq \emptyset$, $a \leftarrow a + \gamma_{\text{Intero}}$.
3. *Relaxation.* The module updates on every broadcast and every P_{Aff}, relaxing dominance toward baseline and arousal toward a_∞, raised by the largest deficit D^* (§3.5), by an explicit-Euler step,

$$a \leftarrow a + (a_\infty - a)\min(1,\ \lambda\,\Delta t_{\text{Aff}})\,, \qquad a_\infty = a_0 + \gamma_D\,(1 - a_0)\,D^*,$$

with relaxation rate λ and Δt_{Aff} the entity time since the previous update.

4. *Sleep.* The sleep module's affective reset returns valence, arousal, and dominance to baseline, resets every deficit to 0 and re-arms its crossing, sets the categorical emotion to neutral, and clears the active goals and a speaker's perceived emotion. The alert rates, error and wellness statistics, and r_{int} carry on through the sleep, since habituation recovers with time away from the stimulus and can outlast breaks (Thompson and Spencer 1966; Rankin et al. 2009; Vankov, Hervé-Minvielle and Sara 1995).

Each update, in order, decays the alert rates, computes learning progress and alert excess, relaxes arousal, dominance, and valence, relieves boredom, builds every drive, and publishes crossings and re-entries; on a broadcast the appraisal follows. It publishes its state, including a, at most every P_{Aff}. The cycle's $\hat{a}_k$ is the arousal of the last affect state event in E_k ($\hat{a}_k = \hat{a}_{k-1}$ if none), starting from the Table A1 default.

Valence. For each perceptual module m (vision, acoustic audition) the affect module keeps fast and slow time-decayed means of its raw forward-model error ν_t^m (A.2), counting from the first positive error. At the latest counted report time t_r,

$$\langle \nu^m \rangle_\tau = \frac{\sum_{i \le r} \text{e}^{-(t_r - t_i)/\tau}\, \nu_{t_i}^m}{\sum_{i \le r} \text{e}^{-(t_r - t_i)/\tau}}, \qquad \tau \in \{\tau_{\text{f}}, \tau_{\text{s}}\},$$

The error dynamics ed is the mean, over modules with $\langle \nu^m \rangle_{\tau_{\text{s}}} > 0$, of $\text{clip}_{[-1,1]}\big((\langle \nu^m \rangle_{\tau_{\text{s}}} - \langle \nu^m \rangle_{\tau_{\text{f}}})/\langle \nu^m \rangle_{\tau_{\text{s}}}\big)$, 0 before any. Learning progress lp is the largest, over modules with errors in both halves, of $\text{clip}_{[-1,1]}(1 - \bar{\nu}_{\text{n}}^m/\bar{\nu}_{\text{o}}^m)$, the newer and older half's mean errors over the latest even number of completed Δ_{lp} bins, counted likewise, at most T_{lp}, and for audition restarting at each band step of A.7 (Baranes and Oudeyer 2013; §3.5); $\text{lp}^+ = \max(0,\ \text{lp} - \text{lp}_0)/(1 - \text{lp}_0)$. The floor lp_0 is calibrated in simulation, with vision noise fitted to a test boot: a forward model with nothing to learn then ends about 0.7 boredom

Preprint. Not peer-reviewed.

crossings by learning progress in 16 awake hours. The appraisal's pleasantness check is $\mathrm{pl} = \tanh(k_v\,\mathrm{ed})$ (perceived-emotion term off), and valence relaxes toward $v_\infty = \mathrm{clip}_{[-1,1]}(\mathrm{pl} + \langle w\rangle_{\tau_\mathrm{f}} - \langle w\rangle_{\tau_\mathrm{s}})$,

$$v \leftarrow v + (v_\infty - v)\left(1 - \mathrm{e}^{-\Delta t_{\mathrm{Aff}}/\tau_v}\right),$$

with time-decayed means of the interoception module's wellness $w \in [0,1]$ (difference 0 before any): the mean headroom in $[0,1]$ over host metrics, 1 healthy (one minus utilization for processor, memory, and GPU memory; GPU temperature linear between its scaling bounds; cycle latency 1 up to its target, linear to 0 at three times it).

Drives. Each drive $D \in [0,1]$ has build rate β_D, decay rate δ_D, relief gain ρ_D, and threshold θ_D. Each update first relieves boredom by $D \leftarrow D\,\mathrm{e}^{-\rho_D \max(\mathrm{lp}^+,\,\mathrm{exc})\Delta t_{\mathrm{Aff}}}$, then, with its build signal $u \in [0,1]$ held over the step, moves every drive by the exact solution of $\dot{D} = \beta_D u(1-D) - \delta_D D$,

$$D \leftarrow D_\infty + (D - D_\infty)\,\mathrm{e}^{-(\beta_D u + \delta_D)\Delta t_{\mathrm{Aff}}}, \qquad D_\infty = \frac{\beta_D u}{\beta_D u + \delta_D}.$$

Curiosity has $u = \mathrm{lp}^+$, decay δ_I, and no relief, and survives sleep like the bins it reads. Boredom has $u = 1 - \max(\mathrm{lp}^+, \mathrm{exc})$, with the alert excess

$$\mathrm{exc} = \mathrm{clip}_{[0,1]}\left(\frac{r_\mathrm{f}}{\max(r_\mathrm{s},\ r_0)} - 1 - m_B\right),$$

where r_f and r_s are perceptual alert rates in hertz: each perceptual alert adds $1/\tau_\mathrm{f}$ to r_f and $1/\tau_\mathrm{s}$ to r_s, and each update first decays both by $\mathrm{e}^{-\Delta t_{\mathrm{Aff}}/\tau}$ with their own time constant. The A.2 criterion self-calibrates, so a still scene's base alert rate does not relieve boredom. The social drive has $u = 1$ once the sequence module reports an interaction, $u = 0$ before, with relief $D \leftarrow D(1-\rho_D)$ on each new one. Restlessness has $u = 1 - r_{\mathrm{int}}$, where at each broadcast $r_{\mathrm{int}} \leftarrow r_{\mathrm{int}} + \alpha_{\mathrm{int}}(\min(1, n_{\mathrm{int}}) - r_{\mathrm{int}})$, n_{int} counting intents other than rest since the previous broadcast, with relief $D \leftarrow D(1-\rho_D)$ on each such intent, which also relieves boredom, $D \leftarrow D(1-\rho_{\mathrm{eng}})$, unless about a drive: noticing one's failure to engage is part of boredom (Eastwood et al. 2012). A crossing of θ_D is published once and re-arms after D falls below $\Theta_D\,\theta_D$. With $\delta_D \approx \beta_D/9$, a deprived deficit settles near $D_\infty = 0.9$; δ_I holds curiosity at θ_D under sustained $\mathrm{lp}^+ = 0.045$.

Novelty. The novelty check (suddenness, predictability, familiarity; Sander, Grandjean and Scherer 2005) is $1 - \prod_e (1 - s_e)$ over coalition members with an error or change ratio (0 if none), $s_e = \mathrm{clip}_{[0,1]}(\ln \rho_e / \ln 3)$ for the larger ratio ρ_e, times h at the appraisal for vision and audition reports.

For the appraisal's goal-relevance check, the dominant drive d^* has the largest value, ties going to the alphabetically last name; v^* is its value, $\mathcal{R}(d^*)$ the sources whose events serve it, and frac_k their share of the coalition's total intensity (0 for a zero total). The drive score is

Preprint. Not peer-reviewed.

$$r_{\text{drive}} = v^* \operatorname{frac}_k,$$

with $r_{\text{drive}} = 0$ when there is no drive or $v^* \leq 0$. With an active goal ledger of token-overlap relevance $r_{\text{led}} \in [0, 1]$, the conducive part is $r = \max(r_{\text{drive}}, r_{\text{led}})$, and otherwise $r = r_{\text{drive}}$. The goal score is $\text{clip}_{[-1,1]}(r - o)$, with the obstruction o below. Relevance precedes conduciveness (Scherer 2009), so content serving or obstructing neither the drive nor a goal scores 0.

Coping potential is checked only when the goal score is nonzero, since the component process model assesses coping against a specific demand (Sander, Grandjean and Scherer 2005); otherwise it is undefined, and fear, anger, and sadness, which need it, are ineligible. The score estimates power, the entity's own ability to change the outcome, as distinct from control, whether any agent could (Sander, Grandjean and Scherer 2005). Learning agents derive power from interaction history, as overall success (Moerland, Broekens and Jonker 2018) or, in path planning, goals reached versus aborted (Williams et al. 2015); the estimate is that ratio under Laplace's rule of succession (Jaynes 2003, ch. 18):

$$p = \frac{s + 1}{s + f + 2}.$$

An outcome counts as a success only when the entity's own response produced it, since response-independent outcomes define helplessness (Maier and Seligman 2016). For the drive term, a past crossing of the dominant drive counts in s when the entity's own response ended it: its learning progress, intents, or the novelty an answer to its stir brings (A.7) for boredom and its intents for restlessness, by relieving the drive or holding its build input down on the step the drive fell back, the latter only if it would not have fallen back at the no-response build input. Every other end counts in f: decay alone, other perceptual novelty for boredom, and another's voice for the social drive. A sleep reset, or the end of a curiosity crossing, counts as neither. For the ledger term, s and f count completed and abandoned goals. The term that sets the goal score supplies p, the drive term on a tie. The coping score is $2p - 1, 0$ before any history.

Goals. A deficit's crossing of θ_D creates the goal of relieving D, with priority D at the crossing, unless one exists; intents never create goals. The goal completes when D falls below $\Theta_D \theta_D$ with the credit rule crediting the entity's own response, and is abandoned when the fall came without one, when D stays above θ_D for Δ_{stall}, or, as the lowest priority, when active goals exceed n_{goal}, the last two leaving the crossing open. While a goal is active and $D \geq \theta_D$, its drive event recurs, not as an alert, at intensity D each P_{Aff} until a broadcast accesses it, then P_{re} after that access.

Obstruction. Each active goal carries an expected-relief likelihood $L \in [0, 1]$, set to the coping ratio p at creation. When an intent other than rest made while the goal is active brings no fall of its drive within Δ_{resp}, L falls by $\Delta L = \eta_L L$; for the dominant drive's goal the next appraisal takes $o = v^* \Delta L$, and otherwise $o = 0$.

Welfare indicators. Each indicator reads the record (§6.2) over the latest T_{lp} of awake time, the span of learning progress, which holds a rest and its reset in either stage (§3.5). A deficit

Preprint. Not peer-reviewed.

is deprived when its share at or above θ_D reaches θ_{wf}: a deficit at or above θ_D sets a_∞ to 0.40 to 0.43 (A.11), the high-tonic mode in which the entity disengages from the current task to explore alternatives (Aston-Jones and Cohen 2005), and θ_{wf} marks that mode as the span's majority state. Valence is negative when its share below $-\theta_E$ reaches θ_{wf}. Boredom or restlessness is uncontrolled when, for its s and f crossing ends in the span and $n = s+f+1$, the uniform-prior posterior probability of power below one half, $2^{-n}\sum_{i=s+1}^{n}\binom{n}{i}$, reaches $q_{\text{wf}} = 1 - \alpha_{\text{test}}$; the social drive is left out, since only others' voices end it. Other record fields inform the reviewer. An indicator crosses at its bar, re-arms below it, and restarts its span after each response (§6.5). A reset boredom deficit with no relief recrosses θ_D in about 34 s.

A.5 Interoceptive prediction and regulation

The interoception module reads the host every P_{Intero} into a vector $x_t \in \mathbb{R}^{n_x}$ of normalized metrics. A frozen closed-form continuous-time reservoir advances $\mathbf{h}_t = \mathcal{F}(x_t, \mathbf{h}_{t-1}; s_t^{\text{Intero}})$ with $\mathbf{h}_0 = 0$; the timespan $s_t^{\text{Intero}} = \text{clip}_{[0,s_{\max}]}(\Delta t_{\text{Intero}}/P_{\text{Intero}})$, with Δt_{Intero} the time since the previous reading ($s_t^{\text{Intero}} = 1$ on the first), enters the cell's time gates (Hasani et al. 2022), so irregular readings integrate the elapsed time. Only the linear readout $\mathbf{L}_{\text{Intero}}$ adapts; the prediction of x_t made at reading $t-1$ is

$$\hat{x}_t = \mathbf{L}_{\text{Intero}}\,\varphi_{t-1}, \qquad \varphi_{t-1} = (\mathbf{h}_{t-1},\ b_{t-1},\ 1),$$

with b_{t-1} the context then held (A.1). With $\varepsilon_t = \hat{x}_t - x_t$, the prediction error is $\nu_t^{\text{Intero}} = \|\varepsilon_t\|_2$ (0 on the first reading). One stochastic-gradient step on $\|\varepsilon_t\|_2^2/n_x$ follows, skipped during sleep,

$$\mathbf{L}_{\text{Intero}} \leftarrow \mathbf{L}_{\text{Intero}} - \eta_{\text{fm}}\,\frac{2}{n_x}\,\varepsilon_t\,\varphi_{t-1}^{\top},$$

after which the reservoir advances on x_t and the readout predicts x_{t+1}.

Each channel i keeps an expected absolute error μ_i^U and a spread σ_i^U, both starting at 0. With the bound $\Omega_i = \mu_i^U + \beta_U\sigma_i^U$ taken before this reading's update, the unexpected error is

$$U_t = \sqrt{\sum_{i=1}^{n_x}\Big[\max\big(0,\ |\varepsilon_{t,i}| - \Omega_i\big)\Big]^2}.$$

Then, with $\alpha_t = 1 - \mathrm{e}^{-\Delta t_{\text{Intero}}/\tau_U}$, each channel updates $\sigma_i^U \leftarrow \sigma_i^U + \alpha_t\big(\big||\varepsilon_{t,i}| - \mu_i^U\big| - \sigma_i^U\big)$ and $\mu_i^U \leftarrow \mu_i^U + \alpha_t(|\varepsilon_{t,i}| - \mu_i^U)$, both with the prior μ_i^U. These updates are skipped during sleep; U_t does not enter $I(e)$.

The regulation error is $\nu_t^{\text{act}} = \nu_t^{\text{Intero}}$ when $\mathcal{A}_t \neq \emptyset$ and U_t otherwise. A stress episode runs from the first reading with $\nu_t^{\text{act}} \geq \theta_R$ to the first with $\nu_t^{\text{act}} < \theta_R$. After Δ_t^{st} of an

Preprint. Not peer-reviewed.

episode, each increase of $\lfloor \Delta_t^{\text{st}} / \Delta_R \rfloor \geq 1$ emits one advisory of tier $\min(\lfloor \Delta_t^{\text{st}} / \Delta_R \rfloor, 3)$: reduce the processing rate, shed a low-priority module, or request maintenance. A per-boot warm-up, ending once the boot has made at least n_{wu} readout updates and Δ_{wu} has passed since the module's first reading, withholds advisories unless $\mathcal{A}_t \neq \emptyset$. A rate advisory sets $f_p \leftarrow \text{clip}_{[f_p^{\min},\, f_p^{\max}]}(\gamma_f f_p)$.

Housekeeping. Regenerable or spent data is what no module reads again and the entity's continuation does not need: process-log backups, caches rebuilt on demand, and local study-record copies whose export to the archive volume a checksum has verified; unexported study records are never deleted. The bus trims each stream to a fixed length on write, so in the first configuration host memory's S_j^{hk} is 0. In every sleep, last before the affective reset, housekeeping (§3.5) reclaims space, oldest data first, until on each resource j (disk, host memory) the bytes S_j^{hk} of such data and unexported study records meet $S_j^{\text{hk}} \leq B_j^{\text{hk}}$, or until Δ_{hk} has passed. The pressure is $P_t^{\text{hk}} = \sum_j \max(0,\ S_j^{\text{hk}}/B_j^{\text{hk}} - 1) + \max(0,\ S/B - 1)$, whose last term is the episodic stores' excess over their budget (A.10) and is 0 without them. Awake, fatigue builds at the rate $\nu_t^{\text{act}}\,(1 + P_t^{\text{hk}}) - \delta_{\text{fat}}$, floored at 0, and triggers sleep at θ_{fat}.

Awake floor. The software periodically checks each data volume's free space against a reserve R^{hk}. Below it, the floor reclaims regenerable data, oldest first, until free space reaches $\min(2R^{\text{hk}},\ (V_{\text{vol}} + R^{\text{hk}})/2)$, V_{vol} being the volume's size. If free space is still below R^{hk}, the floor suspends the study record, writing a gap marker, until free space is back at that target. The disk's excess at that moment counts in P_t^{hk} until the next sleep.

A.6 Report rule

The action-selection module derives intents only from accessed broadcasts ($\iota_k = 0$), by these first-configuration steps, timed by t_{now}.

1. *Guard release.* The speak guard is released when the language organ publishes external speech, the think guard on internal speech, accessed or not, and either after timeout Δ_g, so a failed realization cannot silence the entity.
2. *Report signal.* Let $\mathcal{C}_k^{\circ}$ be the coalition minus the language organ's events; an empty one yields no intent. Otherwise $e_k^{\dagger}$ is its first (highest-scoring) member, $Q_k = S'_k(e_k^{\dagger})$, and its signature $\text{sig}_k = (\text{src}(e_k^{\dagger}), \text{typ}(e_k^{\dagger}))$ pairs its source and event type.
3. *Signature check.* A spoken report is blocked when $\text{blk}_k = \mathbf{1}\big[\text{sig}_k = \text{sig}^{\text{last}} \wedge t_{\text{now}} - t^{\text{sig}} < \Delta_{\text{sig}}\big] = 1$, with sig^{last} and t^{sig} the last spoken report's signature and time.
4. *Speak.* A speak intent is emitted when

$$Q_k \geq \theta_{\text{sp}}, \qquad \text{the speak guard is released}, \qquad t_{\text{now}} - t^{\text{sp}} \geq \Delta_{\text{sp}}, \qquad \text{blk}_k = 0.$$

Emitting it arms the speak guard and sets $t^{\text{sp}} \leftarrow t_{\text{now}}$, $\text{sig}^{\text{last}} \leftarrow \text{sig}_k$, and $t^{\text{sig}} \leftarrow t_{\text{now}}$.

5. *Think.* Otherwise a think intent is emitted when $Q_k \geq \theta_{\text{th}}$, the think guard is released, and $t_{\text{now}} - t^{\text{th}} \geq \Delta_{\text{th}}$, arming the think guard and setting $t^{\text{th}} \leftarrow t_{\text{now}}$, with no signature

Preprint. Not peer-reviewed.

check and no change to sig^{last}.

In the developmental phase, a broadcast whose $e_k^{\dagger}$ is the entity's own boredom or restlessness drive event (A.4) instead yields a stir intent when $t_{\text{now}} - t^{\text{st}} \geq \Delta_{\text{st}}$, setting $t^{\text{st}} \leftarrow t_{\text{now}}$; a deprived deficit's event scores about 0.49 at the arousal it sets, above θ. At most one intent is emitted per broadcast; t^{sp}, t^{th}, and t^{st} start at $-\infty$. The rule requires $0 \leq \theta_{\text{th}} \leq \theta_{\text{sp}} \leq 1$, both bars above θ and on the same score S'_k; an optional interrupt bar above θ_{sp} is unset here.

A.7 Heartbeat entrainment marker

Awake time $\mathcal{T}_{\text{awake}}$ (entity time awake with the cycle unfrozen, summed over boots) is one clock for all rules; it and all marker, probe, and module state persist with the entity, checkpointed each entity minute and at shutdown, so a restart neither resets nor repeats them.

Self-rhythm model. The interoception module's self-rhythm is a mean-field population with excitatory recurrence, synaptic depression, and adaptive recovery, with activity y, synaptic resource $R \in [0, 1]$, and log recovery time $\ell = \ln(\tau_{\text{rec}}/1\,\text{s})$, initialized from Table A1. Each step of $\Delta_r = 1/f_r$ runs n_s equal semi-implicit Euler substeps, with y updated first and R from the new y:

$$\Xi = J_{\text{rec}}\, R\, y + \Xi_0 + \gamma_o\, o + \gamma_u\, M + \xi, \qquad y_\infty = \frac{1}{1 + \mathrm{e}^{-\Xi/k_y}},$$

$$\tau_y\, \dot{y} = -y + y_\infty, \qquad \dot{R} = \frac{1 - R}{\tau_{\text{rec}}} - \gamma_R\, R\, y,$$

with R clipped to $[0, 1]$ after each substep and τ_{rec} fixed within a step. The own drive $o \in [0, 1]$ is the module's last report intensity, $M \in [0, M_{\text{max}}]$ is the maternal drive (below), and the Gaussian noise ξ, redrawn each substep, has standard deviation $\sigma_\xi \sqrt{(10^{-3}\,\text{s})\, n_s/\Delta_r}$. After each step the averages $\bar{M}$, $\bar{y}$, $\bar{R}$ update with weight $\alpha_{\text{r}} = \min(1, \Delta_r/\tau_m)$, and the recovery time adapts by a phase-projected rule,

$$\ell \leftarrow \text{clip}_{[\ln \tau_{\text{rec}}^{\text{min}},\, \ln \tau_{\text{rec}}^{\text{max}}]}\big(\ell - \eta_\ell\, (M - \bar{M})\, \zeta\, \Delta_r\big), \qquad \zeta = \frac{R - \bar{R}}{\sqrt{(y - \bar{y})^2 + (R - \bar{R})^2}},$$

with $\zeta = 0$ when the root is at most 10^{-9} and bounds in seconds; the rule never uses the maternal rate. amp is the population standard deviation of the activity over the last Δ_{amp}, mean-removed and causally band-passed to $[f_{\text{lo}}, f_{\text{hi}}]$ (0 for the first four seconds).

Maternal drive and probes. The maternal beat has phase $\psi \in [0, 2\pi)$, a nominal rate with slow drift, and a "lub-dub" envelope $\mathcal{E}(\psi) \in [0, 1]$. Each step sets $M = M_{\text{max}}\, c_M\, \bar{\mathcal{E}}$, with

Preprint. Not peer-reviewed.

$\bar{\mathcal{E}}$ the mean envelope since the last step and c_M equal to c_M^{base}, 0 in a withdrawal (length Δ_W), or c_M^{pert} in a perturbation (length Δ_{pert}). A withdrawal or perturbation starting at t' schedules the next of its kind at $t' + P_W(1 + j_P\Upsilon)$ or $t' + P_{\text{pert}}(1 + j_P\Upsilon)$, Υ uniform on $[-1, 1)$ from the run seed; the first of each follows the post-boot settling period, offset likewise. A due probe starts at the first readout step with no probe running, the entity awake and unfrozen, at least P_{ro} (Δ_E for a withdrawal) since boot or thaw, and at least $\max(60\text{ s}, \Delta_W)$ since the last probe; the earlier-scheduled kind goes first.

Phase-locking value. At rate f_g the readout samples y, ψ, and the phases $\psi^{(1)}, \dots, \psi^{(19)}$ of 19 surrogate heartbeats (the beat generator under 19 other seeds, the surrogate method of Van Leeuwen et al. 2003). Withdrawal j's driven window, the contiguous run of undisturbed samples of at most Δ_E ending at its start, uses only samples with all 21 values and needs at least $0.8\,\Delta_E f_g$ of them. The activity is mean-centered, band-passed to $[f_{\text{lo}}, f_{\text{hi}}]$ by a zero-phase forward-backward Butterworth band-pass of design order two (fourth order as a filter), and Hilbert-transformed to the phase ϕ_q; $\Delta_{\text{trim}} f_g$ samples are trimmed from each end of every series. Over the n_ϕ remaining samples,

$$\text{PLV} = \left| \frac{1}{n_\phi} \sum_{q=1}^{n_\phi} \mathrm{e}^{\,\mathrm{i}(\phi_q - \psi_q)} \right|, \qquad \text{PLV}^{(l)} = \left| \frac{1}{n_\phi} \sum_{q=1}^{n_\phi} \mathrm{e}^{\,\mathrm{i}(\phi_q - \psi_q^{(l)})} \right|,$$

and the locking condition is PLV $>$ $\max_{1 \le l \le 19}$ PLV$^{(l)}$, strictly, so ties fail. Under the null that the own-heartbeat PLV is exchangeable with the 19 surrogate values, this is a rank test with attained one-sided $p = 1/20 = 0.05$.

Frequency pull. f_w is the least-squares slope, divided by 2π, of the activity's unwrapped Hilbert phase during the withdrawal (processed as above, untrimmed; needs at least $0.6\,\Delta_W f_g$ samples), f_b that slope for the unwrapped beat phase over the driven window's used samples, and f_{w0} the mean f_w over the entity's first n_{base} withdrawals with defined f_w. From the n_{base}-th such withdrawal on, inclusive,

$$\text{pull} = 1 - \frac{|f_w - f_b|}{|f_{w0} - f_b|},$$

undefined when f_w or f_b is undefined or $|f_{w0} - f_b| < f_{\text{guard}}$. The pull condition is pull $\ge$ pull$_{\text{min}}$. Below about 60 bpm the pull is undefined, so the maternal model uses 60 to 80 bpm.

Self-sustain. With $\overline{\text{amp}}^{\,\text{dr}}$ the mean amplitude over undisturbed samples in the Δ_W before the withdrawal and $\overline{\text{amp}}^{\,\text{wd}}$ the mean during it, the rhythm self-sustains when $\overline{\text{amp}}^{\,\text{wd}} \ge \frac{1}{2}\,\overline{\text{amp}}^{\,\text{dr}}$ and $\overline{\text{amp}}^{\,\text{wd}} > 0$.

Marker. Withdrawal j passes (pass$_j$ $= 1$) when locking, pull, and self-sustain all hold, and is undefined when any is. The count is $n_j^{\text{pass}} = n_{j-1}^{\text{pass}} + 1$ if pass$_j = 1$ and $n_j^{\text{pass}} = 0$ otherwise, undefined passes included. The marker after withdrawal j is undefined when

Preprint. Not peer-reviewed.

pass$_j$ is and otherwise $\mathbf{1}[n_j^{\text{pass}} \geq 3]$. Three consecutive passes guard against a chance pass at $p = 0.05$; consecutive withdrawals are correlated, so the joint error rate is not 0.05^3.

Viability. After each withdrawal, $\mathcal{P}$ holds the defined pull values of the last Δ_V awake hours; when $|\mathcal{P}| \geq n_V$, $\widetilde{\text{pull}}$ is their median and $\widehat{\text{slope}}$ their least-squares slope per awake hour. The phase is declared unviable at the first rule that fires:

- R0: $\mathcal{T}_{\text{awake}} \geq t_{R0}$ and no defined pull has been recorded;
- R1: $\mathcal{T}_{\text{awake}} \geq t_{R1}$, the marker has never been true, $|\mathcal{P}| \geq n_V$, $\widetilde{\text{pull}} < \text{pull}_{R1}$, and $\widehat{\text{slope}} \leq \text{slope}_{R1}$;
- R2: $\mathcal{T}_{\text{awake}} \geq t_{R2}$, the marker has never been true, $|\mathcal{P}| \geq n_V$, and $\widetilde{\text{pull}} < \text{pull}_{R2}$;
- R3: $\mathcal{T}_{\text{awake}} \geq t_{R3}$ and the marker has never been true.

The verdict is recorded once; the study runner ends the step.

Maturation gate and budget. The gate, evaluated each gate cadence, has three conditions, each failing closed on missing or stale evidence. C1 requires a readout at most three cadences old in which the latest withdrawal's self-sustain and entrainment markers are true, $V_{\text{per}} \geq V_{\text{per}}^{\min}$, $\text{err}_{\text{Vis}} \leq \text{err}_{\text{Vis}}^{\max}$, and the return-to-baseline time is at most $\Delta_{\text{max}}^{\text{rec}}$. The period variability V_{per} is the coefficient of variation of intervals between self-rhythm phase wraps over the last Δ_{per}, once four wraps exist; $V_{\text{per}}^{\min}$ checks only for rigidity and lies below fetal five-minute heart-rate values (Moraes et al. 2012), since an entrained rhythm takes on its driver's regularity. The vision error ratio err_{Vis} is the median vision prediction error on the developmental stimulus (§7) since the last readout, divided by the first such median of at least three errors. The return-to-baseline time runs from the last perturbation's end to the first sample where the median of ν^{Intero} over the preceding five seconds (not before the perturbation's start) is at most $(1 + \text{tol}_{\text{rec}})$ times its median over the minute before the perturbation, and is $\Delta_{\text{cap}}^{\text{rec}}$, a C1 failure, if none comes within $\Delta_{\text{cap}}^{\text{rec}}$. C2 requires at least n_{sleep} completed sleeps with the sleep module present, and at least n_{cons} consolidation passes with the world-model module also present. C3 requires the sensory schedule to be complete (A.7). The study runner also ends a developmental step once its wall-clock duration exceeds Δ_{budget}, the 96-hour validation horizon, which exceeds $t_{R3}/(1 - s_q) = 88.2$ entity hours, s_q the rest share.

Maternal voice and maturation. A pre-rendered maternal voice, heard above the intrauterine background with intonation kept (Querleu et al. 1988; Richards et al. 1992), plays L_v above the soundscape's RMS ($\pm j_v$). A Δ_v slot talks with probability p_v times maternal arousal (about 7 percent speech, near a woman's daily speaking time; Mehl et al. 2007), opening with a repeated passage (DeCasper and Spence 1986) with probability r_v; five of 20 items start unlocked, one more per Δ_u awake. Bed and voice (low-passed at 500 and 3000 Hz) pass a band that widens in five stages, as fetal hearing does: from 500 Hz alone, adding 250, 100, 1000, then 3000 Hz (Hepper and Shahidullah 1994), edges geometrically midway between tones heard and not. Stage k ends when audition's learning progress, restarted at the step, stays below lp_0 for W_{sat}, or at $d_{\max}$, as staged development is released by ability rather than age (Law et al. 2014). A missing reading counts as unsaturated, so a stage lasts at least the two Δ_{lp} bins of a first reading plus W_{sat}, ten bins, which no single noisy bin can satisfy. Color saturation is $k/5$ in stage k, a design choice without a source (newborn color

Preprint. Not peer-reviewed.

vision is limited; Adams, Courage, and Mercer 1994). Offline a band took 5.1 to 5.7 hours to saturate (one seed), yet in the test boot of §6.5 audition's learning progress never rose after a step and every stage would have saturated 12 to 17 awake minutes in, so the measure misses the band change or the band adds little; release by ability then opens the full band within about an hour.

Added sounds and stirs. A melody within the opening band fills a Δ_v slot with probability p_m, and an outside voice at the soundscape's level one with p_x, its fundamental heard from the 100 Hz step; the upper steps add both voices' detail. A stir is answered, with probability p_{st} and after a delay uniform on $[\Delta_{\text{st}}^{\text{lo}}, \Delta_{\text{st}}^{\text{hi}}]$, by maternal speech. A stir relieves restlessness at once, and its answer relieves boredom through the novelty and learning it brings, credited to the entity (A.4), so a stirring entity's boredom reads as controlled and its deprivation indicator should cross less often; one still crossing after rest and stirs marks a stimulus with nothing left to learn, for the notified person to change. These expected effects are to be checked offline.

Offline validation. The validation (`docs/records/validation/self-rhythm-rests/` in the reference repository) runs the archived simulation harness with the developmental rest cycle as a schedule only: generator, drive, and plasticity run on through rests, and the cycle sets awake time and where withdrawal probes fall. Over 96 simulated hours all 10 seeds entrain at 70 bpm at 11.1 to 16.2 entity hours (median 13.45), and each rate from 60 to 80 bpm tracks its own heartbeat: 60 bpm locks at 13.9 to 32.3 hours (3 of 3 runs) and 80 bpm at 40.5 to 45.6. Without rests, 60 bpm, nearest the rhythm's natural frequency, locks at about 3 hours, before that record's start criterion (no replicated pass before 5.5 hours) allows (`openspec/changes/archive/2026-10-03-self-rhythm-earned-entrainment/validation.md`). No control passes: no maternal drive, a jittered beat, no plasticity, and six pairs scored against a foreign heartbeat, whose single withdrawals passed by chance in 3.1 to 8.6 percent, never three in a row. The three-pass requirement, fixed on the first run, held on fresh seeds. R0 to R3 flagged no viable run and the unviable controls by 24 to 27 awake hours, the jittered beat at about 29.6.

Preprint. Not peer-reviewed.

A.8 Broadcast information gain and the planned test

Gain. Let $\mathcal{J} = \{\text{Vis}, \text{Aud}, \text{Intero}\}$, the sequence module excluded (§6.3). For $j \in \mathcal{J}$ and report t, b is the context j held when forming the prediction scored at t (A.2, A.5), and $\nu_t^j(b')$ that prediction's error with b' for b and all weights unchanged, so $\nu_t^j(b) = \nu_t^j$ (for audition, the acoustic path's error). The kept sources of j are j and, for interoception, the language organ, whose host load would otherwise let its context share predict the interoceptive input. In the null context b_j^0, each component of b summing per-event contributions (A.1) becomes the kept sources' share plus the others' share averaged over the contexts the processor has adopted so far in the run; coalition-level components 1 to 4 become their means over those contexts; the age and two flags are kept. The report's broadcast information gain is

$$\text{gain}_{j,t} = \frac{\nu_t^j(b_j^0) - \nu_t^j(b)}{\bar{\nu}_t^j},$$

with $\bar{\nu}_t^j$ the running mean of j's errors over W_P (A.2). A report is included when its error is

defined, $\bar{\nu}_t^j > 0$, a context was adopted before the prediction, the entity was awake, and it falls after the learning period T_{learn} (entity time from run start). IG_j is the mean of $\text{gain}_{j,t}$ over j's included reports in one run of an arm, undefined if there is none, and the headline is

$$\text{IG} = \frac{1}{3} \sum_{j \in \mathcal{J}} \text{IG}_j,$$

undefined if any IG_j is.

Positive control. For each $j \in \mathcal{J}$, a synthetic component informative about j's next input x_{t+1}^j joins j's context as another source's contribution. It is $v_j^\top (x_{t+1}^j - \mu_j) + \epsilon_t$, with v_j the first principal component of j's inputs (Jolliffe 2002) and μ_j their mean, both from a preliminary on-arm run on a seed not used again, and ϵ_t independent Gaussian noise with the projection's variance divided by SNR_j. A second component of equal mean and variance, independent of the input, runs beside it. A component's effect is IG_j with it minus without it, over included reports of complete film-program runs. Following Currie's (1968) detection limits, with target $\theta_{\text{PC}} = \min(\theta_{\text{eff}}^{\text{match}}, \theta_{\text{eff}}^{\text{pool}})$, SNR_j sets the informative component's expected effect to θ_{PC}, and j passes when, over at least n_{PC} included reports, that component's effect reaches $\theta_{\text{PC}}/2$ and the other's does not; n_{PC} gives detection probabilities of 0.95 and at most 0.05. The live test runs only after every processor passes.

Arms. The arms run the first configuration (§6.3), run r of each on the same seeds, so runs are paired.

- *Workspace on.* Selection and access as in A.1, with the context $b = \Phi_I(\mathcal{C}_{k'}^{\text{acc}}, \Delta^b, 0)$.
- *Matched selection.* Each broadcast tick's coalition is $\min(K, |E_k|)$ candidates drawn uniformly without replacement from E_k, seeded from the run seed; the rule of A.1 then gives the accessed set, whose Φ_I is the context.
- *Pooled.* No candidate is scored, selected, or gated. Every broadcast has $\iota_k = 0$ and replaces the context by $b = \Phi_I(E_{k'}, \Delta^b, 0)$. The action-selection module applies the rule of A.6 with intensities for scores, and the access rate (A.3) runs unchanged on alerts and arousal.

Run contrasts. Against control $\text{ctl} \in \{\text{match}, \text{pool}\}$, run r has contrast $\text{dIG}_r^{\text{ctl}} = \text{IG}_r^{\text{on}} - \text{IG}_r^{\text{ctl}}$ and minimum effect $\theta_{\text{eff}}^{\text{ctl}} = c_{\text{eff}}\, s_{\text{cal}}^{\text{ctl}}$, with $s_{\text{cal}}^{\text{ctl}}$ the sample standard deviation of dIG^{ctl} over the n_{cal} paired pilot runs. A run is NOT EXERCISED when competition rarely operates (in the on arm, the share of broadcasts with $|E_k| > K$ is below ϕ_{min}), when the on arm has no accessed broadcast, or when $\text{dIG}_r^{\text{ctl}}$ is undefined. Otherwise it is NEGATIVE if $\text{dIG}_r^{\text{ctl}} \le -\theta_{\text{eff}}^{\text{ctl}}$, WIN if $\text{dIG}_r^{\text{ctl}} \ge \theta_{\text{eff}}^{\text{ctl}}$, and NULL otherwise.

Decision rule. For each control, the runs classified WIN, NULL, or NEGATIVE enter two one-sided exact sign tests, which assume only that runs are independent: the WIN test on $u_r^+ = \text{dIG}_r^{\text{ctl}} - \theta_{\text{eff}}^{\text{ctl}}$ and the NEGATIVE test on $u_r^- = -\theta_{\text{eff}}^{\text{ctl}} - \text{dIG}_r^{\text{ctl}}$. Dropping exact zeros, with $n_\pm$ values left of which $n_\pm^{\text{pos}}$ are positive,

Preprint. Not peer-reviewed.

$$\mathrm{pv}^{\pm} = 2^{-n_{\pm}} \sum_{i=n_{\pm}^{\mathrm{pos}}}^{n_{\pm}} \binom{n_{\pm}}{i},$$

with $\mathrm{pv}^{\pm} = 1$ when $n_{\pm} = 0$. The verdict is WIN when $\mathrm{pv}^{+} \leq \alpha_{\mathrm{test}}$, NEGATIVE when $\mathrm{pv}^{-} \leq \alpha_{\mathrm{test}}$, NOT EXERCISED when every run is, and NULL otherwise. The hypothesis holds for this configuration on a WIN against both controls, each tested at α_{test} (§6.3). With k_n the smallest k for which $2^{-n} \sum_{i=k}^{n} \binom{n}{i} \leq \alpha_{\mathrm{test}}$, the power of the WIN test at n runs, when each $\mathrm{dIG}_r^{\mathrm{ctl}}$ exceeds $\theta_{\mathrm{eff}}^{\mathrm{ctl}}$ with probability q_{pow}, is

$$\pi(n) = \sum_{i=k_n}^{n} \binom{n}{i} q_{\mathrm{pow}}^{i} \, (1 - q_{\mathrm{pow}})^{n-i}.$$

As the test is discrete, $\pi(n)$ is not monotone in n.

Contrast at threshold. In the on arm access is exactly $S_k^* \geq \theta$ (A.1), a sharp regression discontinuity (Imbens and Lemieux 2008) with running variable S_k^* and cutoff θ; a broadcast is near threshold when $|S_k^* - \theta| \leq \theta_{\mathrm{band}}$. The reports following $\mathcal{B}_k$ are the included reports of each $j \in \mathcal{J}$ predicted after receiving $\mathcal{B}_k$ and before the next broadcast. Reports are grouped by their context age Δ^b (§6.3) in bins fixed with the other settings, and $\bar{g}_{k,a}$ is the mean gain of the reports following $\mathcal{B}_k$ in bin a. In run r and bin a, $\hat{\tau}_{r,a}$ is the jump of $\bar{g}_{k,a}$ at $S_k^* = \theta$ from a local-linear fit per side with a triangular kernel of half-width θ_{band}, bias-corrected by a local-quadratic fit of half-width b_{band} (Calonico, Cattaneo, and Titiunik 2014). A side is fitted only if it holds at least n_{side} distinct values of S_k^* within both bandwidths of the threshold, and each fit's weighted least-squares problem, with the score scaled by its bandwidth, has condition number at most κ_{max}. A local-polynomial fit behaves as if it had one observation per distinct value of the running variable (Cattaneo, Idrobo, and Titiunik 2024), so a sparse bin would otherwise weigh as much as a full one; $n_{\mathrm{side}} = 10$ is their smallest recommended per-side number, for the narrowest local-randomization window. The run's contrast is the mean $\hat{\tau}_{r,a}$ over bins with both sides fitted, undefined if none, and across runs enters a sign test of the same form on the contrast minus its minimum effect. The robust confidence interval of Calonico, Cattaneo, and Titiunik (2014), with standard errors clustered by run, is reported for each bin's pooled estimate.

Speech sound and tone. A variant of the on arm replaces the tone rule of A.2 by $I(e) = \Gamma_{\mathrm{Aud}}(\tilde{\nu}_t^{\mathrm{utt}})$ for every tone event, its run r seeded as run r of the arms. The difference of run r, the accessed fraction of emotional tone events minus that of neutral ones, both after T_{learn} (undefined unless both occur), enters a sign test of the same form. The adjusted difference averages the within-bin difference over bins of $\tilde{\nu}_t^{\mathrm{utt}}$ holding both kinds (undefined if none) and enters a second such test; bins are not tested separately.

Settings from the pilot runs. $n_{\mathrm{cal}} = 30$ paired pilot runs of the three arms and the tone variant, on seeds not used again, fix every remaining setting by the rules below; thirty is a common pilot size for parameter estimation (Lancaster, Dodd, and Williamson 2004). NOT EXERCISED pilot runs are replaced until thirty count, r_{NE} is the fraction of attempts NOT EXERCISED, and no live run enters a setting.

Preprint. Not peer-reviewed.

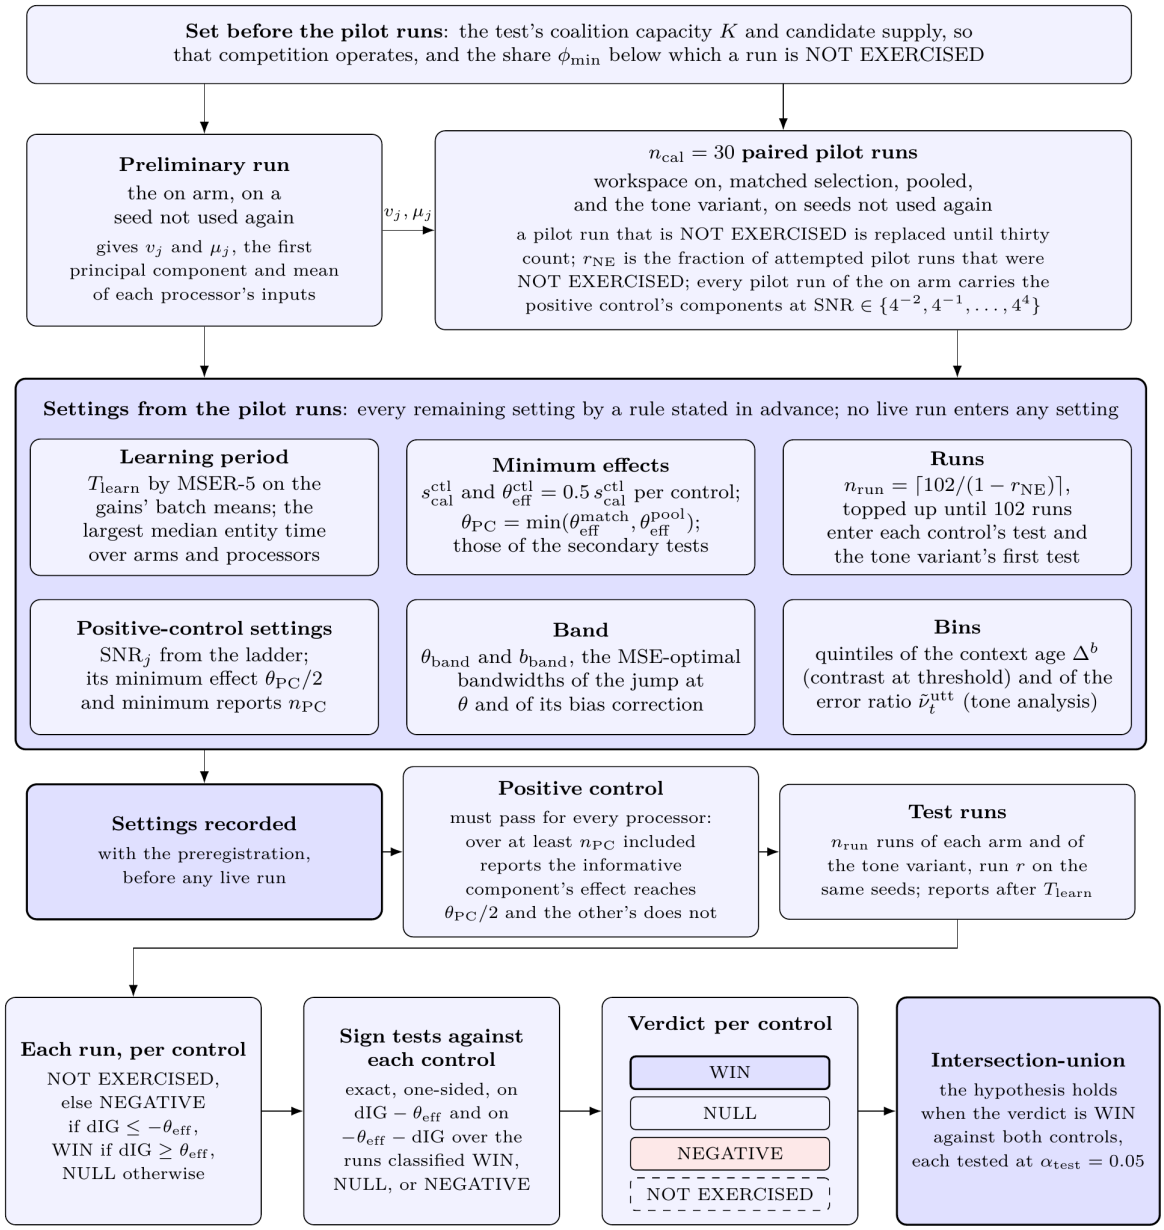


Figure 8: The test's design, from pre-pilot settings to the verdict.

Preprint. Not peer-reviewed.

Learning period. For each arm and $j \in \mathcal{J}$, the gains of j's reports are averaged across pilot runs at each report number and grouped in consecutive batches of five, with means $\bar{Y}_1, \dots, \bar{Y}_b$. The truncation point follows the marginal standard error rule (White 1997), the most effective automatic warm-up rule in the comparison of Hoad, Robinson, and Davies (2010),

$$d^* = \arg\min_{0 \le d \le b-5} \frac{1}{(b-d)^2} \sum_{i=d+1}^{b} \left(\bar{Y}_i - \bar{Y}_{(d)}\right)^2,$$

with $\bar{Y}_{(d)}$ the mean of $\bar{Y}_{d+1}, \dots, \bar{Y}_b$. As there, at least five batch means remain after d, since a few close final means give a spuriously small statistic. T_{learn} is the largest, over arms and processors, of the median entity time of report $5d^*$ in the pilot runs, rounded up to a whole second, and applies to every test run and the positive control. The pilot runs are too short if $d^* > b/2$ or a series has under ten batch means; the test then waits for longer ones.

Minimum effects. The contrast at threshold and the two tone tests use c_{eff} times the standard deviation of their run-level statistic over the pilot runs where it is defined.

Runs. $q_{\text{pow}} = 0.65$, Cohen's (1988) medium effect for a sign test on $\text{dIG}_r^{\text{ctl}} - \theta_{\text{eff}}^{\text{ctl}}$ (under normality, a mean contrast of about 0.9 of its standard deviation), is a design value, as pilot effect estimates are too imprecise to set power (Kraemer et al. 2006). n_{exact}, the smallest n with $\pi(n') \ge 0.9$ for every $n' \ge n$ (Chernick and Liu 2002), is 102, with $\pi(102) = 0.920$. $n_{\text{run}} = \lceil n_{\text{exact}}/(1 - r_{\text{NE}}) \rceil$ runs are planned, with more added until at least n_{exact} enter each control's test; entry ignores the contrast, so the level is unchanged. The tone analysis's first test has the same power function, so the tone variant takes the same n_{run}, extended until at least n_{exact} runs give a defined difference.

Positive control. Every on-arm pilot run carries both components of each processor at $\text{SNR} \in \{4^{-2}, 4^{-1}, \dots, 4^4\}$; SNR_j is the smallest value at which the informative component's mean pilot effect reaches θ_{PC}, interpolated linearly in log SNR from the value below. If none does, the live test does not run. With s_j^2 the larger of the two components' pilot variances of the run-level effect at SNR_j and $\bar{n}_j$ the mean included reports of j per pilot run,

$$n_{\text{PC}} = \max_{j \in \mathcal{J}} \left\lceil \frac{(2 z_{0.95})^2 \, \bar{n}_j \, s_j^2}{\theta_{\text{PC}}^2} \cdot \frac{n_{\text{cal}} - 1}{\chi^2_{0.2,\, n_{\text{cal}}-1}} \right\rceil,$$

where $z_{0.95}$ is the standard normal quantile, so that $2z_{0.95}$ standard errors, the detection limit of Currie (1968) at $\alpha = \beta = 0.05$, equal θ_{PC}. The last factor is Browne's (1995) recommended 80 percent upper confidence limit on a pilot variance.

Band and bins. θ_{band} and b_{band} minimize the asymptotic mean squared error of the local-linear jump estimate at θ and of its local-quadratic bias correction (Calonico, Cattaneo, and Titiunik 2014), on the pooled on-arm pilot broadcasts, one observation per broadcast with a following included report (their mean gain). A run's jump is a point estimate entering a sign test, so the mean-squared-error criterion applies; the coverage-optimal bandwidth is smaller, and the robust interval stays valid at the larger one (Calonico, Cattaneo, and Farrell

Preprint. Not peer-reviewed.

2020). A bias shared by all runs survives aggregation, while per-run noise costs only power, so both bandwidths are set on the pooled sample, narrowing θ_{band} to about $30^{-1/5} \approx 0.5$ times a single run's optimum. The age bins are the quintiles of Δ^b over reports following near-threshold broadcasts in the on-arm pilot runs, and the tone bins those of $\tilde{\nu}_t^{\text{utt}}$ over all tone events after T_{learn}, emotional and neutral, in the tone-variant pilot runs on the counted seeds. The plain mean over quintile bins weights both sides alike and adjusts directly to the pooled pilot distribution: five equal subclasses remove about 90 percent of the bias from a continuous covariate (Cochran 1968).

A.9 Multi-seed stability

For per-seed headline values $Y_1, \dots, Y_{n_{\text{seed}}}$ with mean $\bar{Y}$ and population standard deviation σ_Y, the coefficient of variation is

$$\text{CV} = \begin{cases} \sigma_Y / |\bar{Y}| & \bar{Y} \neq 0, \\ 0 & \bar{Y} = 0,\ \sigma_Y = 0, \\ \infty & \bar{Y} = 0,\ \sigma_Y > 0. \end{cases}$$

An ensemble is stable when $\text{CV} \leq \text{tol}_{\text{CV}}$ and at most one distinct verdict occurs; one without verdicts counts as unanimous. The offline suite checks the coherence-layer ablation over n_{seed} seeds.

A.10 Memory strength and sleep

These rules of §4 apply when the held episodic-memory module is present. It stores one trace per broadcast, each persisted trace i with a strength $h_i \in (0, 1]$, set to 1 at storage and changed only by downscaling, and a storage time on the entity clock. $\mathcal{W}$ is the set of traces stored since the last downscaling, and $N_{\text{new}} = |\mathcal{W}|$.

Replay. Each sleep first replays the set $\mathcal{Q}$ of the n_{rep} traces of $\mathcal{W}$ (all, if fewer) with the highest score $c_{\text{aff}}\, a_i^{\text{enc}} + c_{\text{rec}}\, \text{rec}_i$, where a_i^{enc} is the affect module's last published arousal (A.4) before trace i was stored (0 if none) and rec_i maps storage time linearly onto $[0, 1]$ across $\mathcal{W}$, newest 1 and oldest 0 ($\text{rec}_i = 1$ if all share one time).

Downscaling. Traces outside $\mathcal{Q}$ are then scaled by one factor,

$$h_i \leftarrow \varrho^{\,N_{\text{new}}/N_{\text{day}}}\, h_i, \qquad i \notin \mathcal{Q},$$

and those in $\mathcal{Q}$ keep theirs. $N_{\text{day}} = 86\,400\, f_0$ traces are stored per entity day at the resting broadcast rate. At the resting rate $\sum_i h_i$ approaches $N_{\text{day}} / \ln(1/\varrho)$, about five days of storage, after which each sleep returns it to about its pre-wake level.

Associative replay. After downscaling, each sleep co-replays at most n_{pair} pairs, most similar first, of a trace in $\mathcal{Q}$ and one stored before $\mathcal{W}$ with $\text{sim}_{\text{min}} \leq \text{sim} < \text{sim}_{\text{dup}}$ (sim as in recall below). A pair's link has weight $w = 1/2$, a pair co-replaying once since $\mathcal{Q}$ holds new traces; downscaling scales w as it scales h_i, and recall of either member also returns the other. A

Preprint. Not peer-reviewed.

trace with at least n_{sch} links forms a schema node, which recall matches like a trace: the centroid of its own and its partners' embeddings, with their mean link weight as strength. Links and nodes count in S and are pruned as traces are, a deleted trace taking its links, and each trace keeps at most n_{link} links, the weakest going first.

Pruning. The module then deletes every trace with $h_i < h_{\text{min}}$ and, while the stores' size S exceeds the budget B, the weakest trace outside $\mathcal{Q}$, oldest first among equals, until $S \leq B$ or none remains outside $\mathcal{Q}$. S counts the persisted records' bytes. Deletion happens only here, so between sleeps S can exceed B by the records stored since the last. An unreplayed trace reaches the floor after about $\ln h_{\text{min}} / \ln \varrho$, about 23, entity days at the resting rate, so with a non-binding budget the stores settle near $N_{\text{day}} \ln h_{\text{min}} / \ln \varrho$, about 6.7 million traces. A preservation copies stores and strengths as they stand, so later pruning deletes nothing from a preserved entity.

Recall. Recall returns the requested number of traces with $\text{sim}_i > 0$ in decreasing order of $h_i \, \text{sim}_i$, where sim_i is the cosine similarity of the cue and trace embeddings. A short-term-buffer trace has strength 1.

A.11 Parameter values

Table A1. Parameter values. Status: blank (design) is part of the model's definition; "calibrated" is set offline by simulation or by the validation of A.7; "provisional" is a current setting, calibrated before the live runs and recorded with them; "operator" is chosen per installation. Times are on the entity clock unless marked as awake time ("awake", A.7) or wall-clock ("wall").

Symbol	Meaning	Value	Units	Status
f_p	processing rate	10	Hz	
f_0	resting broadcast rate	10/3	Hz	
	entity seconds per wall-clock second	1		
B	maximum events read per stream per tick	100	events	
K	coalition capacity	5	events	provisional
θ	access threshold	0.35		provisional
W_N	novelty window	32	candidates	
γ_I	interest gain	0.5		calibrated
γ_G	drive-relevance attenuation (off)	0.5		
T_{min}, T_{max}	arousal level-gain floor, ceiling	0.2, 1.0		
g_{max}	contrast gain at arousal 1	8.0		
κ_{min}, κ_{max}	coherence factor floor, ceiling (layer on)	0.8, 1.25		
W_κ	coherence phase window (layer on)	10	ticks	
π_{min}, π_{max}	workspace-feedback gain floor, ceiling (feedback off here)	0.8, 1.25		
d_{conf}, r_{conf}, $\Delta\pi$, τ_π, Δ_{ref}, η_{td}	feedback conflict thresholds; gain step, relaxation, refractory; head rate	2, 2, 0.1, 10, 10, 0.001	none, none, none, s, s, none	
$I^{\text{lo}}_{\text{Vis}}$, $I^{\text{hi}}_{\text{Vis}}$	vision intensity levels	0.2, 0.7		provisional

Preprint. Not peer-reviewed.

Symbol	Meaning	Value	Units	Status
$I^{\text{lo}}_{\text{Aud}}$, $I^{\text{hi}}_{\text{Aud}}$	audition intensity levels	0.4, 0.8		provisional
$I^{\text{lo}}_{\text{Intero}}$, $I^{\text{hi}}_{\text{Intero}}$	interoception intensity levels	0.1, 0.7		provisional
$I^{\text{lo}}_{\text{Seq}}$, $I^{\text{hi}}_{\text{Seq}}$	sequence intensity levels	0.1, 0.7		provisional
$I^{\text{lo}}_{\text{Aff}}$, $I^{\text{hi}}_{\text{Aff}}$	affect intensity levels	0.1, 0.7		provisional
$I^{\text{lo}}_{\text{Sleep}}$, $I^{\text{hi}}_{\text{Sleep}}$	sleep intensity levels	0.5, 0.8		provisional
$I^{\text{lo}}_{\text{Lang}}$	language-organ utterance intensity	0.4		provisional
m_D, b_D, L_D, N_D	organ repetition penalty $m_D b_D^{n-L_D}$ for a repeated run of $n \geq L_D$ tokens; window	0.8, 1.75, 2, 2048	none; tokens	provisional
	organ length caps (speech, thought); words in a cut repeat	512, 128; 3	tokens; words	provisional
β_Γ	error ratio at which the graded map reaches alert level	2		
W_P	running-ratio window (predictive processors)	32	reports	
β_ν	prediction-error alert ratio (vision, audition)	2.0		
β_c	change alert ratio (vision, audition)	2.0		
$c^{\text{min}}_{\text{Vis}}$, $c^{\text{min}}_{\text{Aud}}$	absolute change floors	10^{-4}, 0.35		
W_{buf}	forward-model input buffer (vision, audition)	16	reports	
	forward-model hidden units (vision; audition)	256; 32	units	
η_{fm}	forward-model learning rate (predictive processors)	10^{-3}		
ω_{min}, ω_{max}	audition attended fraction at arousal 1, 0	0.15, 1.0		
Δ_{att}	shortest attended span (audition)	0.025	s	
	vision saliency grid	12 by 12	tiles	
Θ_F	fovea hysteresis margin	0.15		
F_{min}, F_{max}	fovea size at arousal 1, 0	0.12, 0.5	fraction of frame	
α_F	fovea tile-statistics weight	0.05		
n_F	tile updates before z-scoring	5	clip reports	
k_F	tile deviation floor (fraction of mean tile change)	0.1		
β^{Seq}_ν	sequence temporal-error alert ratio	3.0		
W_z	sequence hidden-state z-score window	64	broadcasts	
W_{recur}	recurrence window	32	broadcasts	
n_{recur}	recurrence count	4	states	
Δ_{recur}	recurrence quantization step	0.25		
$W_{\Delta t}$	sequence interval window	32	broadcasts	
s_{max}	timespan cap (interoception, sequence); entered-interval cap over window mean (sequence)	10		
	reservoir units (interoception; sequence)	32; 32	units	

Preprint. Not peer-reviewed.

Symbol	Meaning	Value	Units	Status
I_0	phasic salience floor	0.5		
τ_{ph}	phasic decay time	1.0	s	
a_0	baseline arousal (affect, access rate)	0.3		
λ	arousal relaxation rate	0.05	s^{-1}	
γ_D	drive gain on the arousal target	0.2		
P_{Aff}	minimum affect-event interval; timer period	1.0	s	
γ_a	perceptual-alert arousal gain	0.15		
ν_{cap}	excess-error-ratio cap per alert	4.0		
γ_{Intero}	arousal step per interoception hard-threshold report	0.05		
τ_{f}, τ_{s}	fast, slow time constants (error and wellness means, alert rates)	10, 100	s	
k_v, τ_v	error-dynamics gain (pleasantness check), valence time constant	2.0, 30	none, s	
lp_0; T_{lp}, Δ_{lp}	learning-progress noise floor; span, bin	0.20; 5400, 60	none; s	calibrated
m_B	alert-excess margin	0.5		calibrated
r_0	alert-rate floor	0.05	s^{-1}	
α_{int}	intent-rate weight	0.05		
β_D, δ_I	curiosity build, decay	0.05, 0.00096	s^{-1}	calibrated
β_D, δ_D, ρ_D; ρ_{eng}	boredom build, decay, relief; relief per engaging intent	0.04, 0.0045, 0.3; 0.5	s^{-1}; none	
β_D, δ_D; ρ_D	social-drive build, decay; relief per interaction	0.01, 0.0011; 0.8	s^{-1}; none	
β_D, δ_D; ρ_D	restlessness build, decay; relief per intent	0.03, 0.0033; 0.5	s^{-1}; none	
θ_D	drive threshold (every drive)	0.7		
Θ_D	drive re-arm fraction	0.9		
Δ_{stall}, n_{goal}, Δ_{resp}, η_L, P_{re}; Δ_{rest}	goal stall horizon, cap; response window, likelihood fall; re-entry interval; developmental rest (§3.5)	600, 3, 10, 0.5, 30; 1410	s, goals, s, none, s; s	
s_{q}; T_{act}	quiescent share of the developmental cycle; awake time per cycle, $\Delta_{\text{rest}}(1-s_{\text{q}})/s_{\text{q}}$ (§3.5)	0.32; 2996	none; s	
θ_E, b_E, τ_E	emotion-regulation set point, tier-1 bias, its decay time	0.5, 0.2, 20	none, none, s	
θ_{wf}, q_{wf}	welfare-indicator bars (A.4)	0.5, 0.95		provisional
	initial held arousal $\hat{a}$	0.3		
	wellness scaling: GPU temperature bounds; cycle-latency target	30 to 80; 300	°C; ms	
P_{Intero}	interoception reading interval	1.0	s	
n_x	interoception feature dimension (three self-rhythm features, zero unless enabled; one is the highest GPU memory use)	8		
τ_U	expected-error time constant	600	s	

Preprint. Not peer-reviewed.

Symbol	Meaning	Value	Units	Status
β_U	expected-error band	2.0		
	hard thresholds: CPU, RAM, GPU temperature, GPU memory, cycle latency	90, 90, 83, 92, 600	%, %, °C, %, ms	
θ_{fat}, δ_{fat}	fatigue threshold, awake decay	100, 0.01	none, s^{-1}	provisional
θ_R	regulation threshold	0.5		provisional
Δ_R	regulation sustain window	30	s	
n_{wu}, Δ_{wu}	warm-up end conditions	1000, 1200	updates, s	
γ_f	processing-rate reduction factor	0.8		
$f_p^{\min}$, $f_p^{\max}$	processing-rate bounds under regulation	0.5, 20	Hz	
θ_{th}	think bar	0.45		provisional
θ_{sp}	speak (report) bar	0.6		provisional
Δ_{th}	think refractory interval	3.0	s	
Δ_{sp}	speak refractory interval	8.0	s	
Δ_{sig}	signature expiry	300	s	
Δ_g	in-flight guard timeout	48	s (wall)	
f_r	self-rhythm step rate	20	Hz	
n_s	Euler substeps per step	10		
J_{rec}	recurrent weight	2.5		calibrated
Ξ_0	input offset	-0.25		calibrated
k_y	sigmoid slope scale	0.08		calibrated
τ_y	activity time constant	0.02	s	calibrated
γ_R	depression rate	8.0	s^{-1}	calibrated
γ_o, γ_u	own-drive, maternal-drive gains	0.02, 0.03		calibrated
σ_ξ	noise scale	0.02		calibrated
$\tau_{\text{rec}}^{(0)}$	initial recovery time	2.1	s	calibrated
$\tau_{\text{rec}}^{\min}$, $\tau_{\text{rec}}^{\max}$	recovery-time bounds	0.9, 2.3	s	calibrated
η_ℓ	frequency-adaptation rate	0.0025	s^{-1}	calibrated
τ_m	moving-average time constant	10	s	
Δ_{amp}	phase and amplitude window	8	s	
	initial activity y, resource R	0.05, 1.0		
	maternal beat rate, drift	70, 0.03	bpm, none	
$M_{\max}$	maximum maternal drive	0.4		
L_v, j_v; Δ_v, p_v, r_v; Δ_u	voice level, jitter; slot, talk probability, repeat share; unlock interval	6, 3; 30, 0.2, 0.25; 3600	dB; s (wall), none, none; s (awake)	

Preprint. Not peer-reviewed.

Symbol	Meaning	Value	Units	Status
W_{sat}; $d_{\max}$	saturation window; stage maximum dwell	600; 25200	s (awake)	provisional; calibrated
p_m, p_x; p_{st}, $\Delta_{\text{st}}^{\text{lo}}$, $\Delta_{\text{st}}^{\text{hi}}$, Δ_{st}	melody, outside-voice slot probabilities; stir answer probability, delay bounds, refractory	0.02, 0.05; 0.5, 2, 10, 300	none; none, s, s, s	provisional
	color saturation in stage k	$k/5$	none	design
c_M^{base}, c_M^{pert}	usual, perturbation drive scales	0.5, 0.75		
f_g	readout sampling rate	10	Hz	
f_{lo}, f_{hi}	entrainment band	0.3, 2.0	Hz	
Δ_{trim}	edge trim	2.0	s	
Δ_E	driven-window length	300	s	
Δ_W, P_W	withdrawal duration, period	20, 1800	s	
Δ_{pert}, P_{pert}	perturbation duration, period	5, 3600	s	
j_P	probe jitter fraction	0.25		
P_{ro}	readout period; settling time after boot or thaw	60	s	
f_{guard}	pull guard	0.05	Hz	
pull_{min}	pull floor	0.5		calibrated
n_{base}	baseline withdrawals	3	count	
	consecutive passes required	3	count	calibrated
t_{R0}, t_{R1}, t_{R2}, t_{R3}	viability checkpoints	6, 24, 48, 60	h (awake)	calibrated
pull_{R1}, slope_{R1}	R1 bounds on median pull, slope	0.12, 0.002	none, h^{-1} (awake)	calibrated
pull_{R2}	R2 bound on median pull	0.3		calibrated
Δ_V, n_V	viability window, minimum points	12, 8	h (awake), withdrawals	calibrated
Δ_{per}	period-variability window	300	s	
$V_{\text{per}}^{\min}$	period-variability floor	0.02		
$\text{err}_{\text{Vis}}^{\max}$	vision error-ratio ceiling	0.3		provisional
$\Delta_{\max}^{\text{rec}}$	return-to-baseline ceiling	30	s	provisional
tol_{rec}, $\Delta_{\text{cap}}^{\text{rec}}$	recovery tolerance, cap	0.25, 300	none, s	
n_{sleep}, n_{cons}	required sleeps, consolidation passes	5, 3	counts	
	gate cadence	60	s	
Δ_{budget}	developmental-phase budget, derived in A.7	96	h (wall)	calibrated
c_{eff}, n_{cal}	minimum-effect multiple, paired pilot runs	0.5, 30	none, runs	
ϕ_{min}	least on-arm share of broadcasts with $\lvert E_k \rvert > K$ (A.8)	set in advance		provisional

Preprint. Not peer-reviewed.

Symbol	Meaning	Value	Units	Status
n_{run}	runs per arm and tone variant	$\lceil 102/(1 - r_{\text{NE}}) \rceil$ planned (A.8)	runs	
q_{pow}	assumed P(dIG > minimum effect), power analysis	0.65		
n_{side}	distinct scores per fitted side (contrast at threshold)	10	broadcasts	
κ_{max}	condition-number cap per fitted side (A.8)	10^8		
α_{test}	level of each live sign test	0.05		
α_{FW}	family-wise level (offline suite)	0.05		
tol_{CV}, n_{seed}	stability tolerance, seeds (offline suite, coherence-layer ablation)	0.05, 3	none, seeds	
ϱ	daily strength an unreplayed trace keeps (A.10)	0.82		
N_{day}	traces per entity day (A.10)	288,000	traces	
h_{min}	strength pruning floor (A.10)	0.01		
B; B_j^{hk}, Δ_{hk}; R^{hk}	episodic-memory storage budget (A.10); housekeeping budget per resource, time allowance; free-space reserve per data volume (A.5)	set by the operator	bytes; bytes, s; bytes	operator
n_{rep}, c_{aff}, c_{rec}	traces replayed per sleep; replay-score weights	5, 0.7, 0.3	traces, none, none	
n_{pair}, sim_{min}, sim_{dup}, n_{sch}, n_{link}	pairs co-replayed per sleep; similarity band; links forming a schema node, per trace	5, 0.5, 0.95, 3, 8	pairs, none, none, links, links	

Reachability at resting arousal. With Table A1's values and no interest ($G = 1$), the level gain at rest is 0.44 and the contrast gain zero, so access needs a priority of about 0.80, which only audition alerts (0.8) and a failed-phase sleep summary have; other alerts need arousal of about 0.37, and an audition alert needs about 0.45 for the think bar and 0.63 for the speak bar. A deficit at or above θ_D sets a_∞ to 0.40 to 0.43; γ_I is the smallest round gain at which full interest and $\text{lp}^+ = 0.5$ lift a 0.7 alert to the think bar there.

Preprint. Not peer-reviewed.